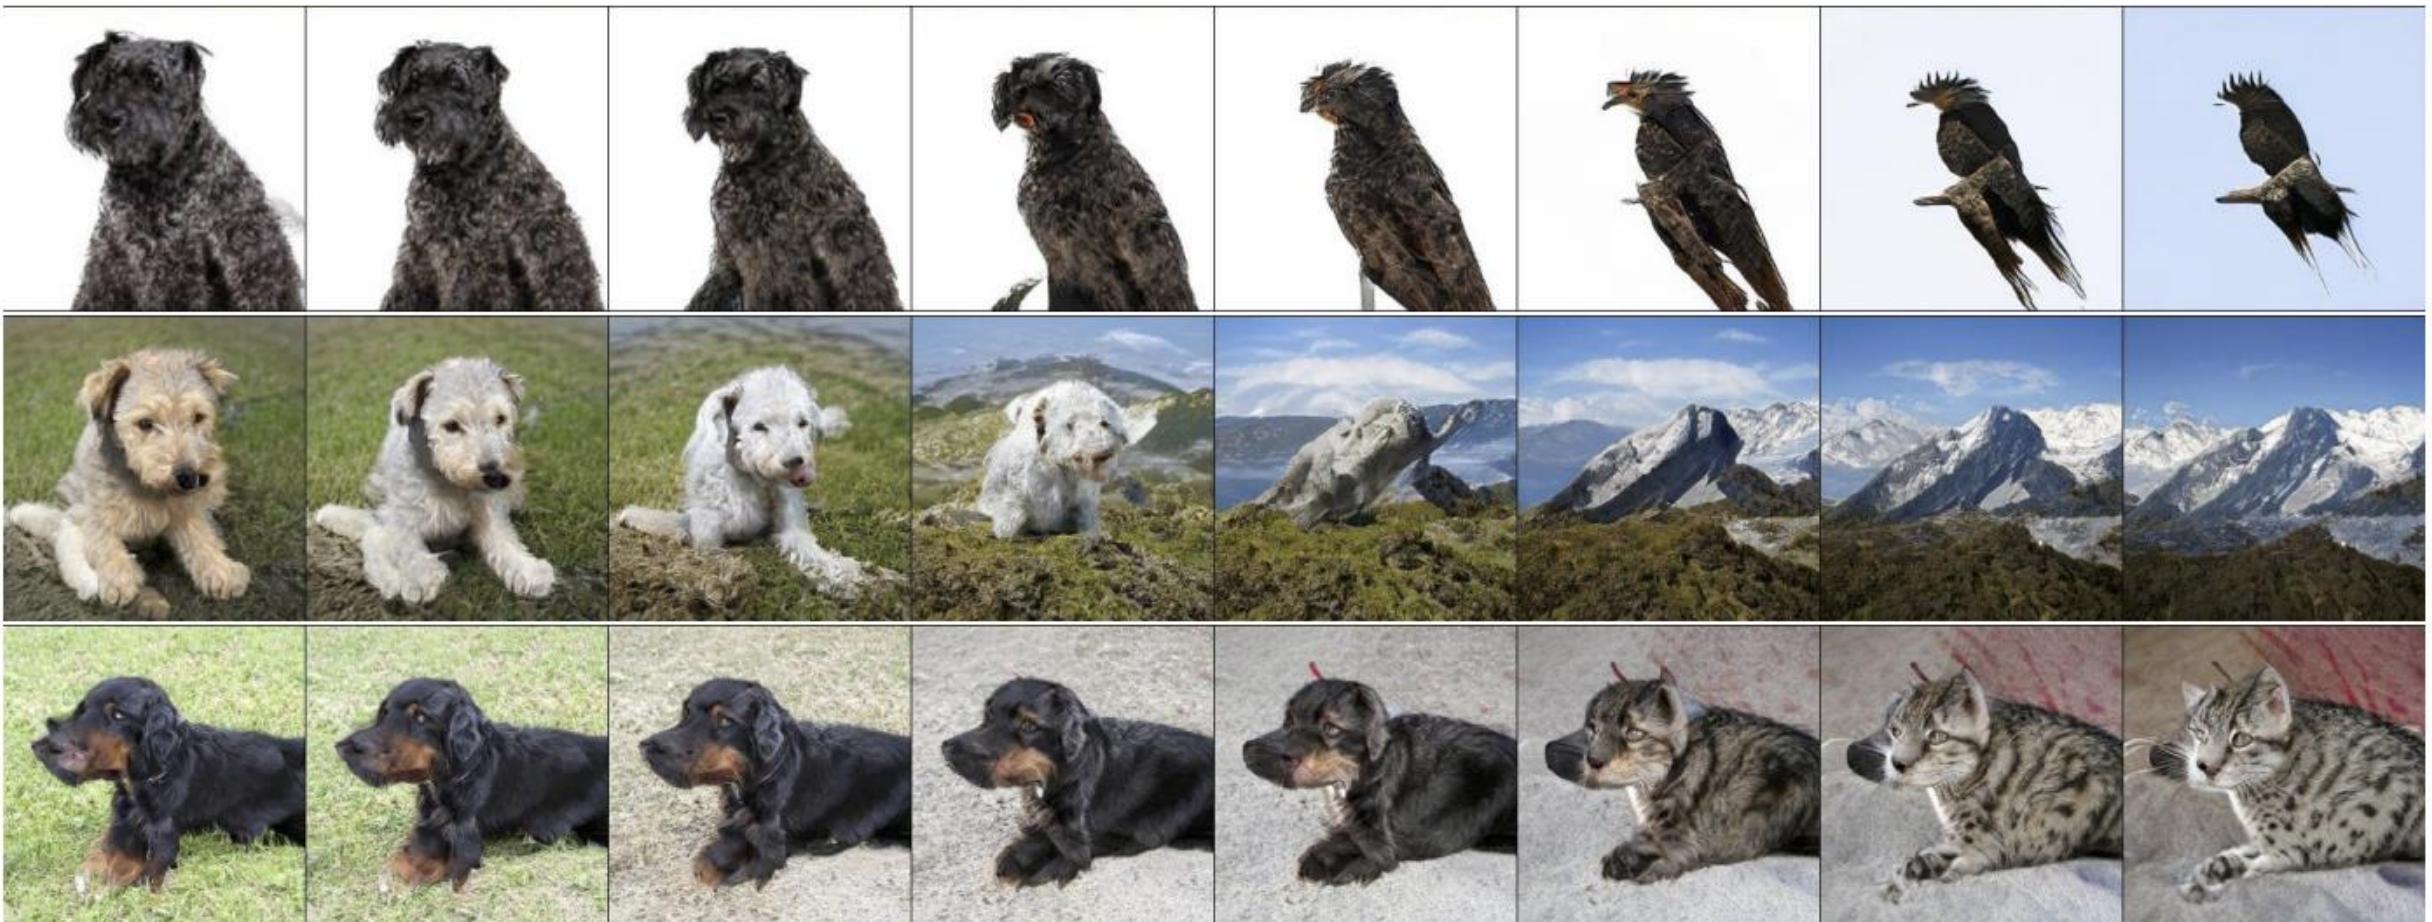


Generative adversarial networks



[BigGAN](#) (2018)

Outline

- Unsupervised Learning
- Generative tasks
- Original GAN formulations
 - NSGAN
 - DCGAN
- Other popular formulations
 - WGAN, WGAN-GP
 - LSGAN
- State-of-the-art architectures
 - Progressive GAN, StyleGAN
- Evaluating GANs
- Visualizing and controlling GANs

Supervised vs Unsupervised Learning

Supervised Learning

Data: (x, y)

x is data, y is label

Goal: Learn a function to map $x \rightarrow y$

Examples: Classification,
regression, object detection,
semantic segmentation,
image captioning, etc.

Supervised vs Unsupervised Learning

Supervised Learning

Data: (x, y)

x is data, y is label

Goal: Learn a function to map $x \rightarrow y$

Examples: Classification,
regression, object detection,
semantic segmentation,
image captioning, etc.



→ **Cat**

Classification

Supervised vs Unsupervised Learning

Supervised Learning

Data: (x, y)

x is data, y is label

Goal: Learn a function to map $x \rightarrow y$

Examples: Classification,
regression, object detection,
semantic segmentation,
image captioning, etc.



DOG, DOG, CAT

Object Detection

Supervised vs Unsupervised Learning

Supervised Learning

Data: (x, y)

x is data, y is label

Goal: Learn a function to map $x \rightarrow y$

Examples: Classification,
regression, object detection,
semantic segmentation,
image captioning, etc.



GRASS, CAT,
TREE, SKY

Semantic Segmentation

Supervised vs Unsupervised Learning

Supervised Learning

Data: (x, y)

x is data, y is label

Goal: Learn a function to map $x \rightarrow y$

Examples: Classification,
regression, object detection,
semantic segmentation,
image captioning, etc.



A cat sitting on a suitcase on the floor

Image Captioning

Supervised vs Unsupervised Learning

Unsupervised Learning

Data: x

Just data, no labels!

Goal: Learn some underlying
hidden structure of the data

Examples: Clustering,
dimensionality reduction, feature
learning, density estimation, etc.

Supervised vs Unsupervised Learning

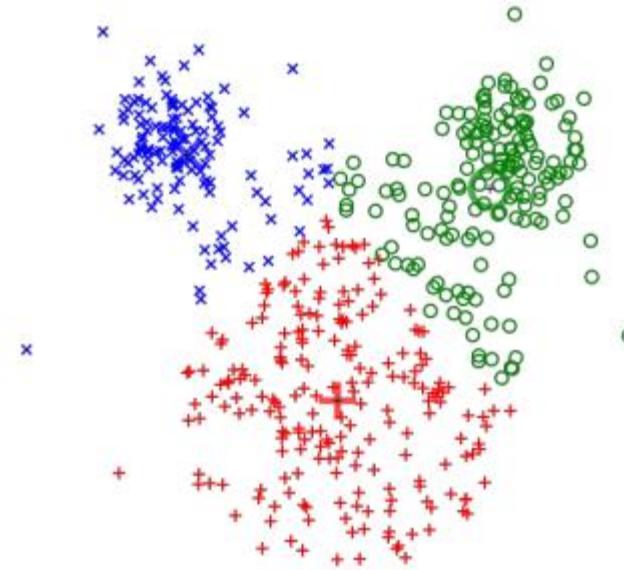
Unsupervised Learning

Data: x

Just data, no labels!

Goal: Learn some underlying hidden structure of the data

Examples: Clustering, dimensionality reduction, feature learning, density estimation, etc.



K-Means Clustering

Supervised vs Unsupervised Learning

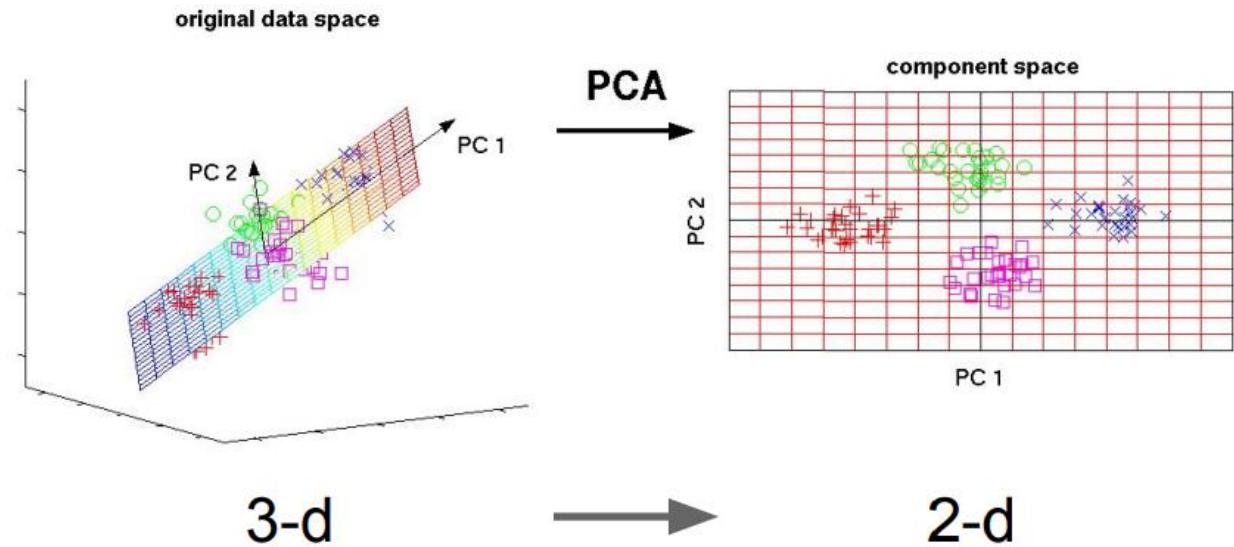
Unsupervised Learning

Data: x

Just data, no labels!

Goal: Learn some underlying hidden structure of the data

Examples: Clustering, dimensionality reduction, feature learning, density estimation, etc.



(Principal Component Analysis)
Dimensionality Reduction

Supervised vs Unsupervised Learning

Unsupervised Learning

Data: x

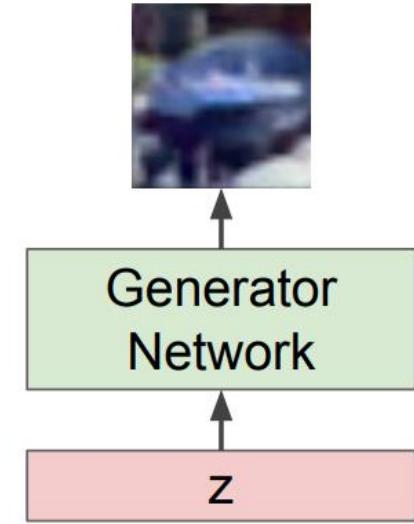
Just data, no labels!

Goal: Learn some underlying hidden structure of the data

Examples: Clustering, dimensionality reduction, feature learning, density estimation, etc.

Output: Sample from training distribution

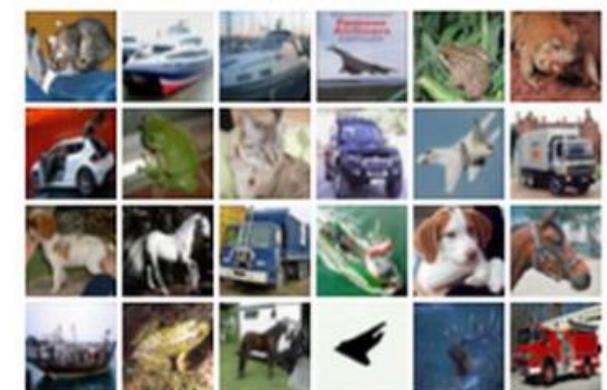
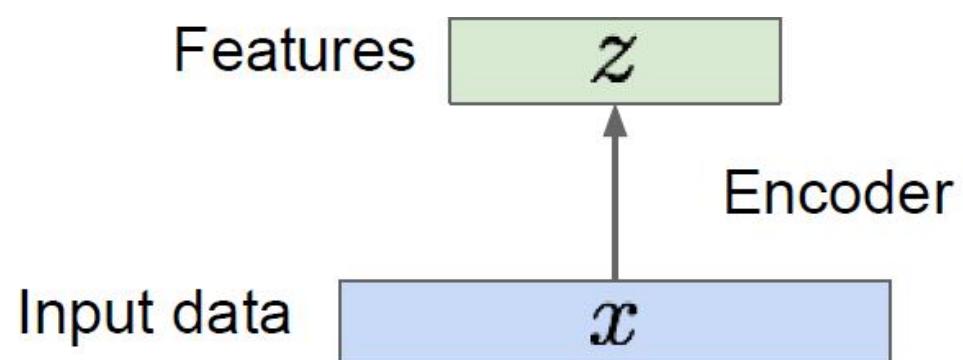
Input: Random noise



**Generative Adversarial Networks
(Distribution learning)**

Autoencoders

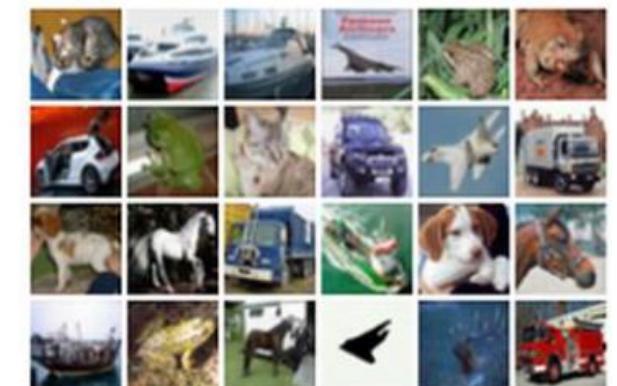
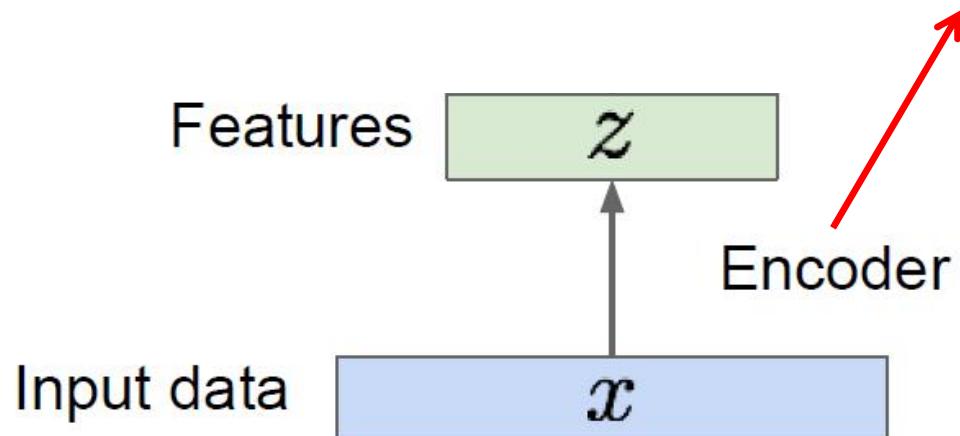
Unsupervised approach for learning a lower-dimensional feature representation from unlabeled training data



Autoencoders

Unsupervised approach for learning a lower-dimensional feature representation from unlabeled training data

Originally: Linear +
nonlinearity (sigmoid)
Later: Deep, fully-
connected
Later: ReLU CNN

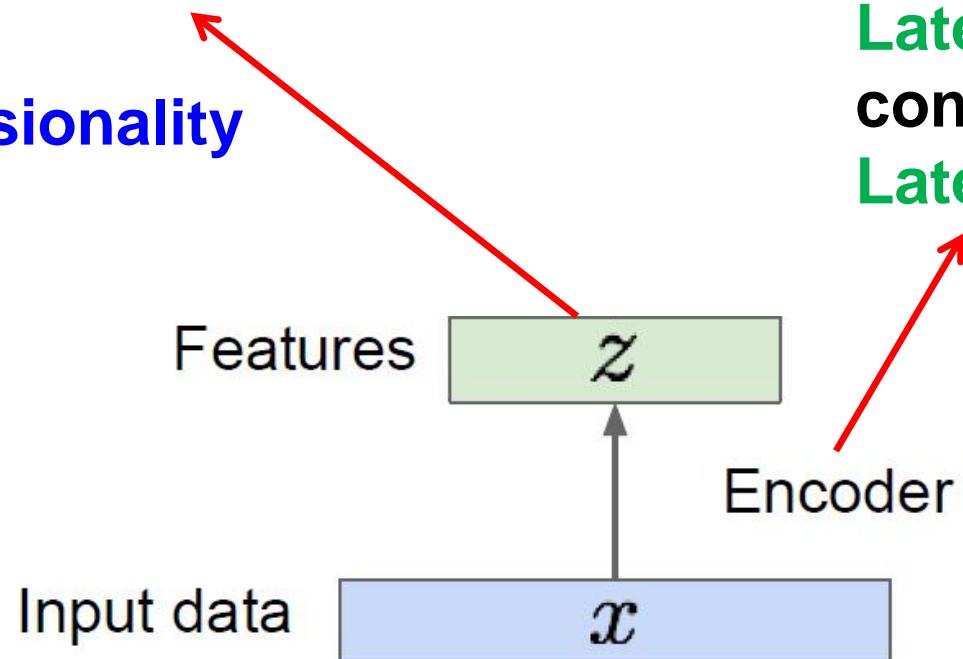


Autoencoders

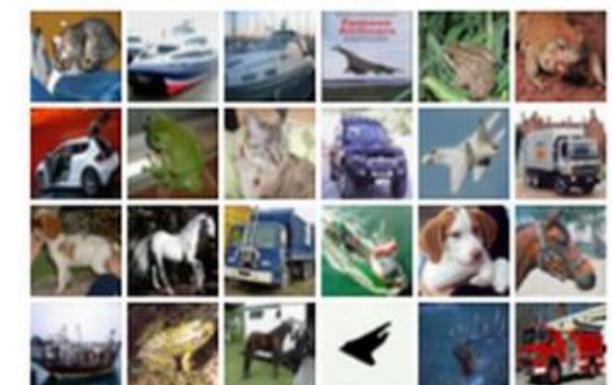
Unsupervised approach for learning a lower-dimensional feature representation from unlabeled training data

**Z usually smaller than X
(Dimensionality Reduction)**

Q: Why dimensionality reduction?



Originally: Linear +
nonlinearity (sigmoid)
Later: Deep, fully-
connected
Later: ReLU CNN



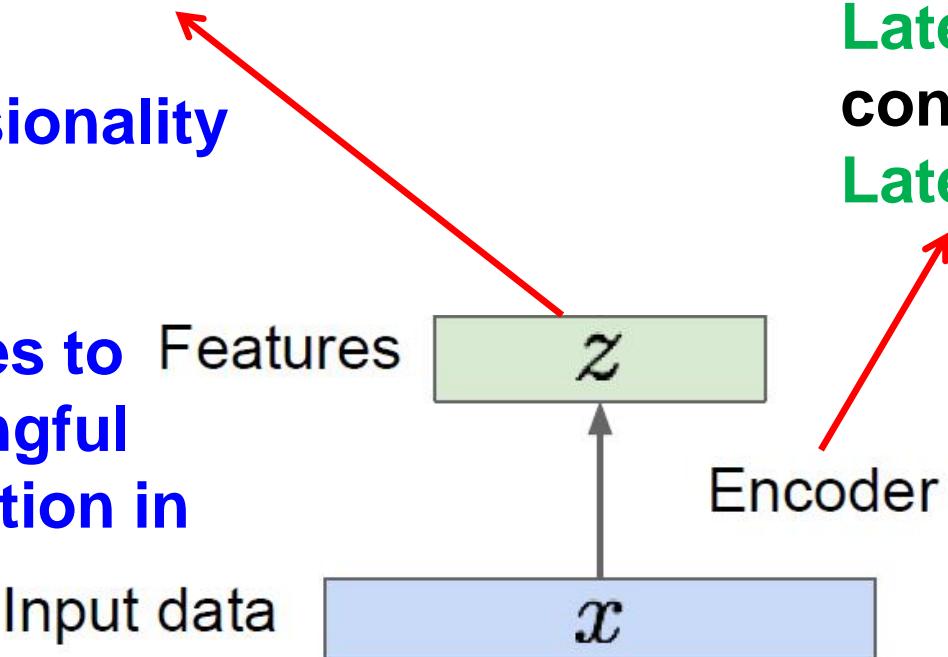
Autoencoders

Unsupervised approach for learning a lower-dimensional feature representation from unlabeled training data

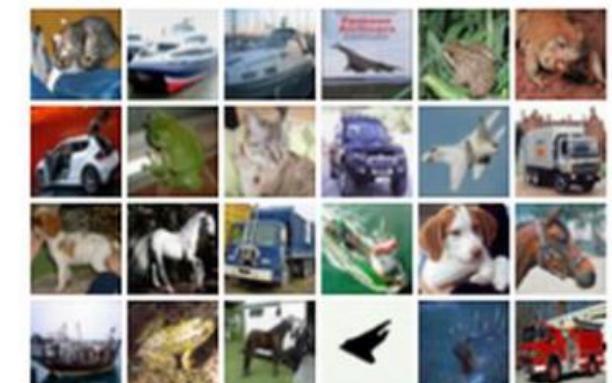
**Z usually smaller than X
(Dimensionality Reduction)**

Q: Why dimensionality reduction?

A: Want features to capture meaningful factors of variation in data

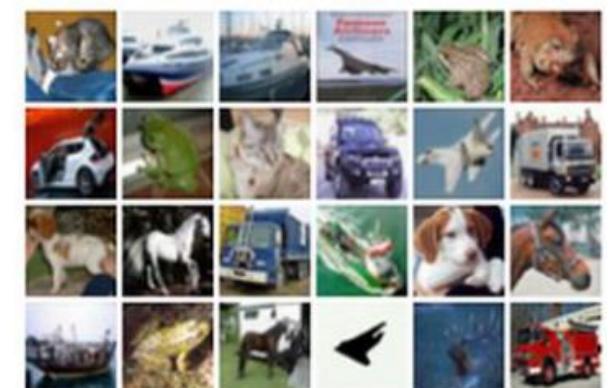
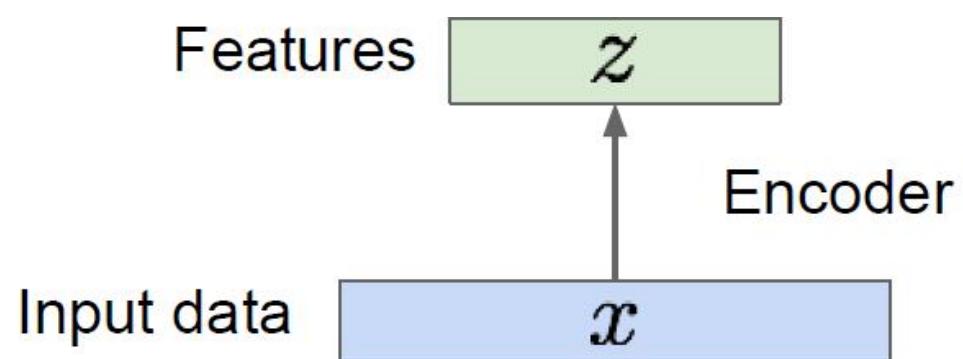


Originally: Linear + nonlinearity (sigmoid)
Later: Deep, fully-connected
Later: ReLU CNN



Autoencoders

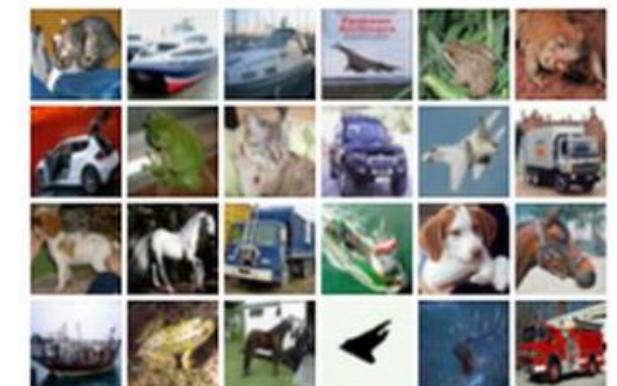
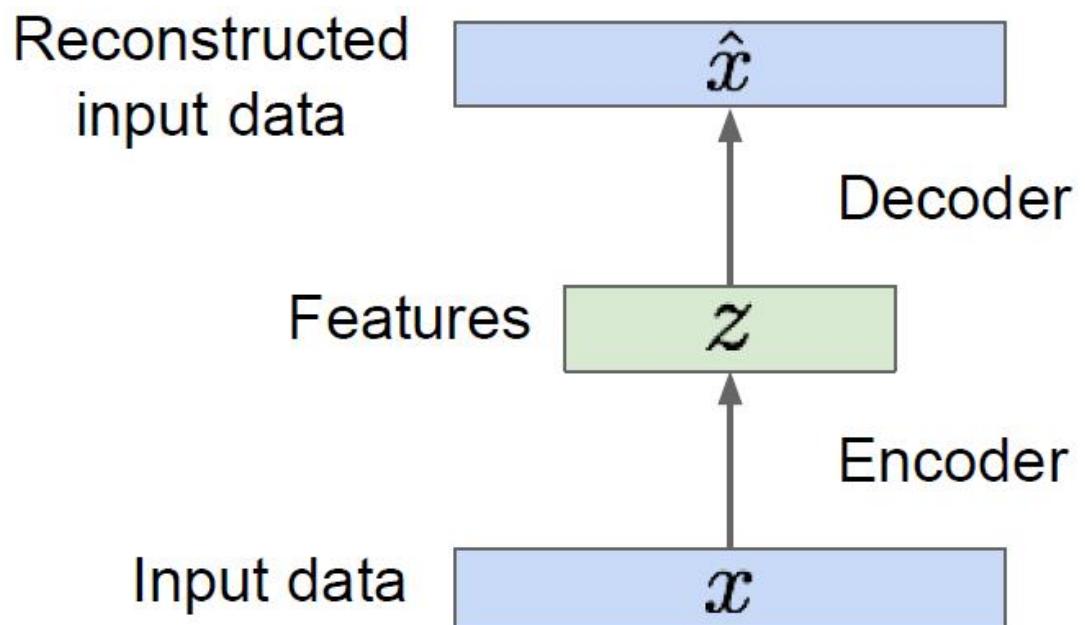
How to learn this feature representation?



Autoencoders

How to learn this feature representation?

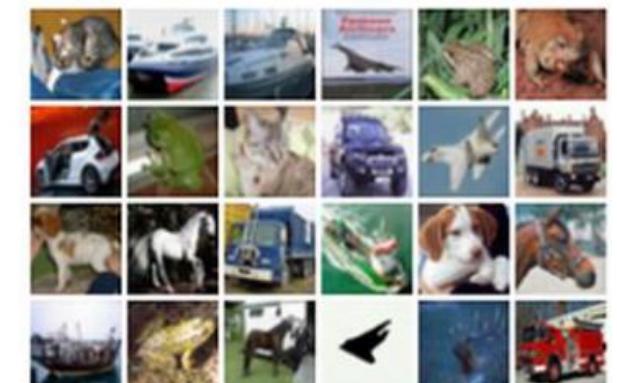
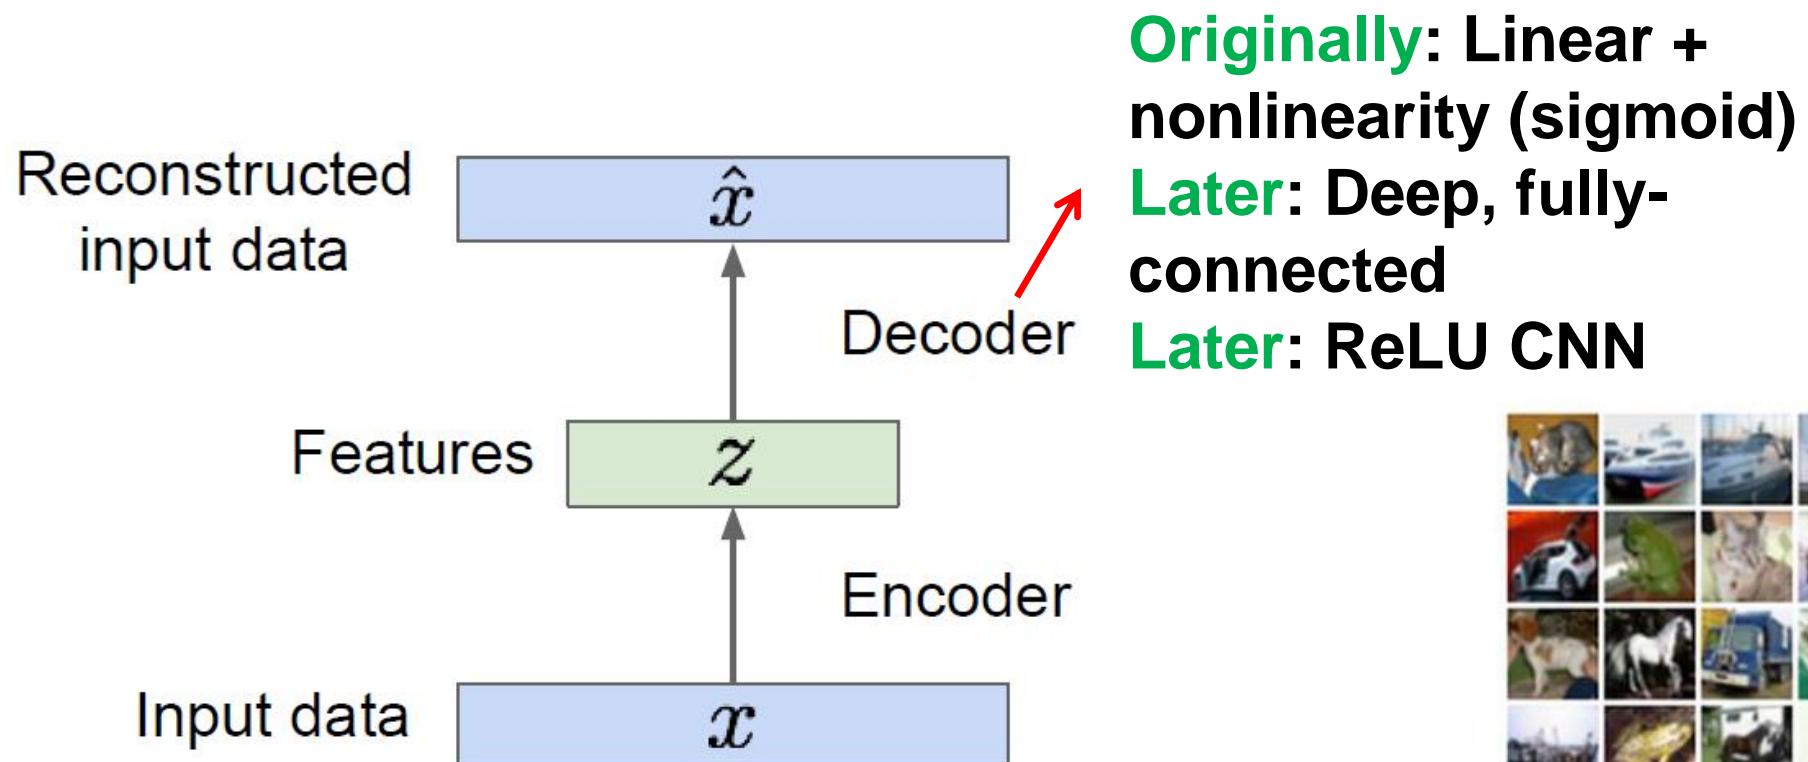
Train such that features can be used to reconstruct original data “Autoencoding” - encoding itself



Autoencoders

How to learn this feature representation?

Train such that features can be used to reconstruct original data “Autoencoding” - encoding itself



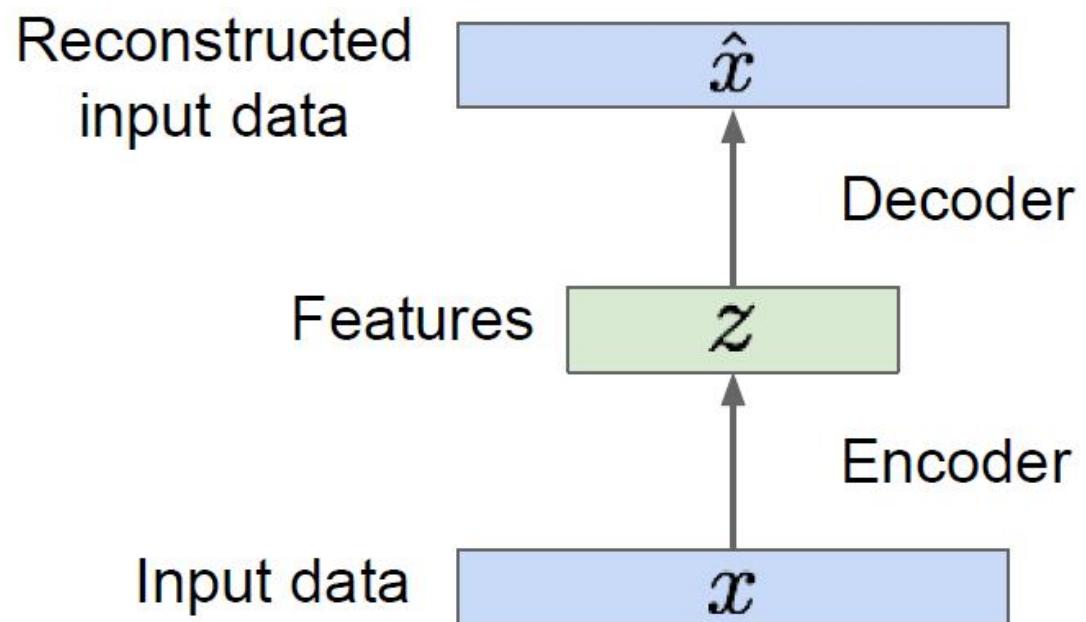
Autoencoders

Reconstructed Data

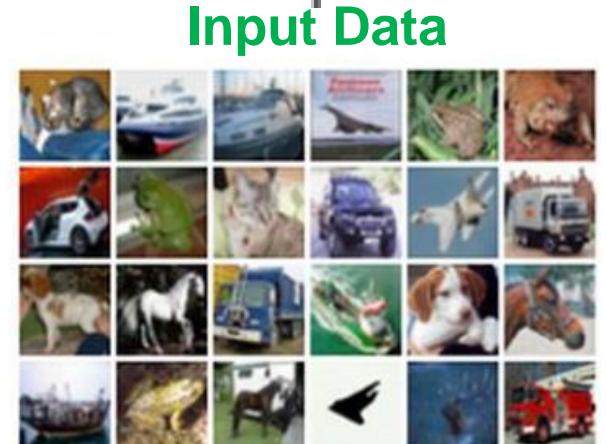


How to learn this feature representation?

Train such that features can be used to reconstruct original data “Autoencoding” - encoding itself



Decoder: 4-layer upconv
Encoder: 4-layer conv

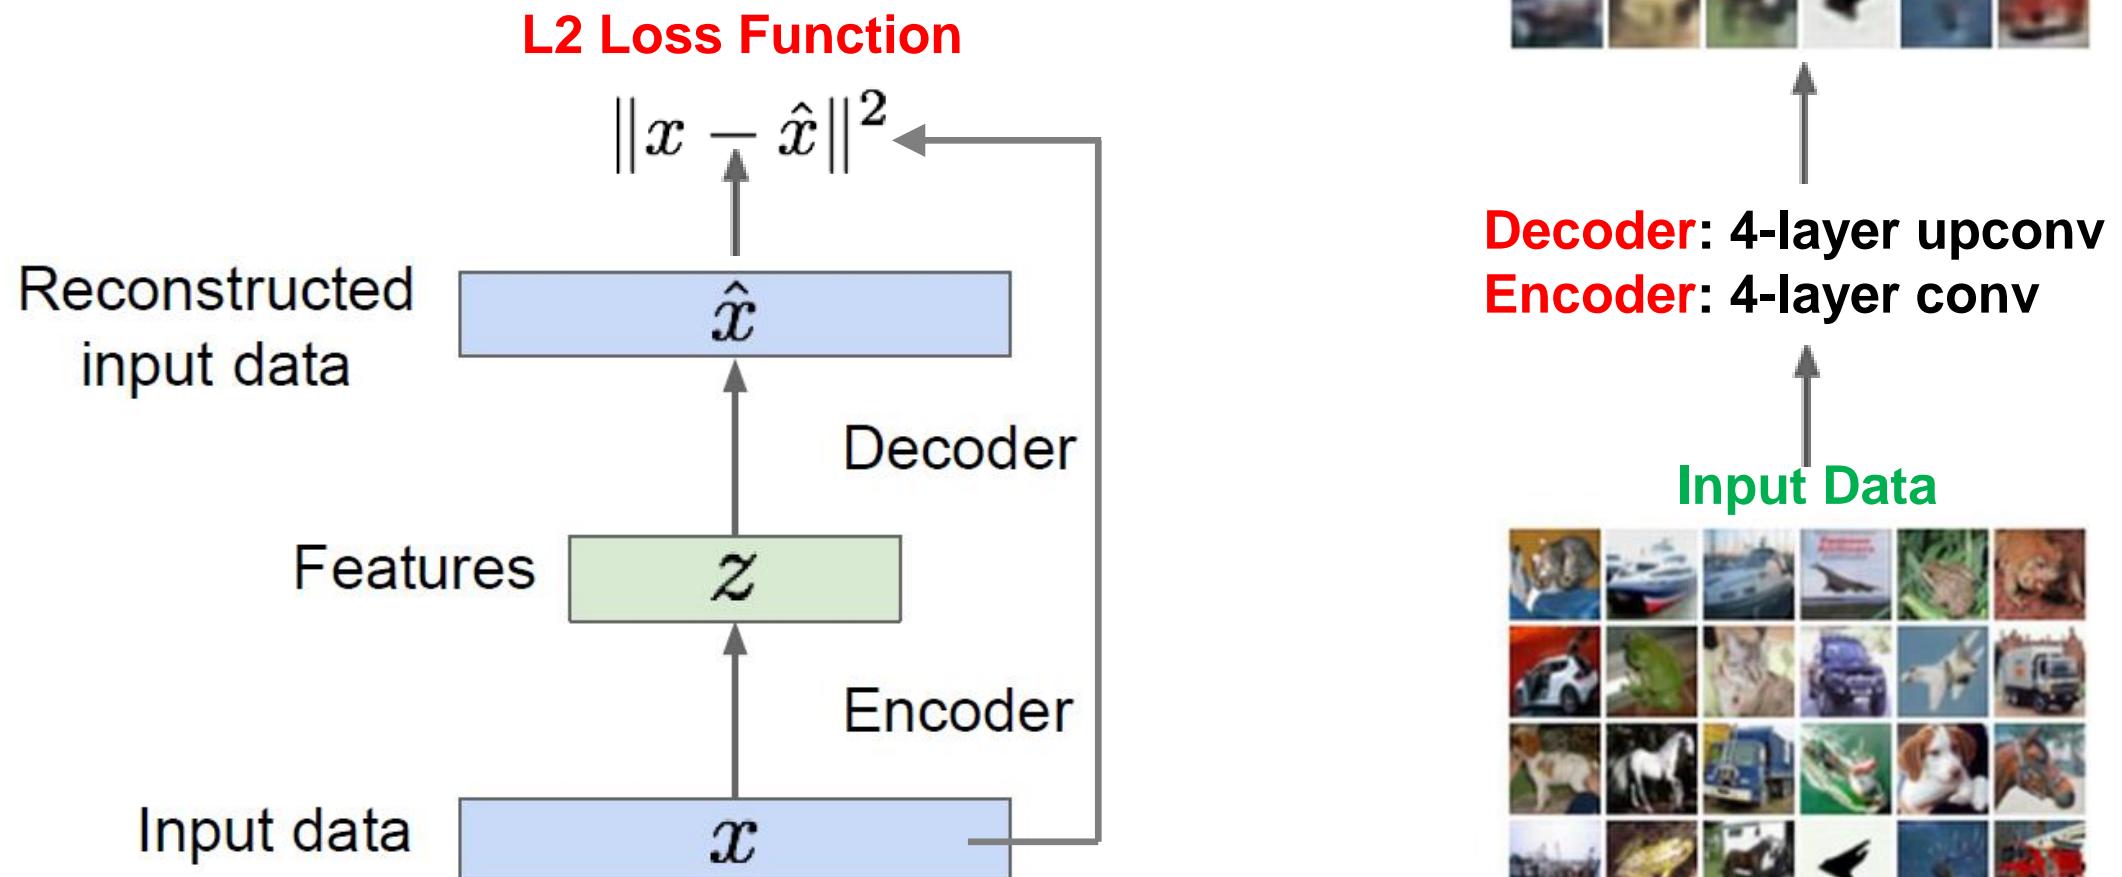


Autoencoders

Reconstructed Data



Train such that features can be used to reconstruct original data



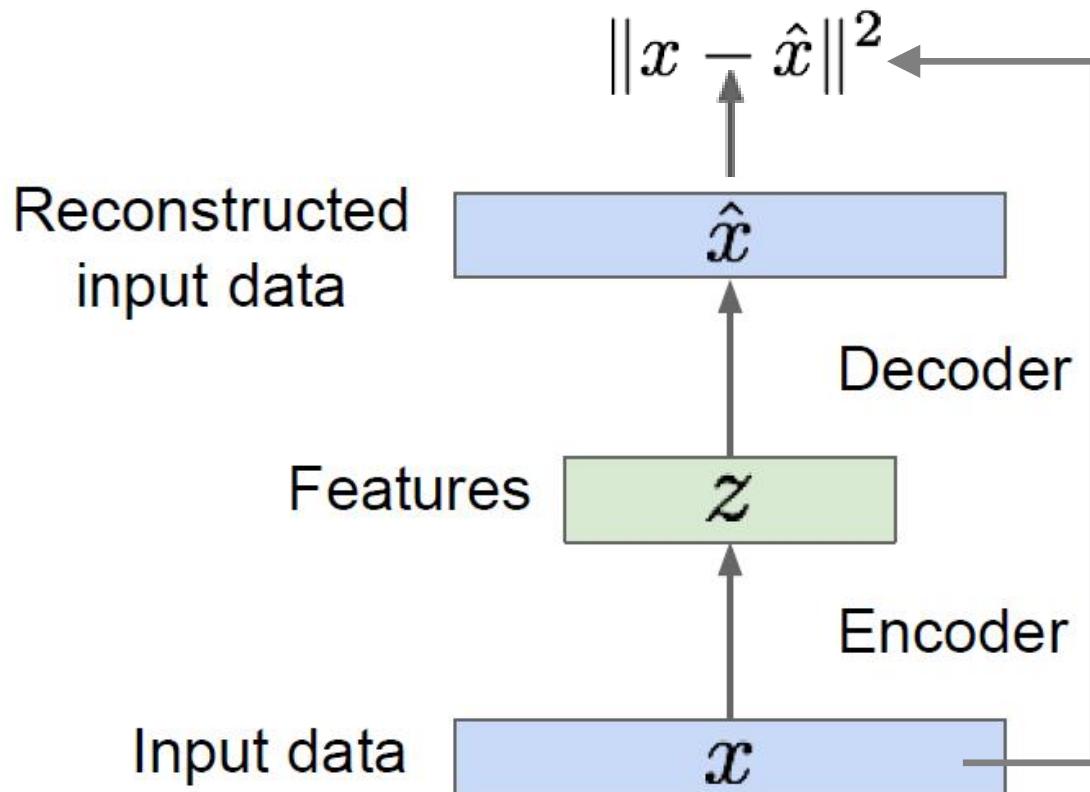
Autoencoders

Reconstructed Data



Doesn't use labels!

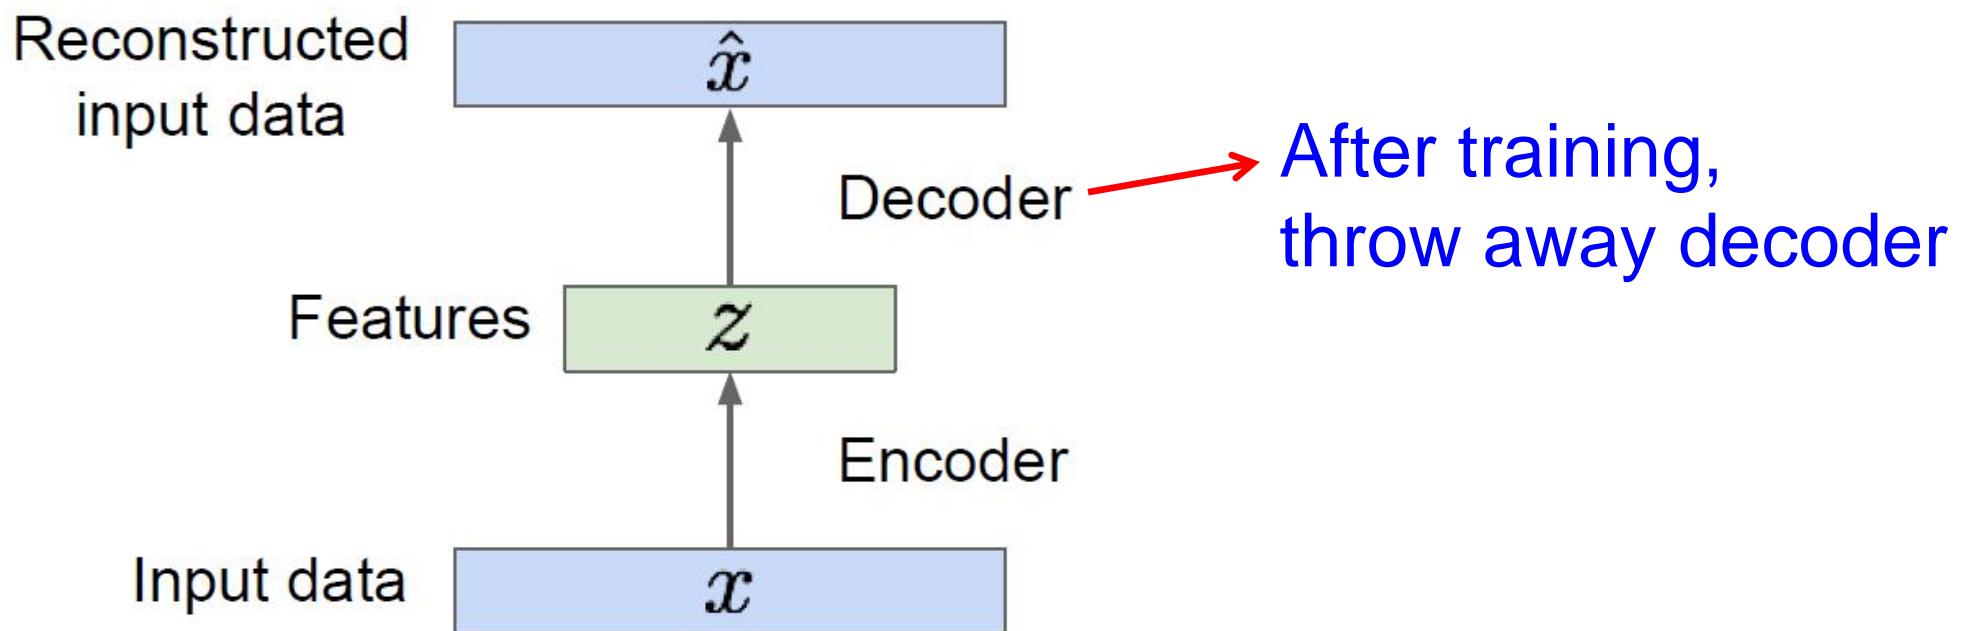
L2 Loss Function



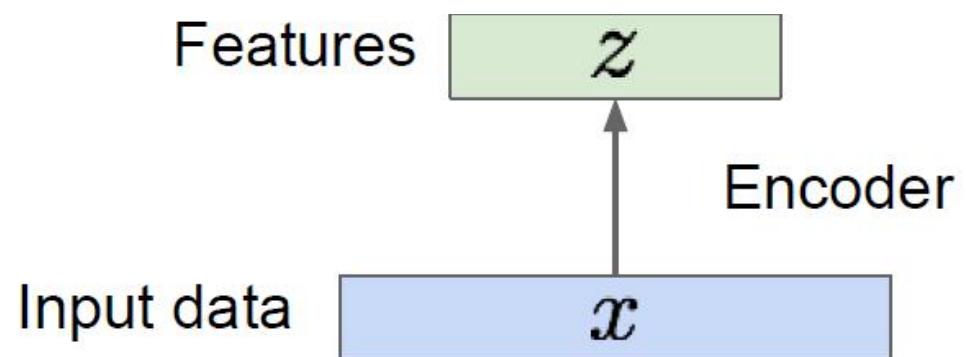
Decoder: 4-layer upconv
Encoder: 4-layer conv



Autoencoders

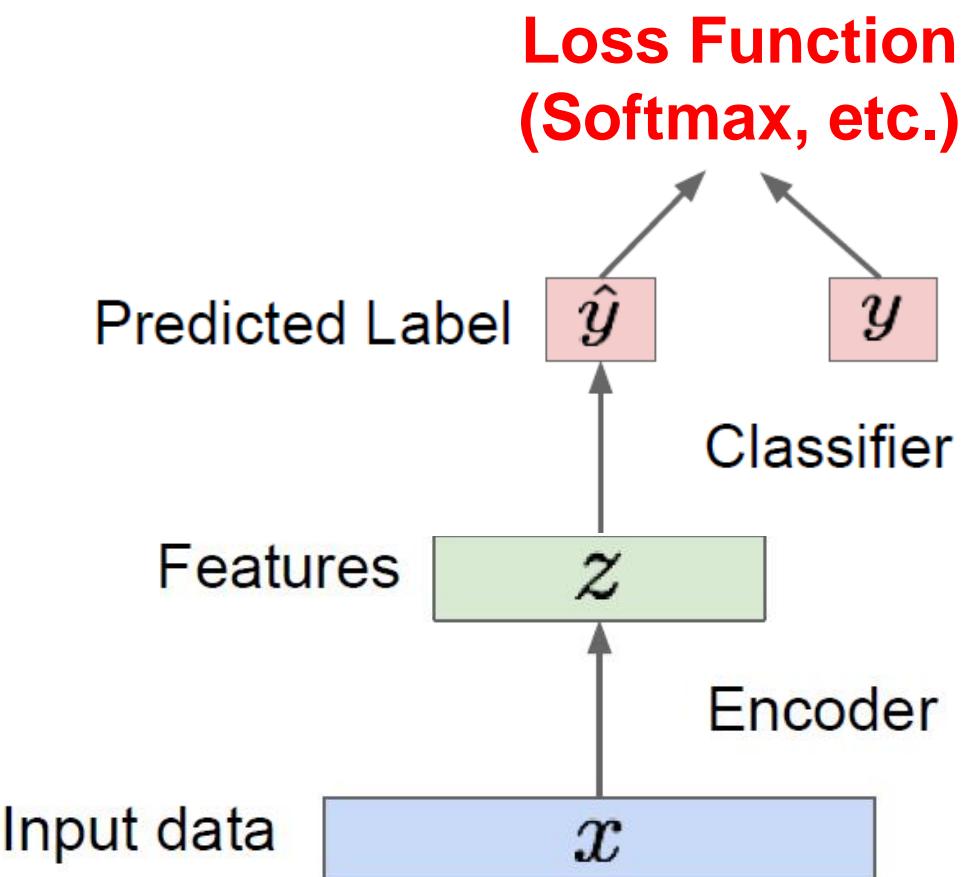


Autoencoders



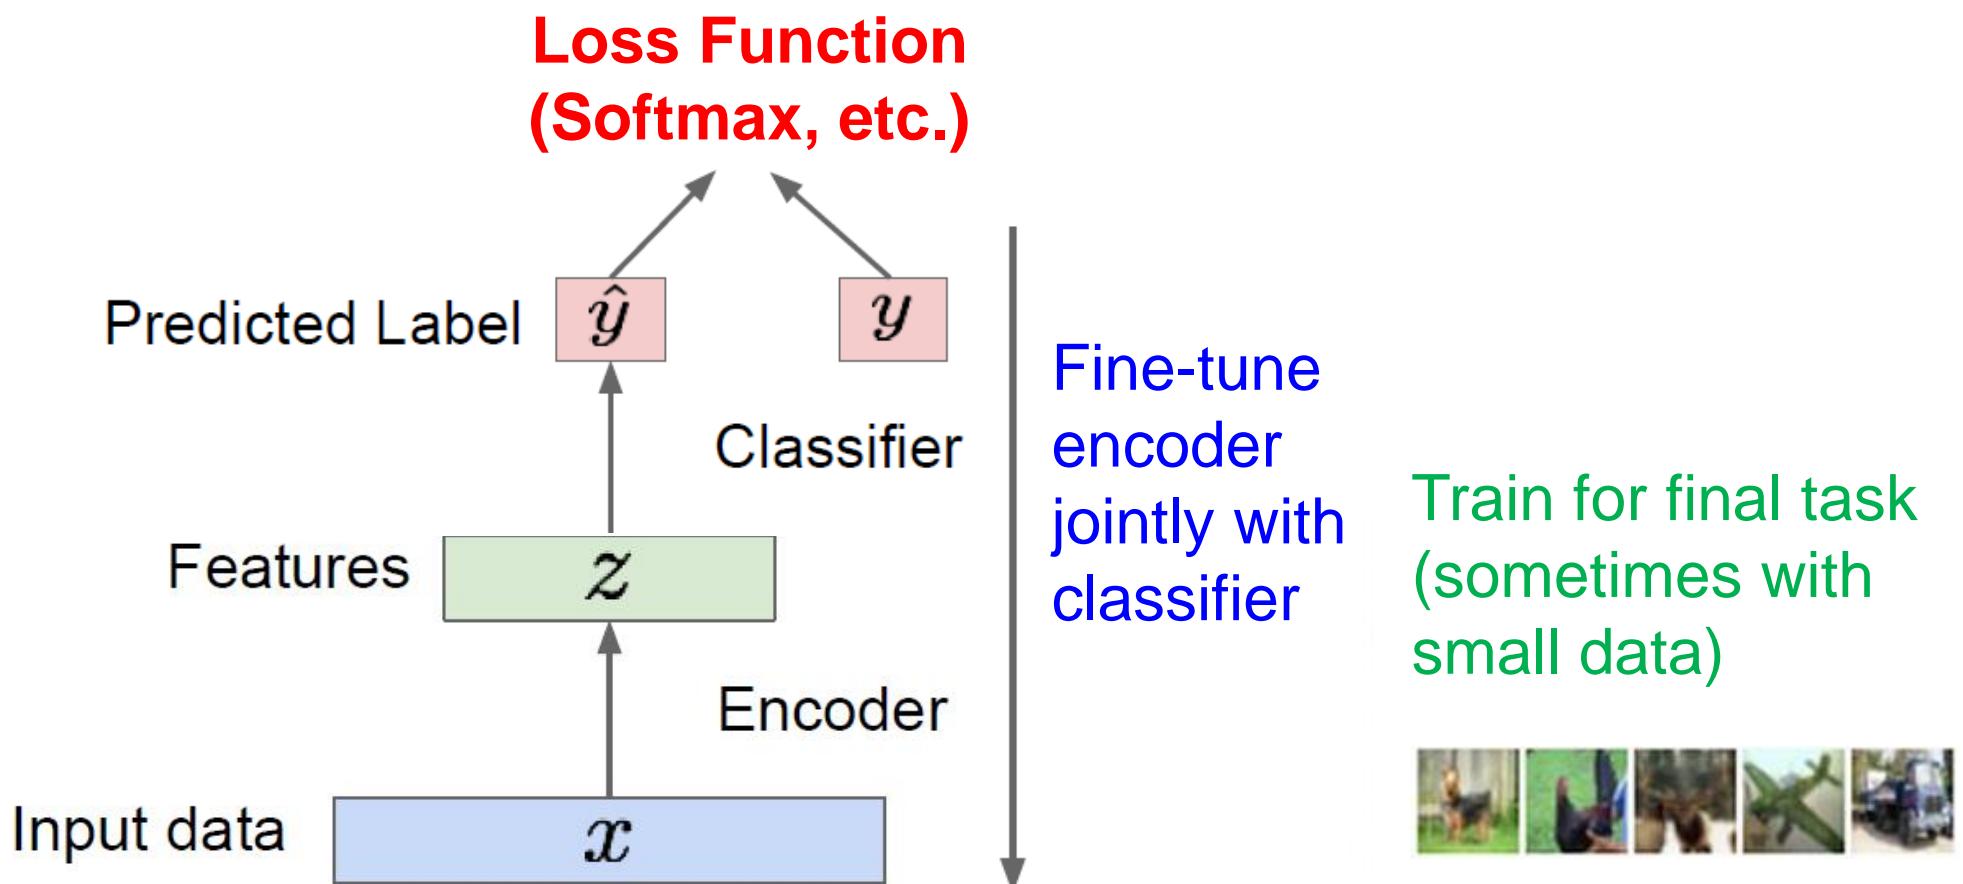
Autoencoders

Encoder can be used to initialize a supervised model



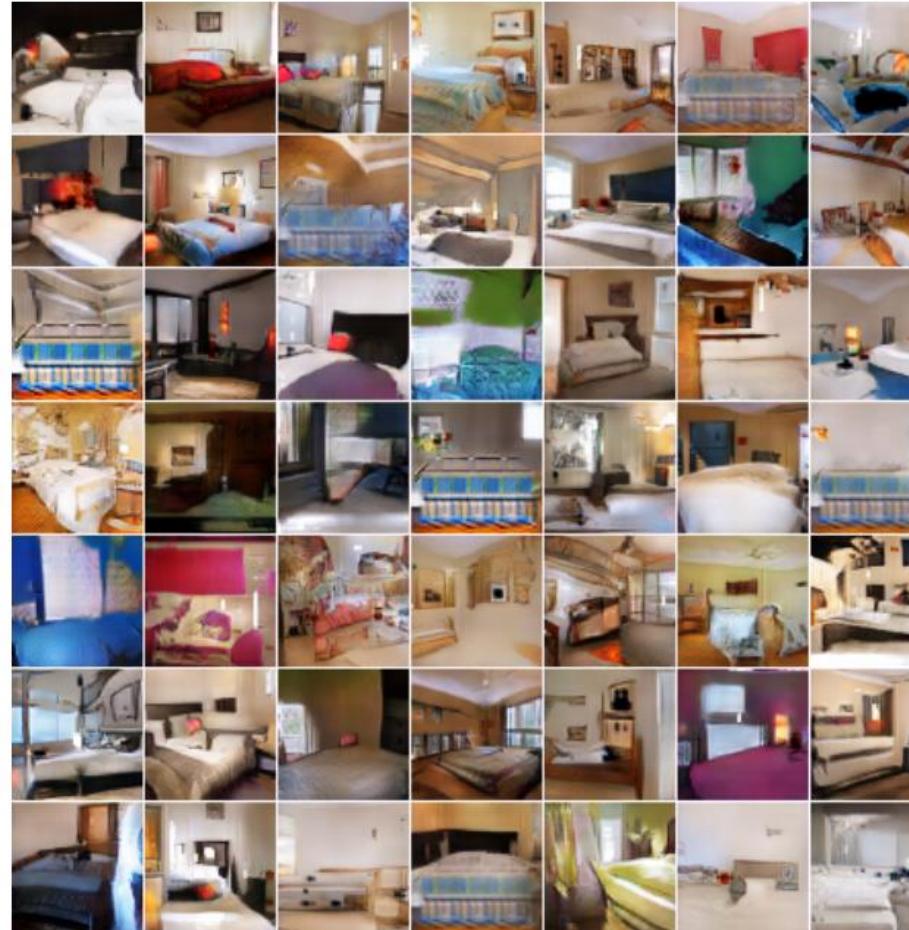
Autoencoders

Encoder can be used to initialize a supervised model



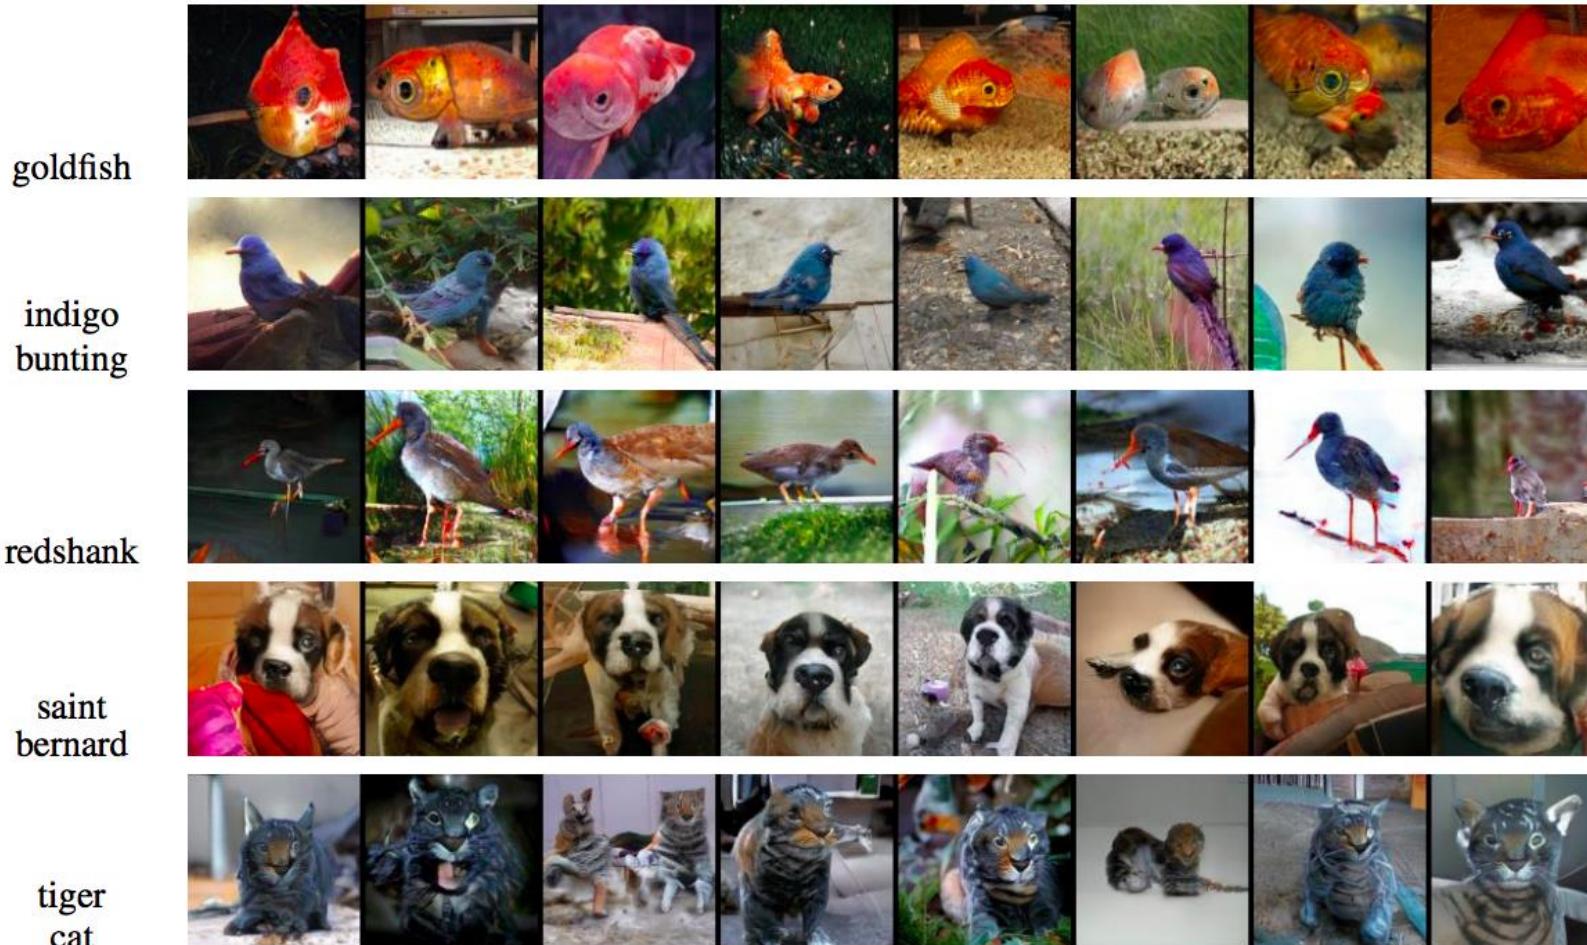
Generative tasks

- Generation (from scratch): learn to sample from the distribution represented by the training set
 - *Unsupervised learning* task



Generative tasks

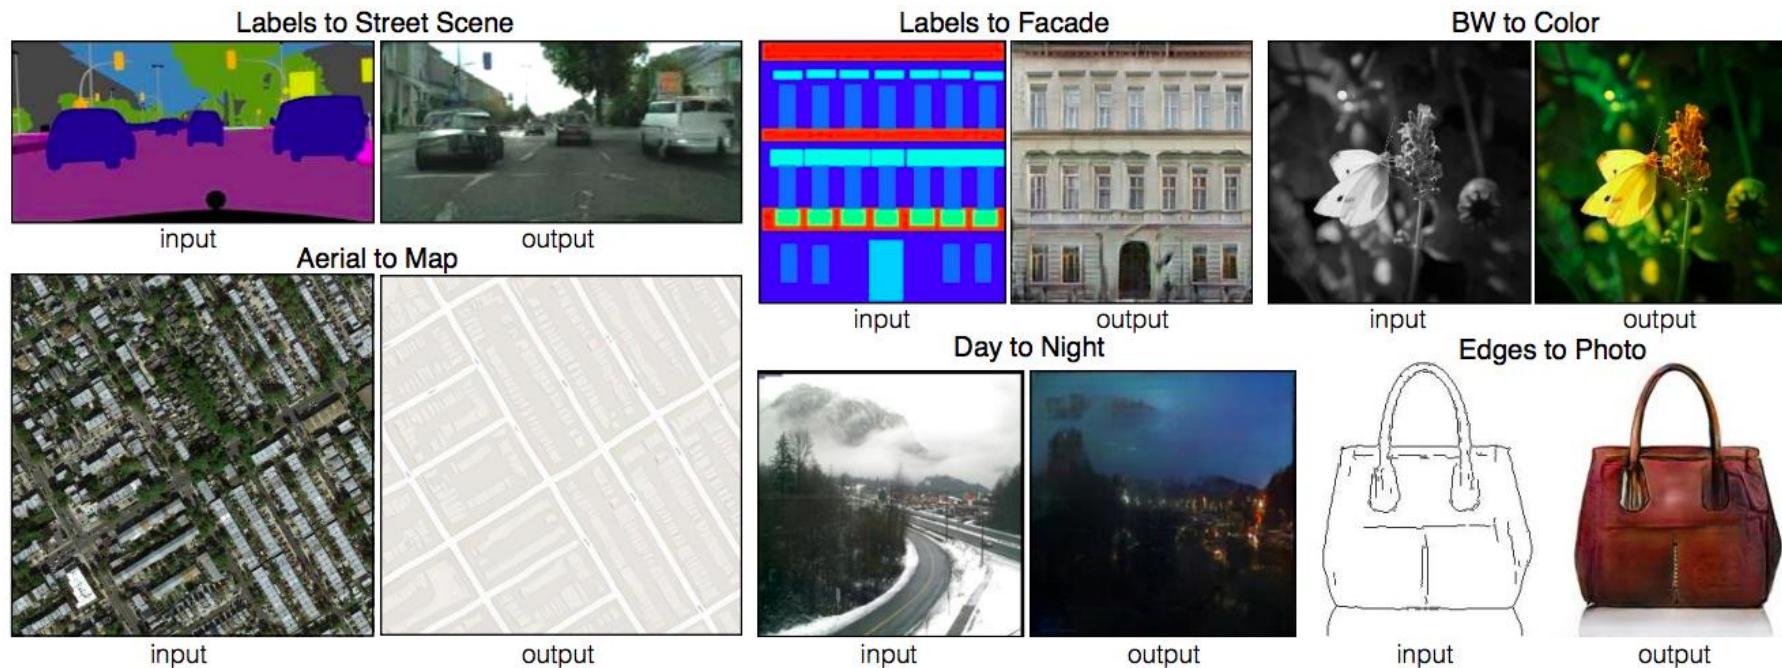
- Generation conditioned on class label



[Figure source](#)

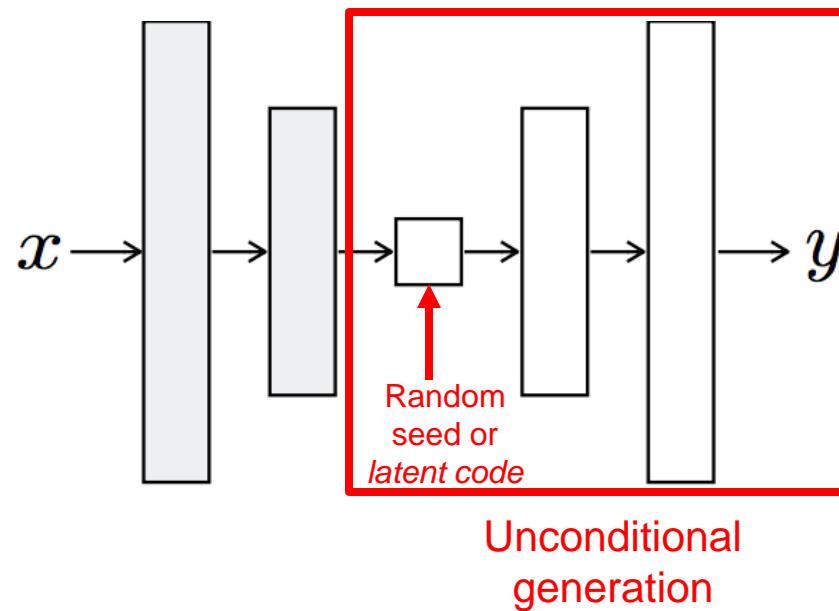
Generative tasks

- Generation conditioned on image (image-to-image translation)



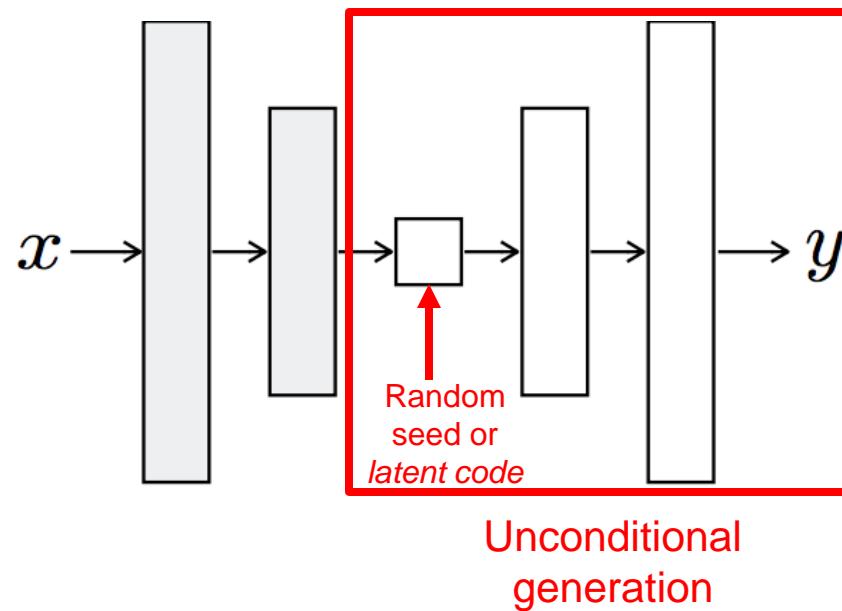
Designing a network for generative tasks

1. We need an architecture that can generate an image
 - Recall upsampling architectures for dense prediction



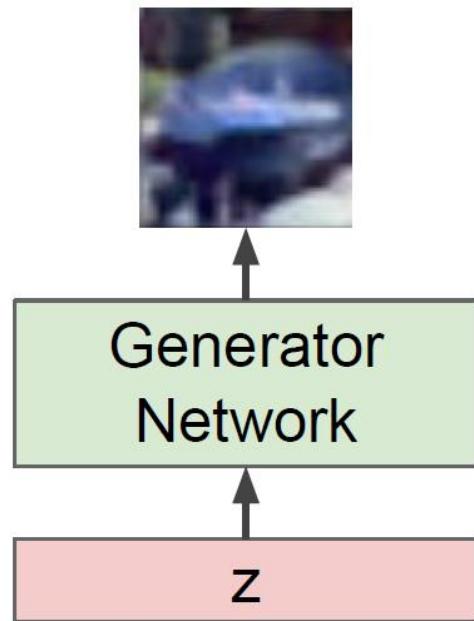
Designing a network for generative tasks

1. We need an architecture that can generate an image
 - Recall upsampling architectures for dense prediction
 - Sample from a simple distribution, e.g. random noise.
 - Learn transformation to training distribution.



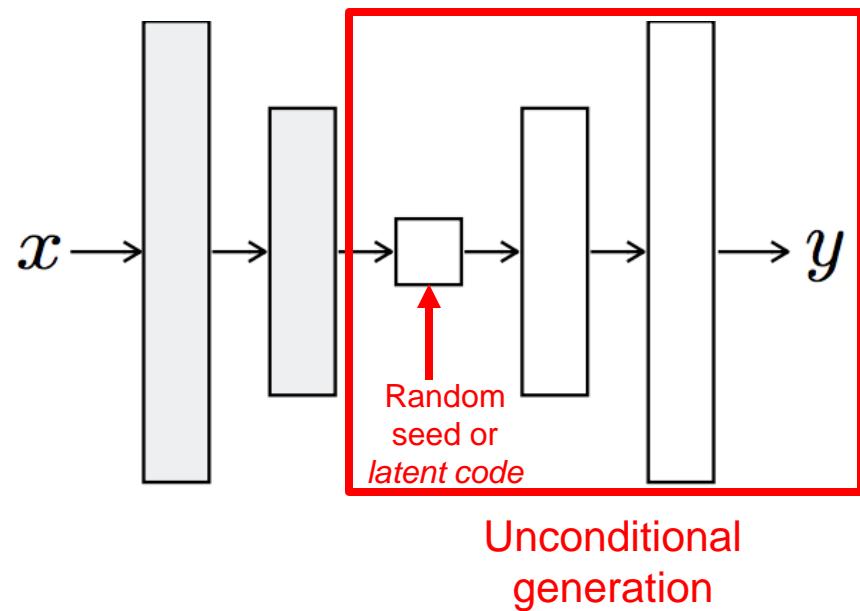
Output: Sample from
training distribution

Input: Random noise



Designing a network for generative tasks

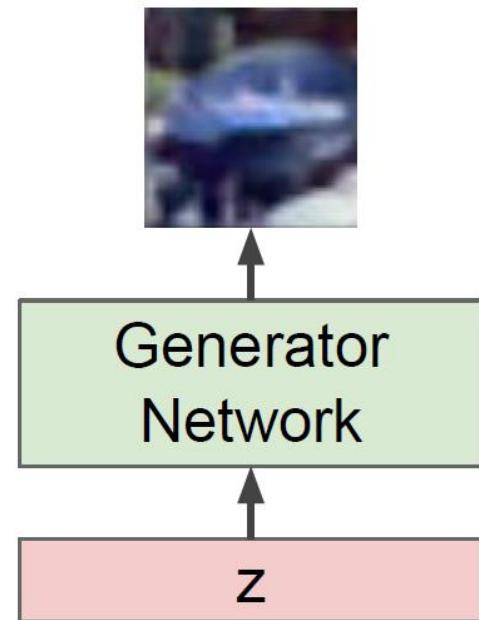
1. We need an architecture that can generate an image
 - Recall upsampling architectures for dense prediction
 - Sample from a simple distribution, e.g. random noise.
 - Learn transformation to training distribution.



Output: Sample from
training distribution

**A neural network can be
used to represent
this complex
transformation?**

Input: Random noise



Designing a network for generative tasks

1. We need an architecture that can generate an image
 - Recall upsampling architectures for dense prediction
 - Sample from a simple distribution, e.g. random noise.
 - Learn transformation to training distribution.

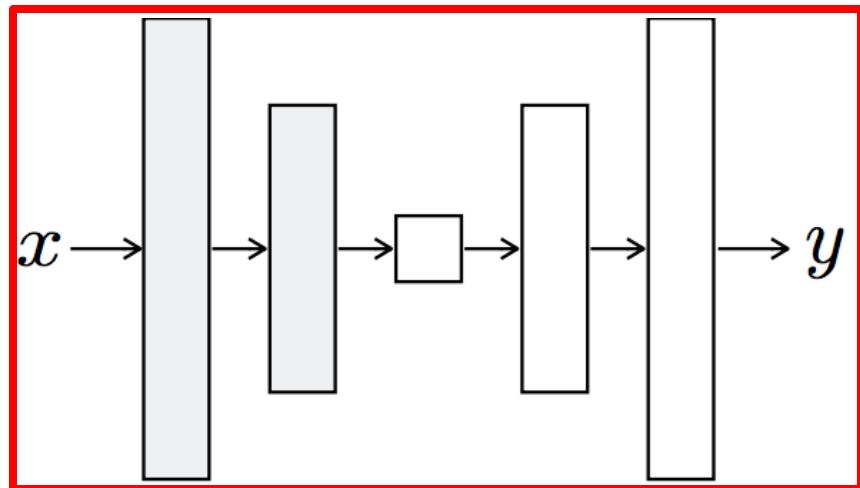
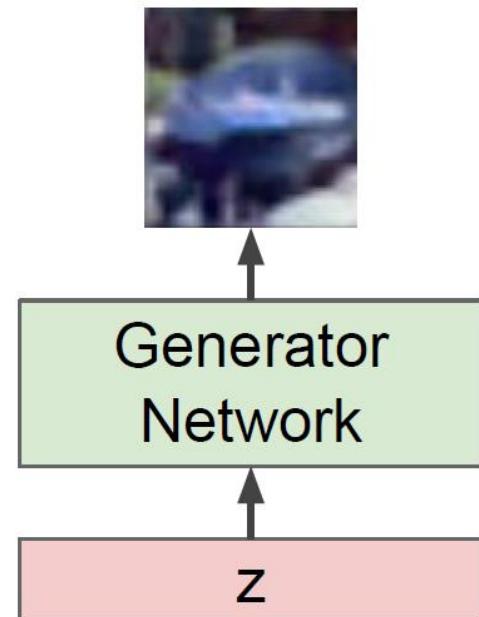


Image-to-image translation

Output: Sample from
training distribution

**A neural network can be
used to represent
this complex
transformation?**

Input: Random noise



Designing a network for generative tasks

1. We need an architecture that can generate an image
 - Recall upsampling architectures for dense prediction
2. We need to design the right loss function

Learning to sample



Training data $x \sim p_{\text{data}}$



Generated samples $x \sim p_{\text{model}}$

We want to learn p_{model} that matches p_{data}

Generative adversarial networks

- Train two networks with opposing objectives:
 - **Generator:** learns to generate samples
 - **Discriminator:** learns to distinguish between generated and real samples

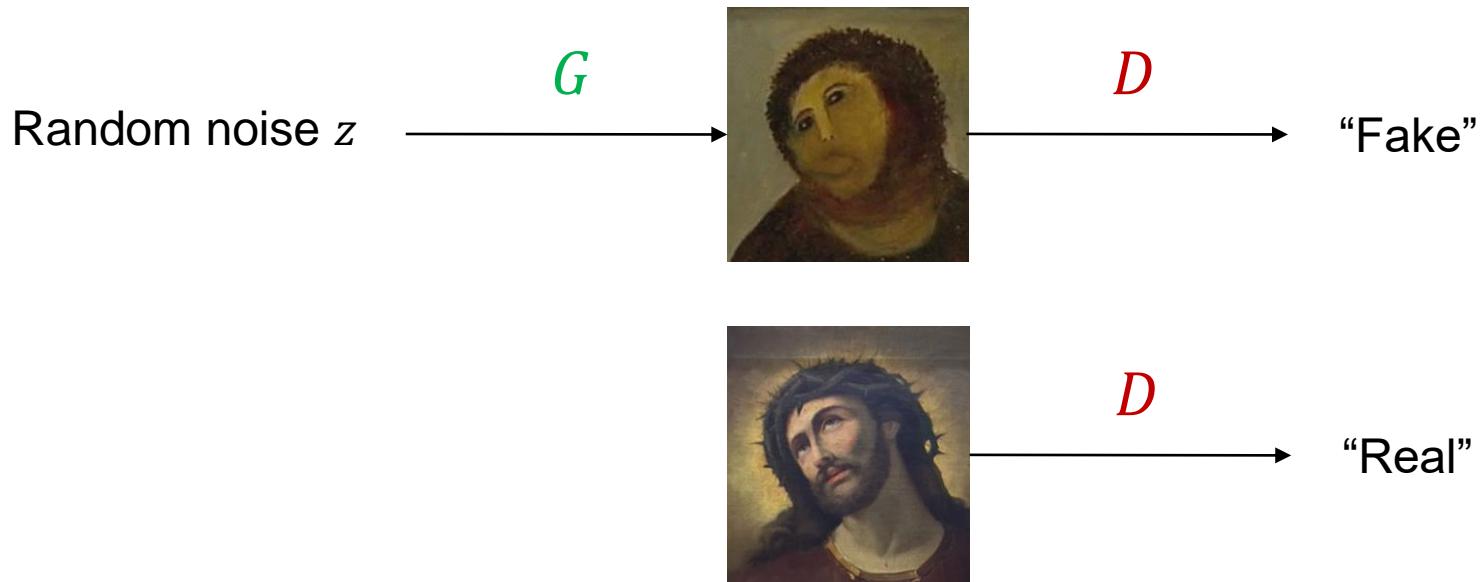
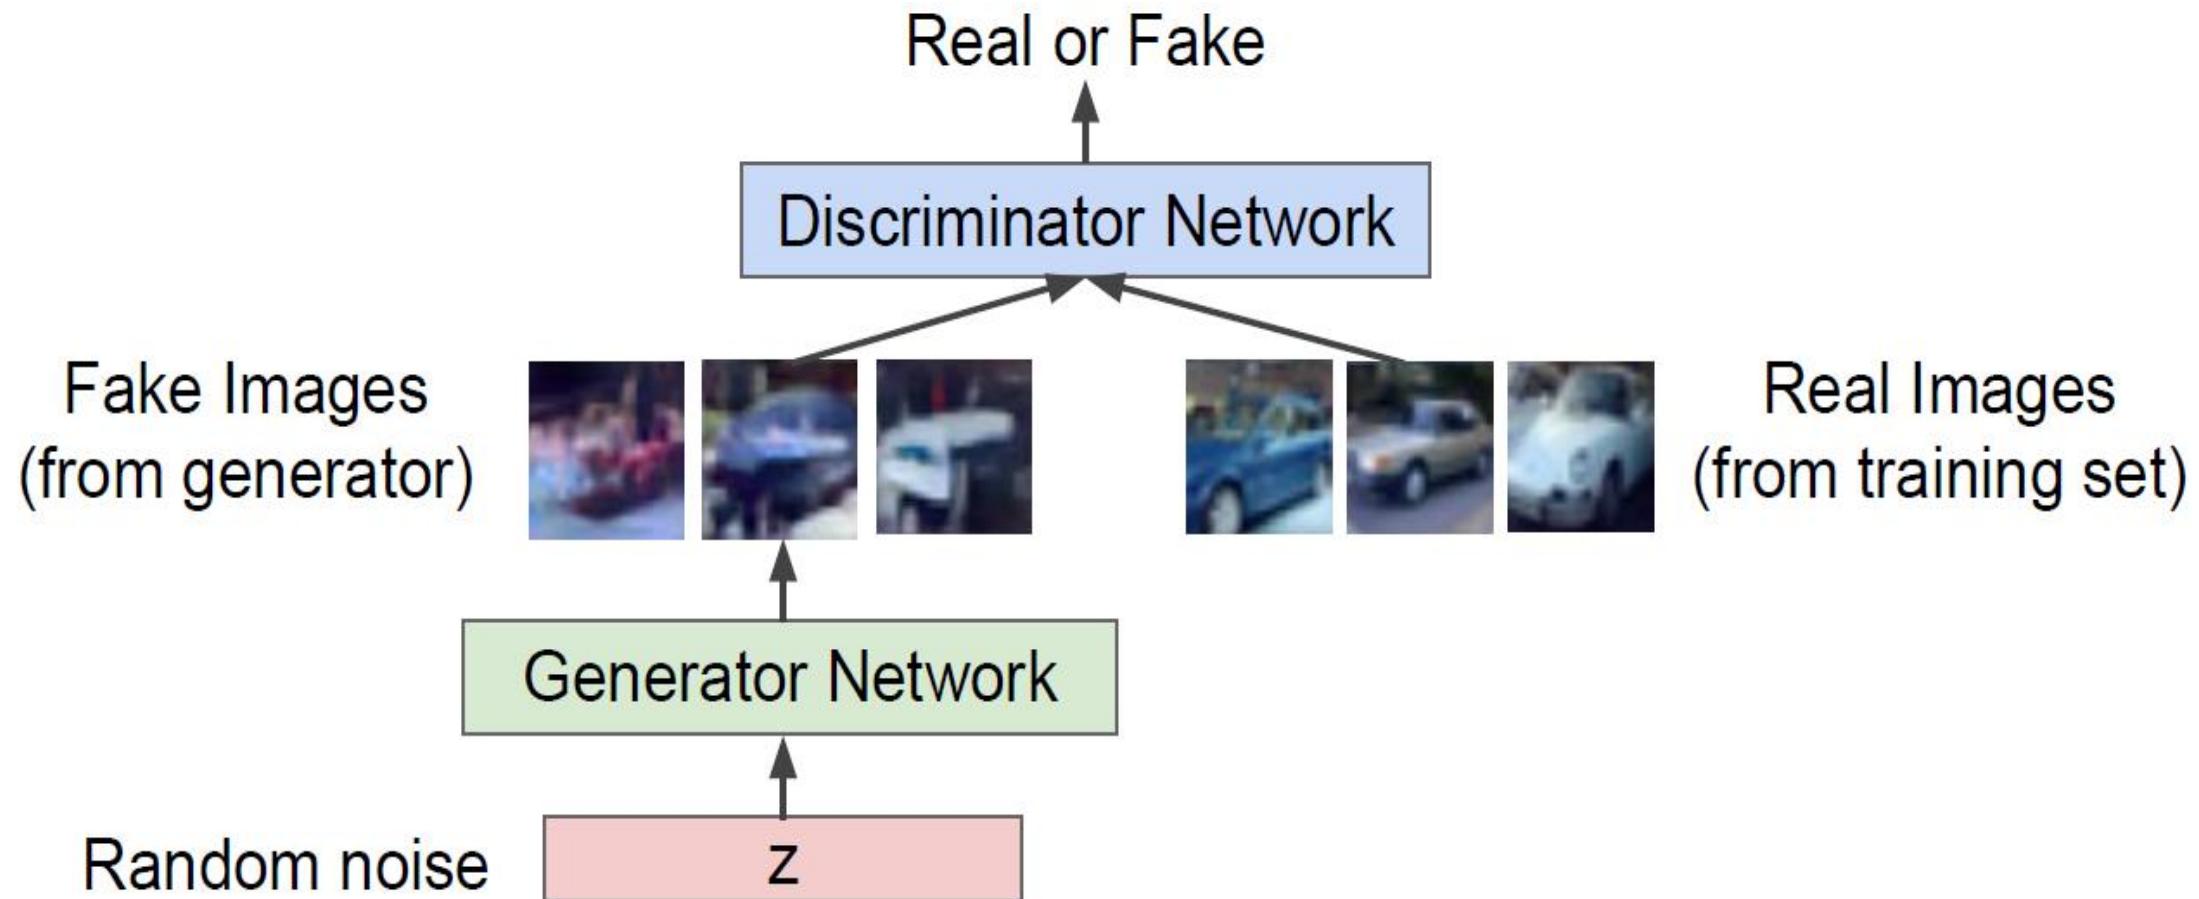


Figure adapted
from [F. Fleuret](#)

Generative adversarial networks

Generator network: try to fool the discriminator by generating real-looking images

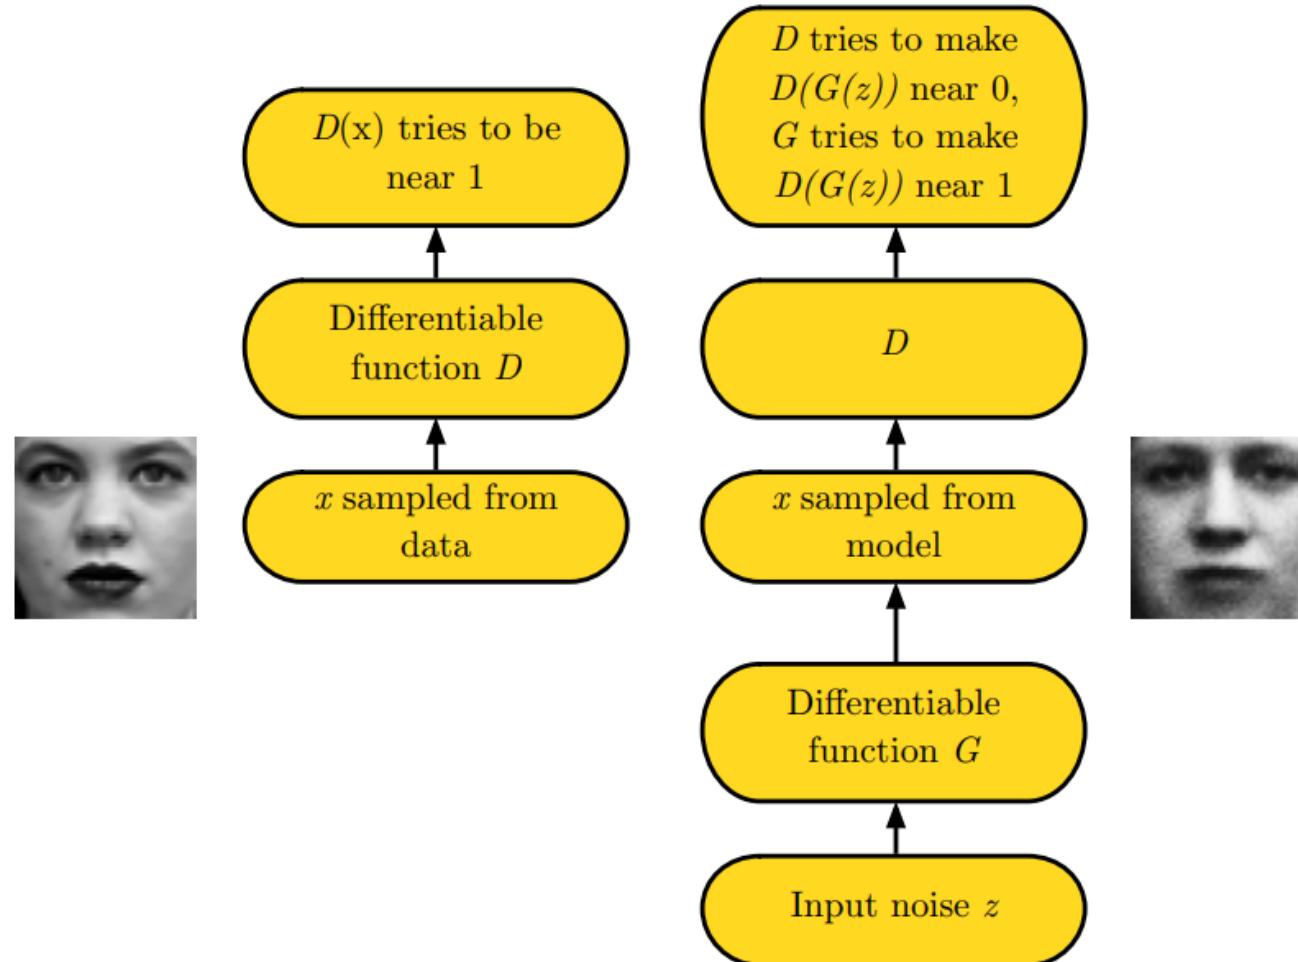
Discriminator network: try to distinguish between real and fake images



Generative adversarial networks

Generator network: try to fool the discriminator by generating real-looking images

Discriminator network: try to distinguish between real and fake images



GAN objective

- The discriminator $D(x)$ should output the probability that the sample x is real
 - That is, we want $D(x)$ to be close to 1 for real data and close to 0 for fake
- Expected conditional log likelihood for real and generated data:

$$\begin{aligned} & \mathbb{E}_{x \sim p_{\text{data}}} \log D(x) + \mathbb{E}_{x \sim p_{\text{gen}}} \log(1 - D(x)) \\ &= \mathbb{E}_{x \sim p_{\text{data}}} \log D(x) + \mathbb{E}_{z \sim p} \log(1 - D(G(z))) \end{aligned}$$

We seed the generator with noise z
drawn from a simple distribution p
(Gaussian or uniform)

GAN objective

$$V(G, D) = \mathbb{E}_{x \sim p_{\text{data}}} \log D(x) + \mathbb{E}_{z \sim p} \log(1 - D(G(z)))$$

- The discriminator wants to correctly distinguish real and fake samples:

$$D^* = \arg \max_D V(G, D)$$

- The generator wants to fool the discriminator:

$$G^* = \arg \min_G V(G, D)$$

- Train the generator and discriminator jointly in a *minimax game*

GAN objective: Theoretical properties

$$V(G, D) = \mathbb{E}_{x \sim p_{\text{data}}} \log D(x) + \mathbb{E}_{z \sim p} \log(1 - D(G(z)))$$

- Assuming unlimited capacity for generator and discriminator and unlimited training data:
 - The objective $\min_G \max_D V(G, D)$ is equivalent to *Jensen-Shannon divergence* between p_{data} and p_{gen} and global optimum (*Nash equilibrium*) is given by $p_{\text{data}} = p_{\text{gen}}$

GAN objective: Theoretical properties

$$V(G, D) = \mathbb{E}_{x \sim p_{\text{data}}} \log D(x) + \mathbb{E}_{z \sim p} \log(1 - D(G(z)))$$

- Assuming unlimited capacity for generator and discriminator and unlimited training data:
 - The objective $\min_G \max_D V(G, D)$ is equivalent to *Jensen-Shannon divergence* between p_{data} and p_{gen} and global optimum (*Nash equilibrium*) is given by $p_{\text{data}} = p_{\text{gen}}$
 - If at each step, D is allowed to reach its optimum given G , and G is updated to decrease $V(G, D)$, then p_{gen} will eventually converge to p_{data}

GAN training

$$V(G, D) = \mathbb{E}_{x \sim p_{\text{data}}} \log D(x) + \mathbb{E}_{z \sim p} \log(1 - D(G(z)))$$

- Alternate between

- *Gradient ascent* on discriminator:

$$D^* = \arg \max_D V(G, D)$$

- *Gradient descent* on generator (minimize log-probability of discriminator being right):

$$\begin{aligned} G^* &= \arg \min_G V(G, D) \\ &= \arg \min_G \mathbb{E}_{z \sim p} \log(1 - D(G(z))) \end{aligned}$$

- In practice, do *gradient ascent* on generator (maximize log-probability of discriminator being wrong):

$$G^* = \arg \max_G \mathbb{E}_{z \sim p} \log(D(G(z)))$$

Non-saturating GAN loss (NSGAN)

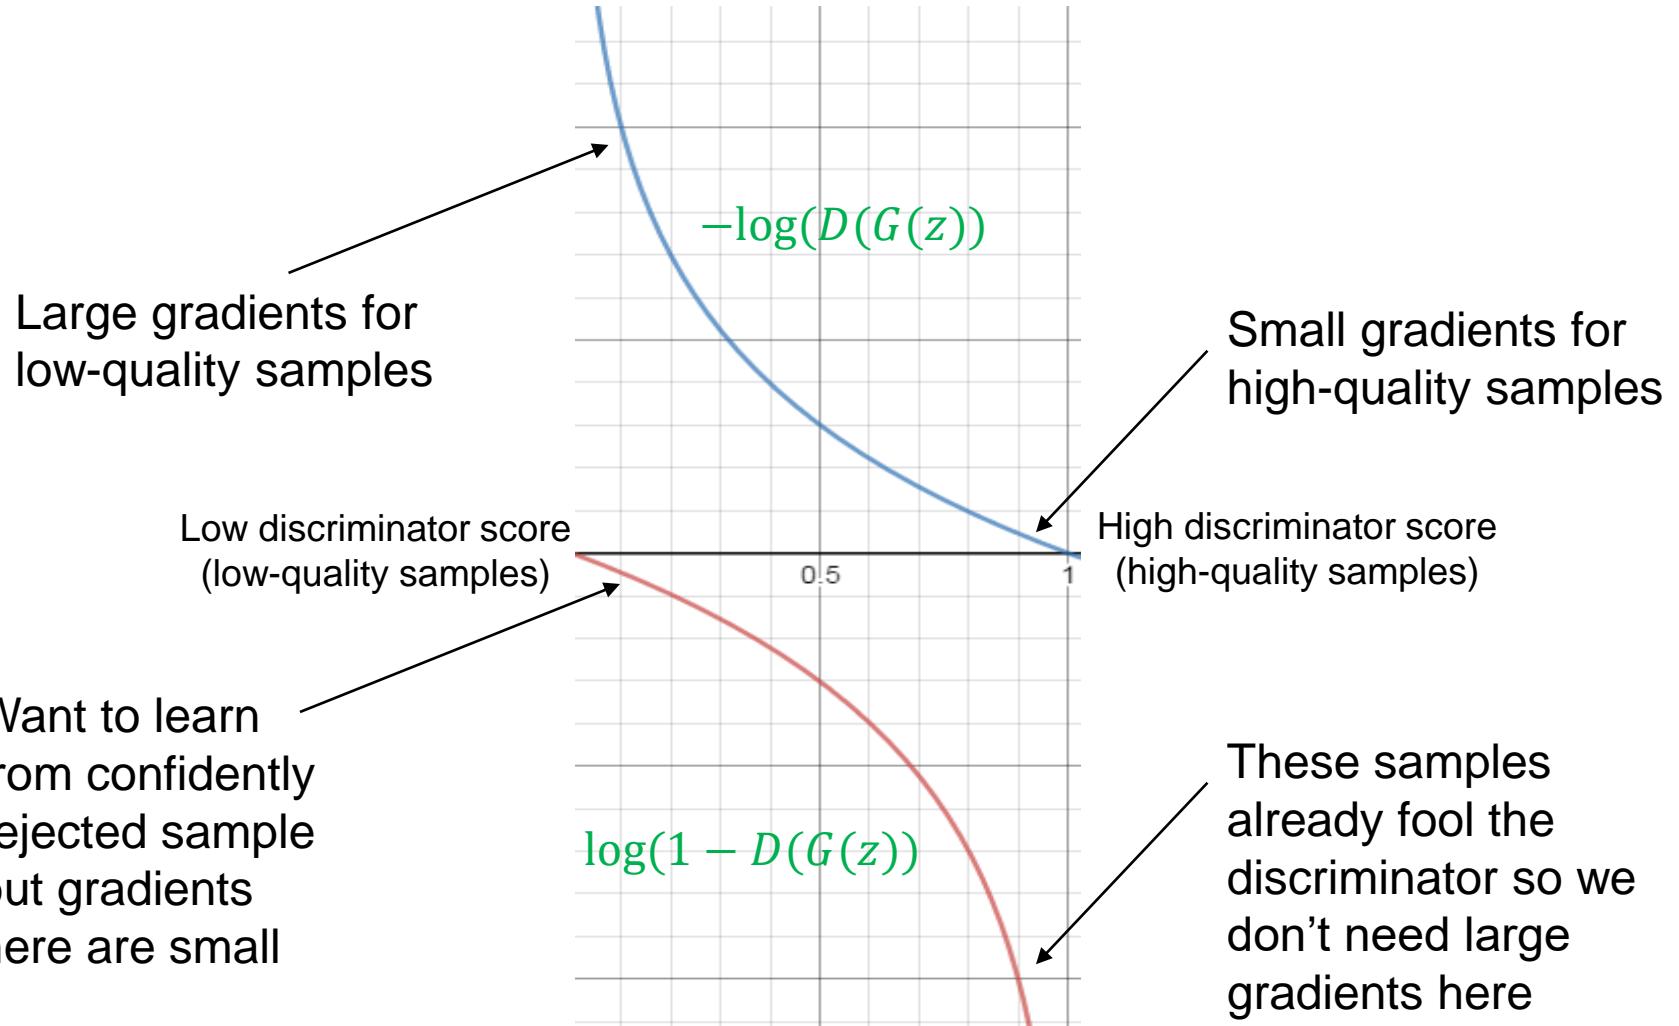
$$\min_{w_G} \mathbb{E}_{z \sim p} \log(1 - D(G(z))) \quad \text{vs.} \quad \max_{w_G} \mathbb{E}_{z \sim p} \log(D(G(z)))$$

Minimize log-probability of
discriminator being right

Maximize log-probability of
discriminator being wrong

Non-saturating GAN loss (NSGAN)

$$\min_{w_G} \mathbb{E}_{z \sim p} \log(1 - D(G(z))) \quad \text{vs.} \quad \max_{w_G} \mathbb{E}_{z \sim p} \log(D(G(z)))$$



NSGAN training algorithm

- Update discriminator:
 - Repeat for k steps:
 - Sample mini-batch of noise samples z_1, \dots, z_m and mini-batch of real samples x_1, \dots, x_m
 - Update parameters of D by stochastic gradient ascent on
$$\frac{1}{m} \sum_m [\log D(x_m) + \log(1 - D(G(z_m)))]$$
 - Update generator:
 - Sample mini-batch of noise samples z_1, \dots, z_m
 - Update parameters of G by stochastic gradient ascent on
$$\frac{1}{m} \sum_m \log D(G(z_m))$$
- Repeat until happy with results

NSGAN training algorithm

- Update discriminator:
 - Repeat for k steps:
 - Sample mini-batch of noise samples z_1, \dots, z_m and mini-batch of real samples x_1, \dots, x_m
 - Update parameters of D by stochastic gradient ascent on
$$\frac{1}{m} \sum_m [\log D(x_m) + \log(1 - D(G(z_m)))]$$

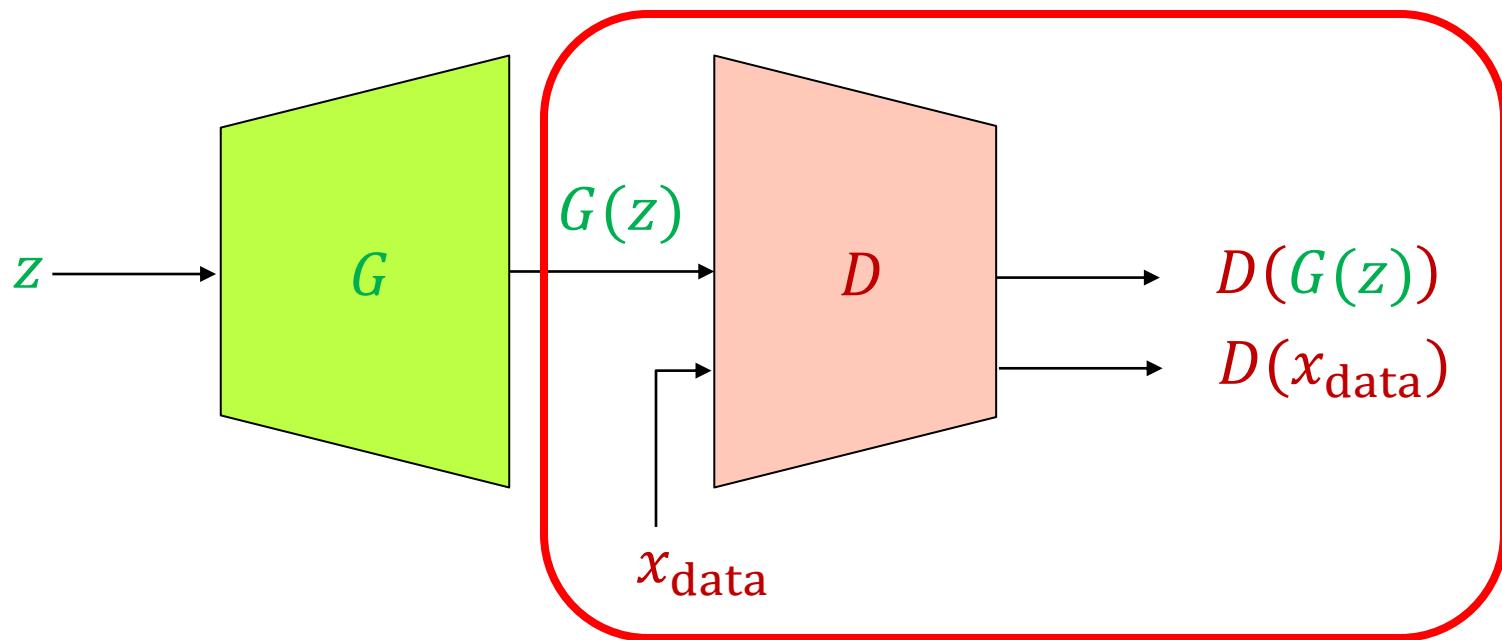
Some find $k=1$ more stable, others use $k > 1$, no best rule.

Recent work (e.g. Wasserstein GAN) alleviates this problem, better stability!

- Update generator:
 - Sample mini-batch of noise samples z_1, \dots, z_m
 - Update parameters of G by stochastic gradient ascent on
$$\frac{1}{m} \sum_m \log D(G(z_m))$$
- Repeat until happy with results

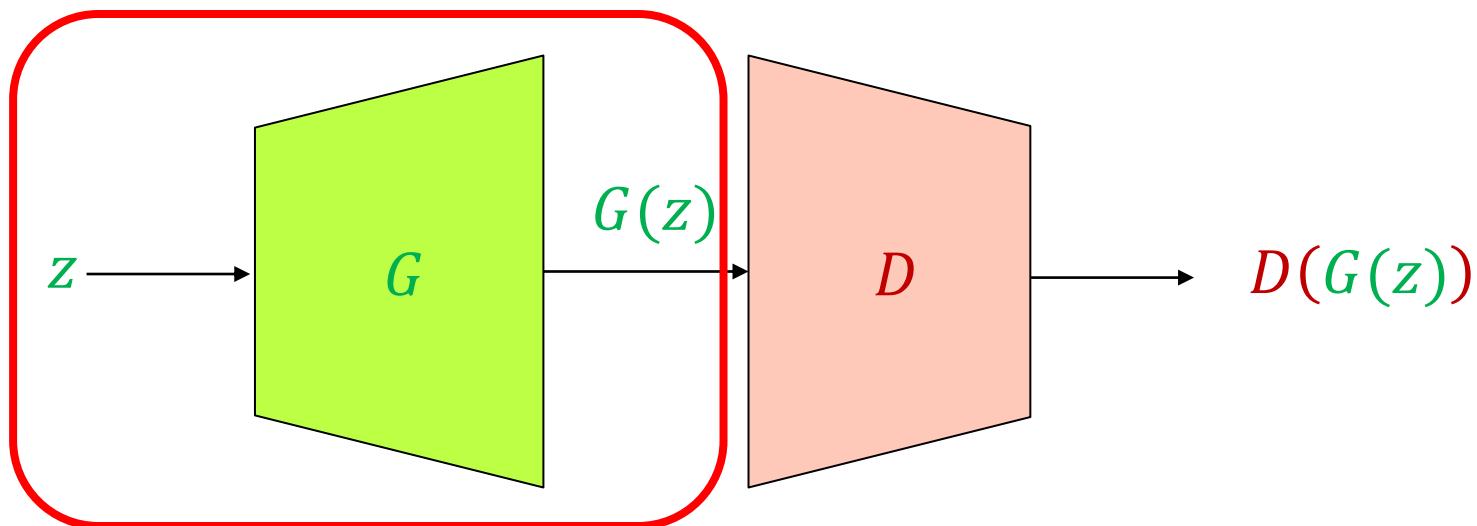
GAN: Conceptual picture

- Update discriminator: push $D(x_{\text{data}})$ close to 1 and $D(G(z))$ close to 0
 - The generator is a “black box” to the discriminator



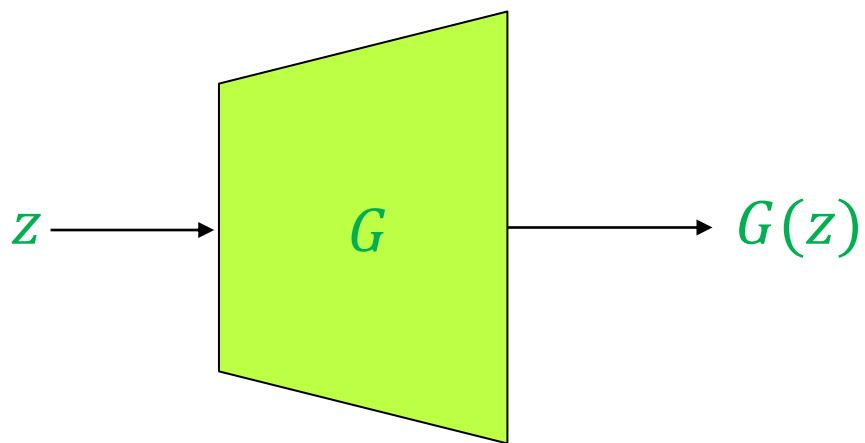
GAN: Conceptual picture

- Update generator: increase $D(G(z))$
 - Requires back-propagating through the composed generator-discriminator network (i.e., the discriminator cannot be a black box)
 - The generator is exposed to real data only via the output of the discriminator (and its gradients)



GAN: Conceptual picture

- Test time – the discriminator is discarded



Original GAN results

MNIST digits



Toronto Face Dataset



Nearest real image for
sample to the left

I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair,
A. Courville, Y. Bengio, [Generative adversarial nets](#), NIPS 2014

Original GAN results

CIFAR-10 (FC networks)



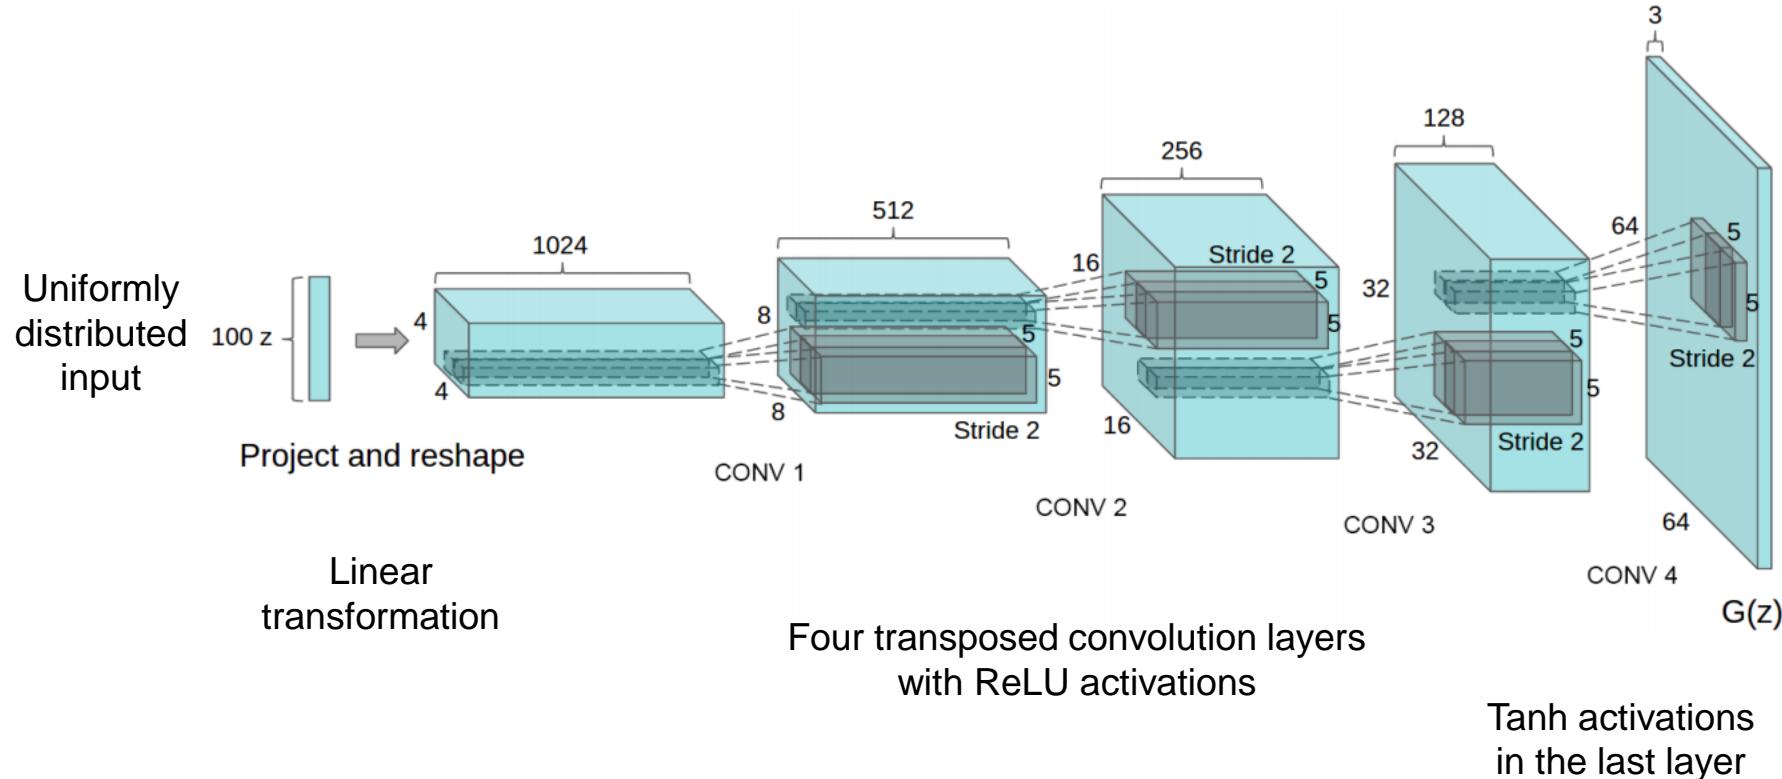
CIFAR-10 (conv networks)



I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair,
A. Courville, Y. Bengio, [Generative adversarial nets](#), NIPS 2014

DCGAN

- Early, influential convolutional architecture for generator



DCGAN

- Early, influential convolutional architecture for generator
- Discriminator architecture:
 - Don't use pooling, only strided convolutions
 - Use Leaky ReLU activations (sparse gradients cause problems for training)
 - Use only one FC layer before the softmax output
 - Use batch normalization after most layers (in the generator also)

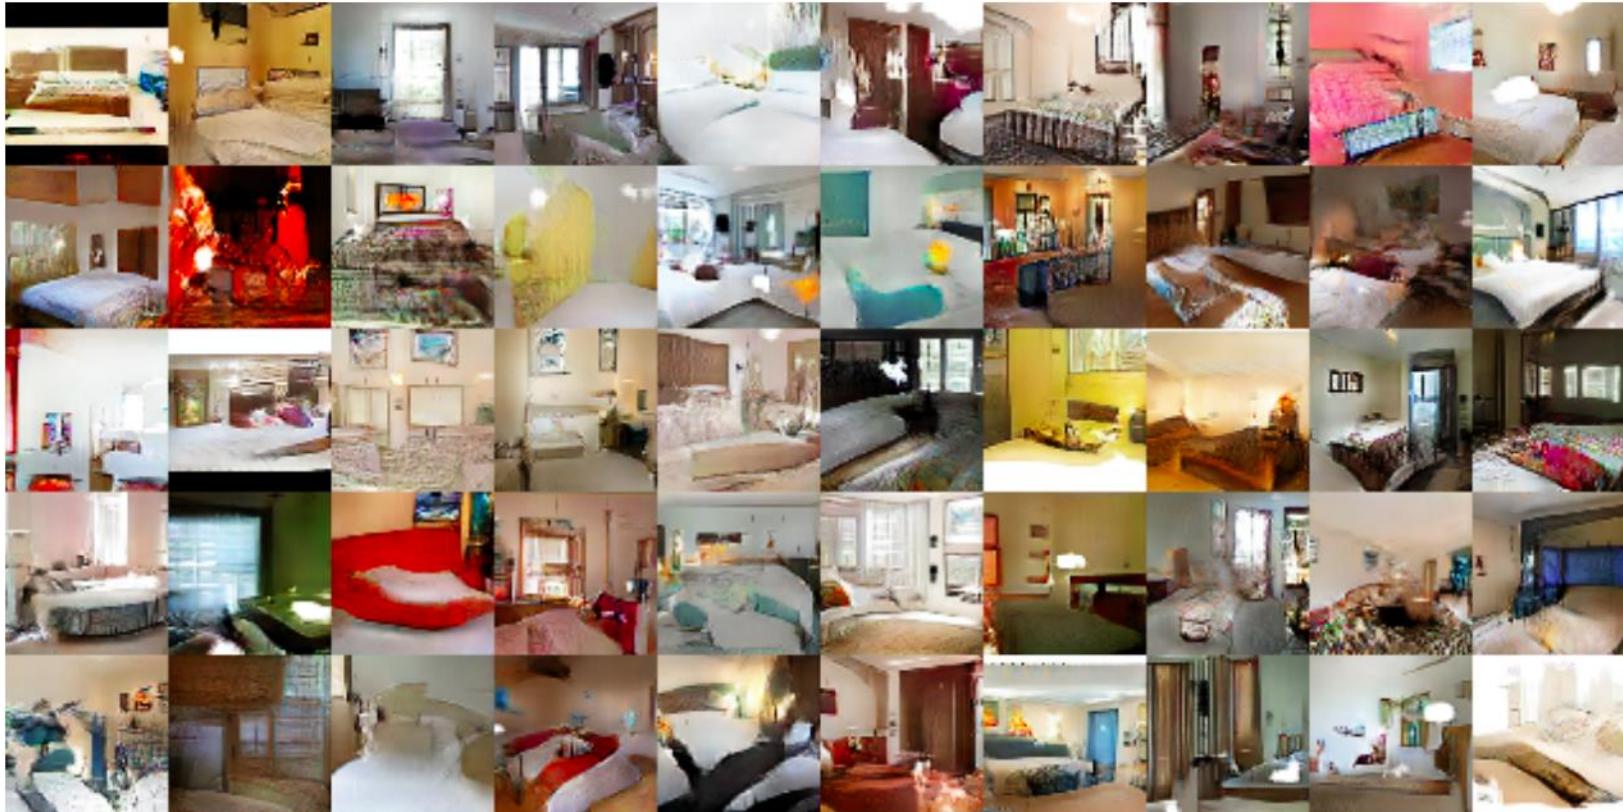
DCGAN results

Generated bedrooms after one epoch



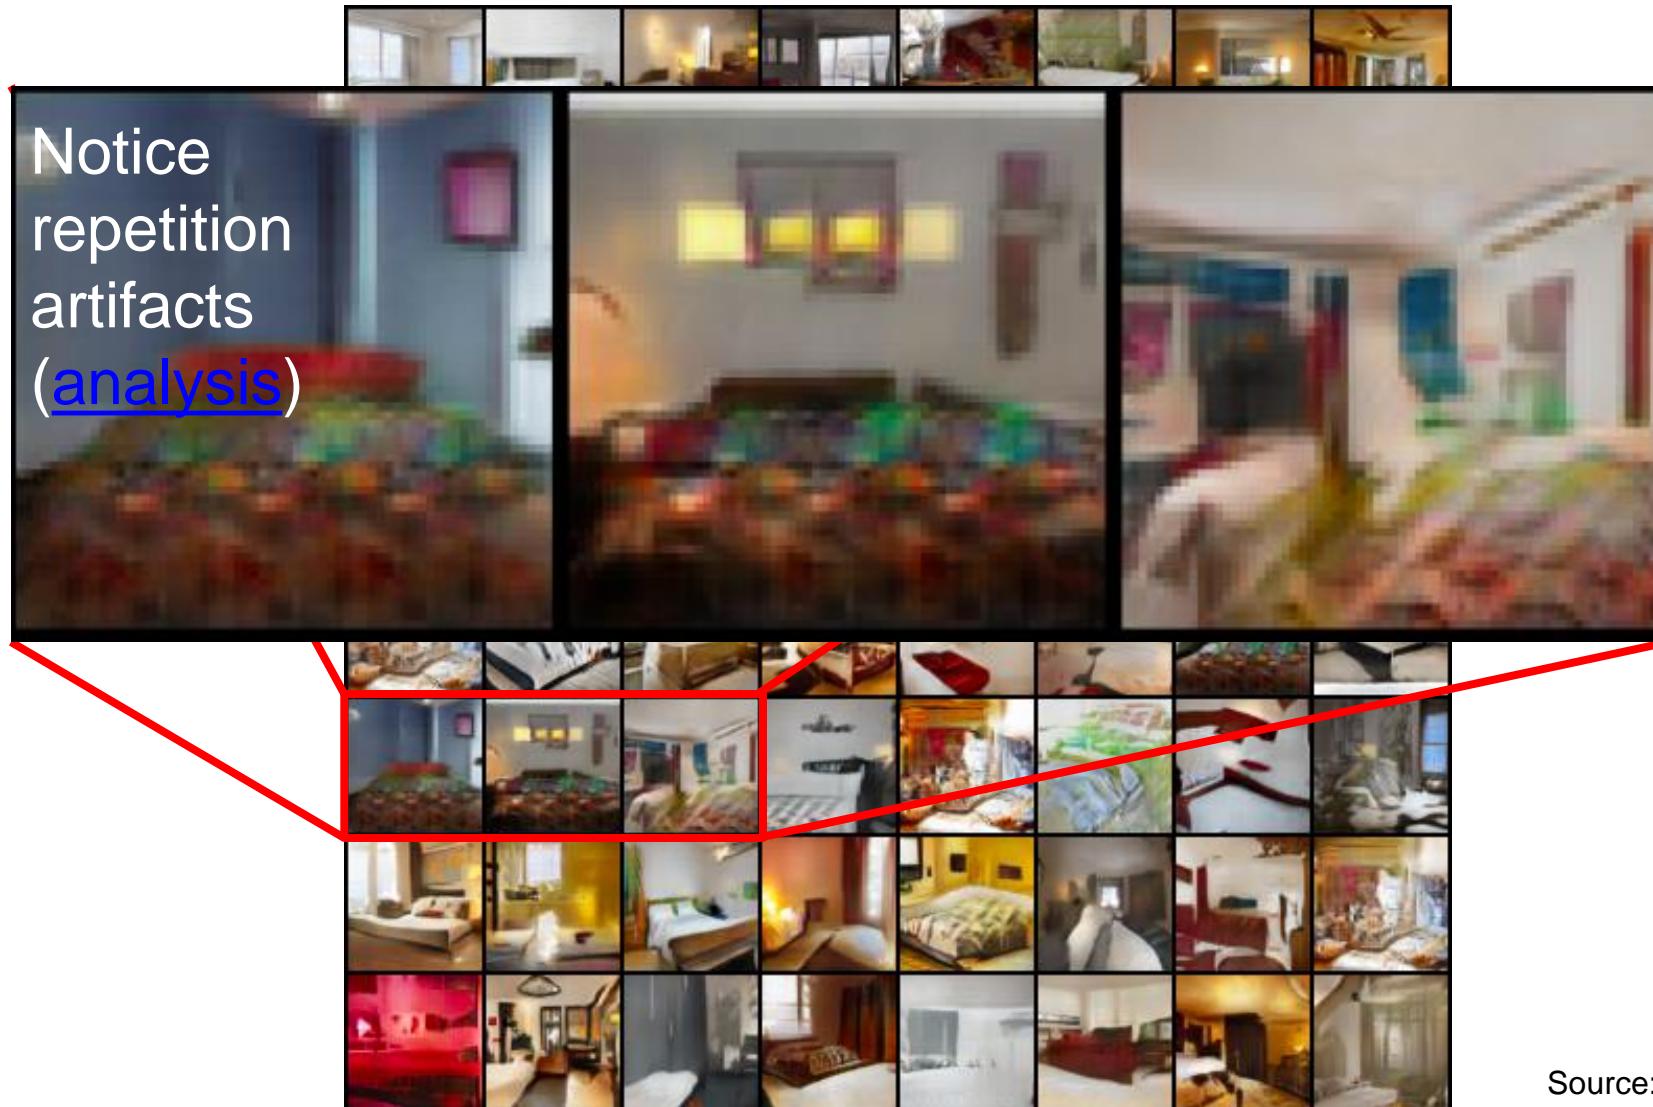
DCGAN results

Generated bedrooms after five epochs



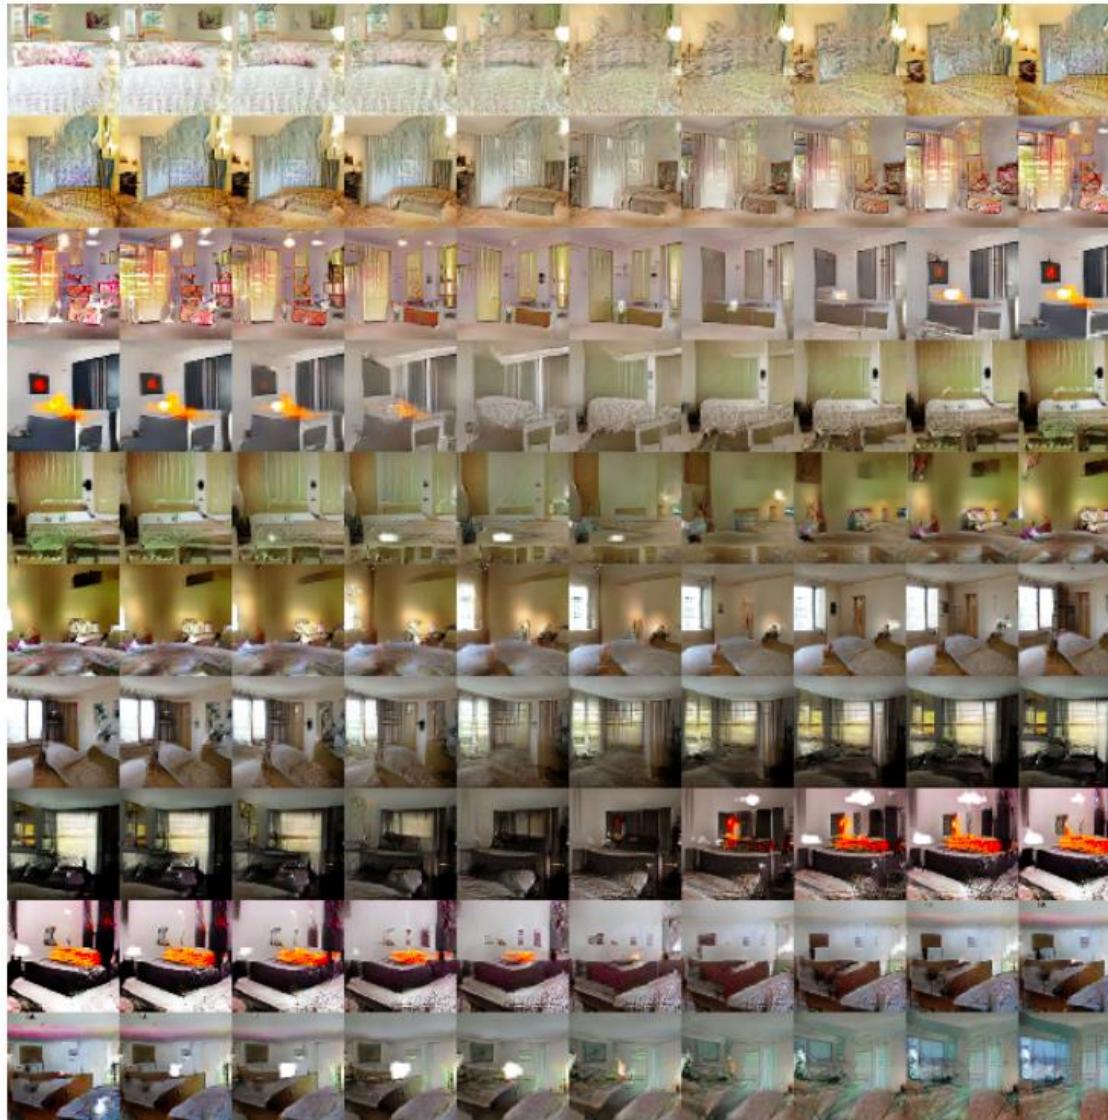
DCGAN results

More bedrooms



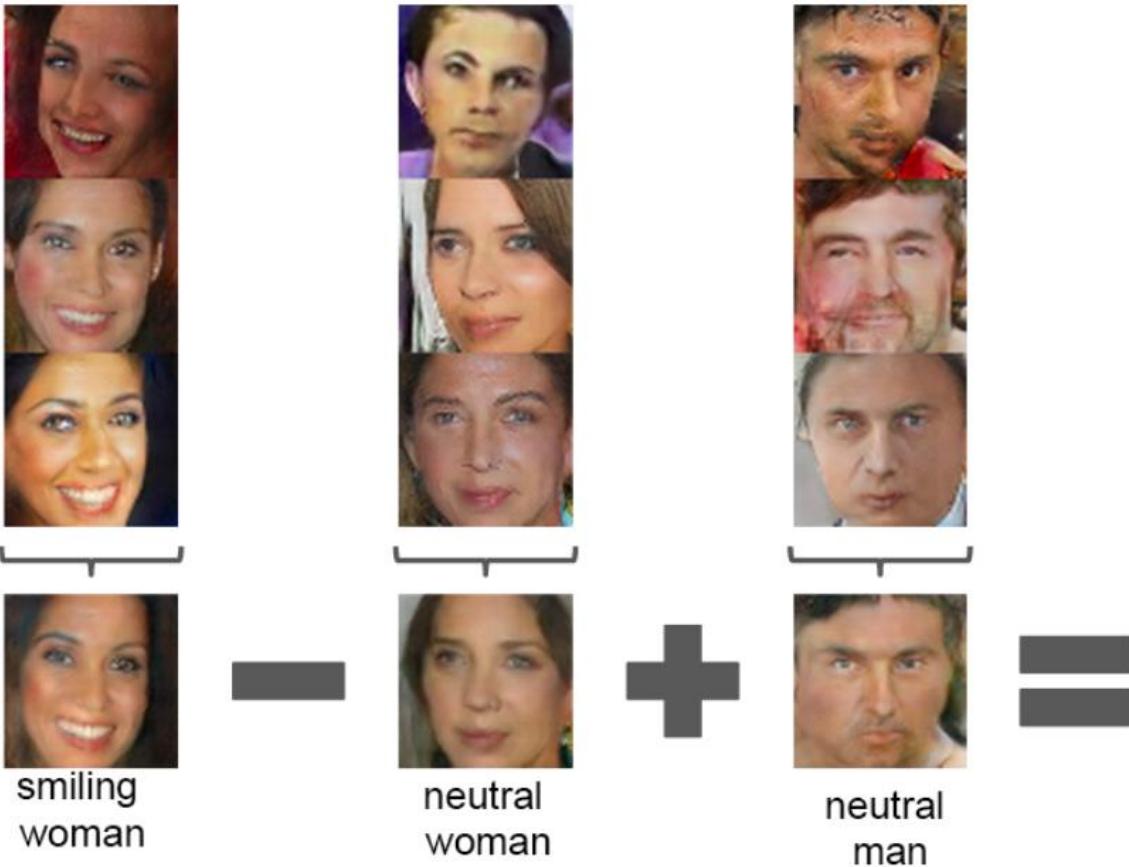
DCGAN results

Interpolation between different points in the z space



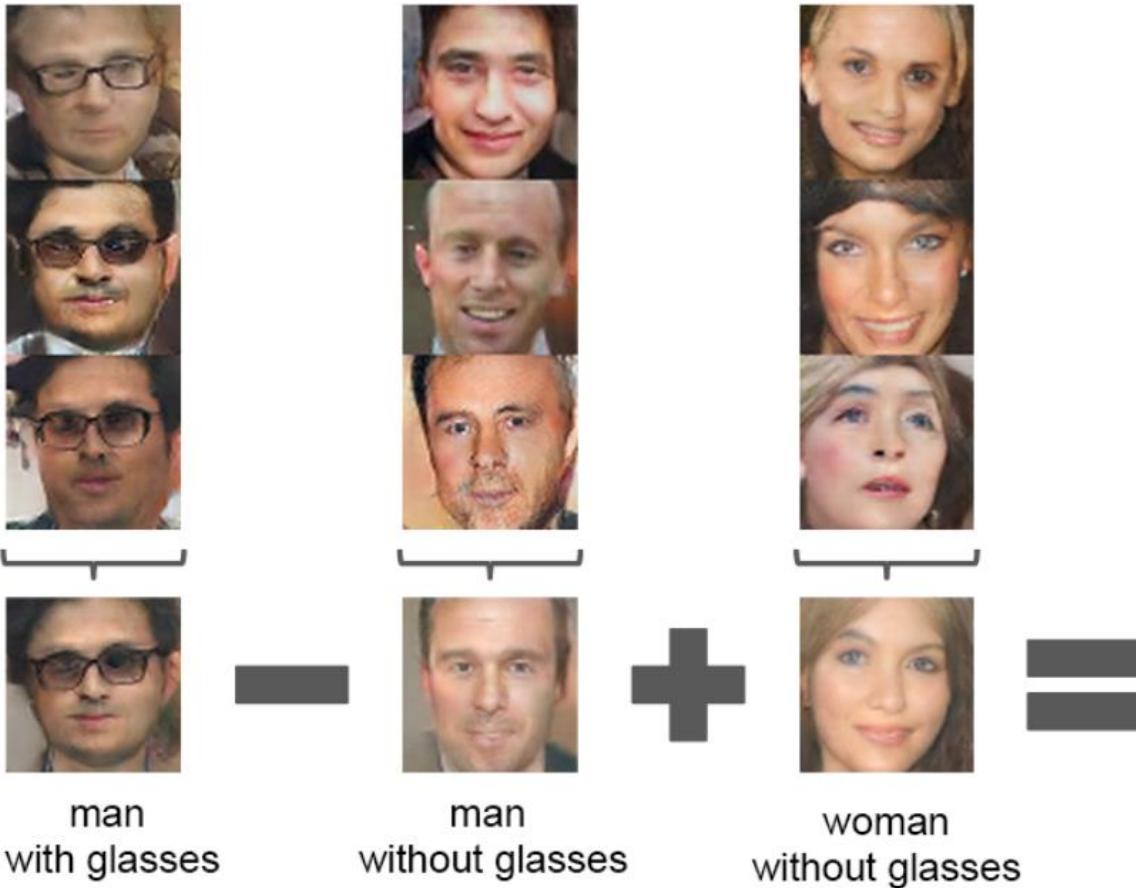
DCGAN results

- Vector arithmetic in the z space



DCGAN results

- Vector arithmetic in the z space



DCGAN results

- Pose transformation by adding a “turn” vector



Outline

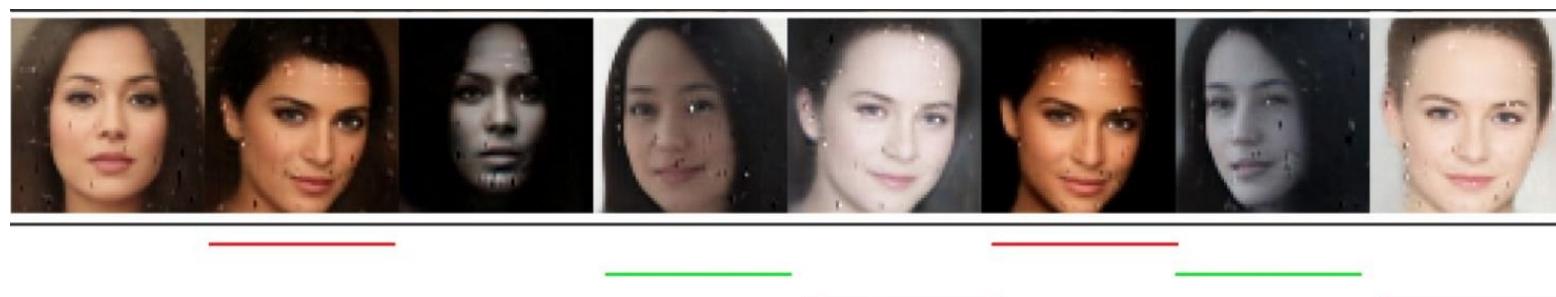
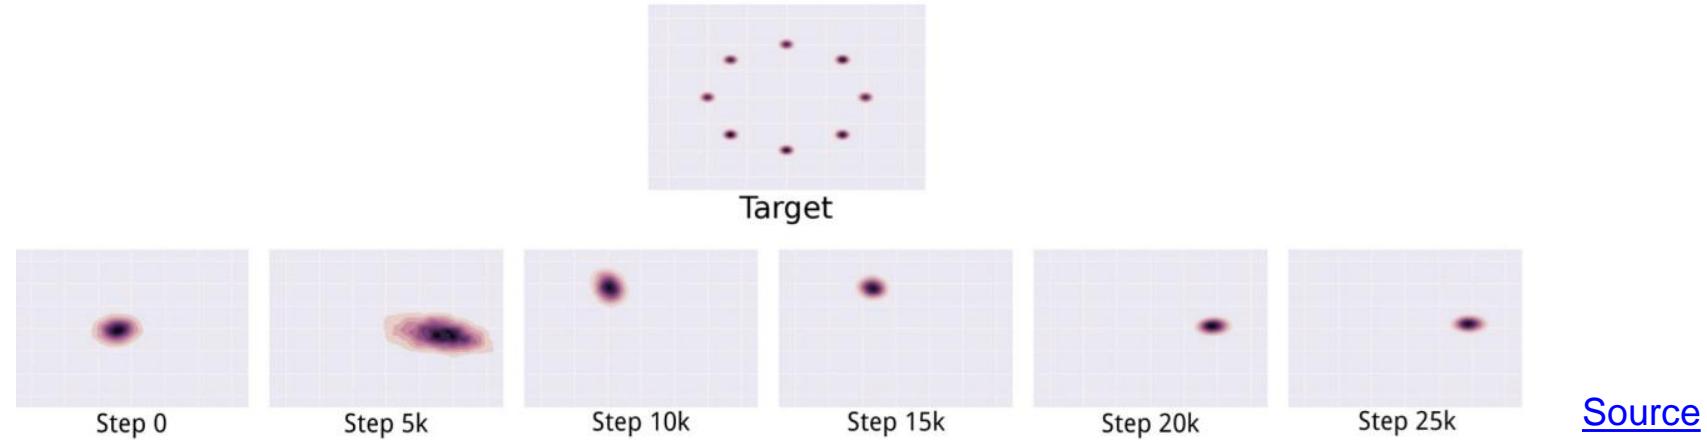
- Generative tasks
- Original GAN formulations
 - NSGAN
 - DCGAN
- Other popular formulations
 - WGAN, WGAN-GP
 - LSGAN

Problems with GAN training

- Stability
 - Parameters can oscillate or diverge, generator loss does not correlate with sample quality
 - Behavior very sensitive to hyperparameter selection

Problems with GAN training

- Mode collapse
 - Generator ends up modeling only a small subset of the training data



Some popular GAN flavors

- WGAN and improved WGAN (WGAN-GP)
- LSGAN

Wasserstein GAN (WGAN)

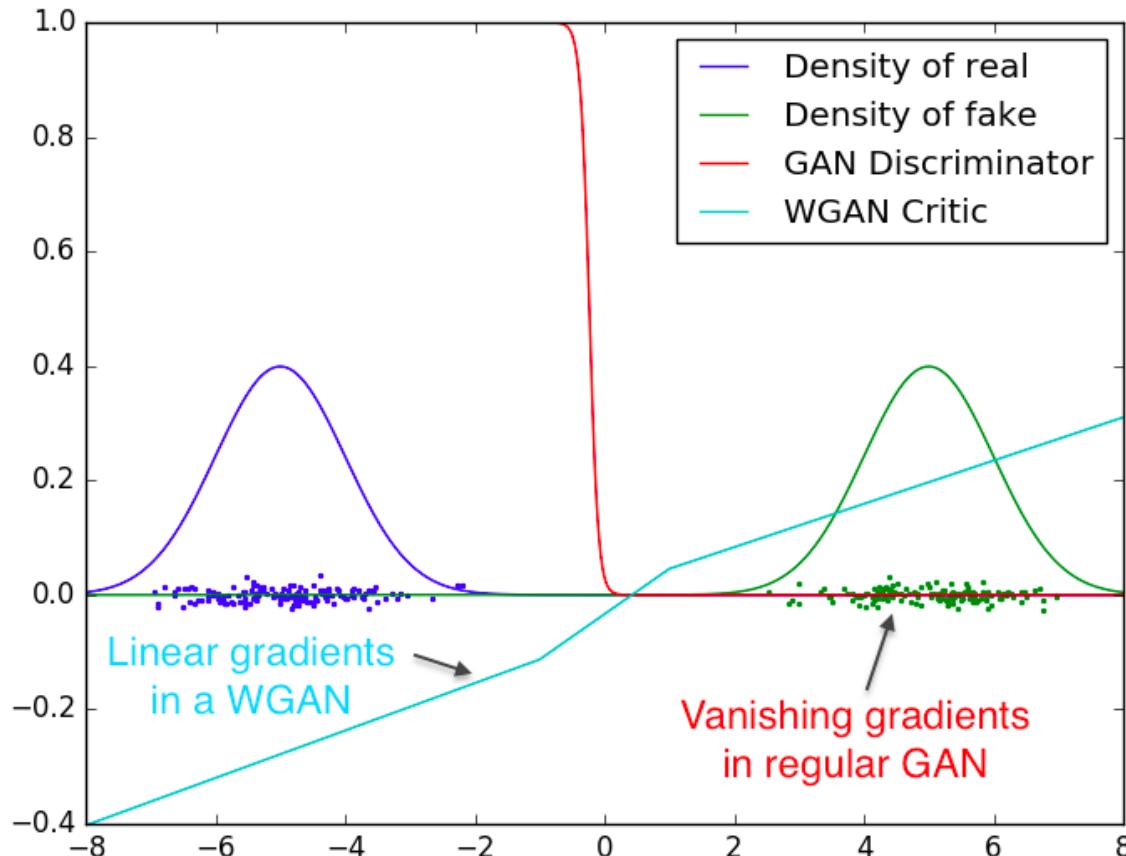
- Motivated by *Wasserstein or Earth mover's distance*, which is an alternative to JS divergence for comparing distributions
 - In practice, use linear activation instead of sigmoid in the discriminator and drop the logs from the objective:

$$\min_G \max_D [\mathbb{E}_{x \sim p_{\text{data}}} D(x) - \mathbb{E}_{z \sim p} D(G(z))]$$

- Due to theoretical considerations, important to ensure smoothness of discriminator
- This paper's suggested method is clipping weights to fixed range $[-c, c]$

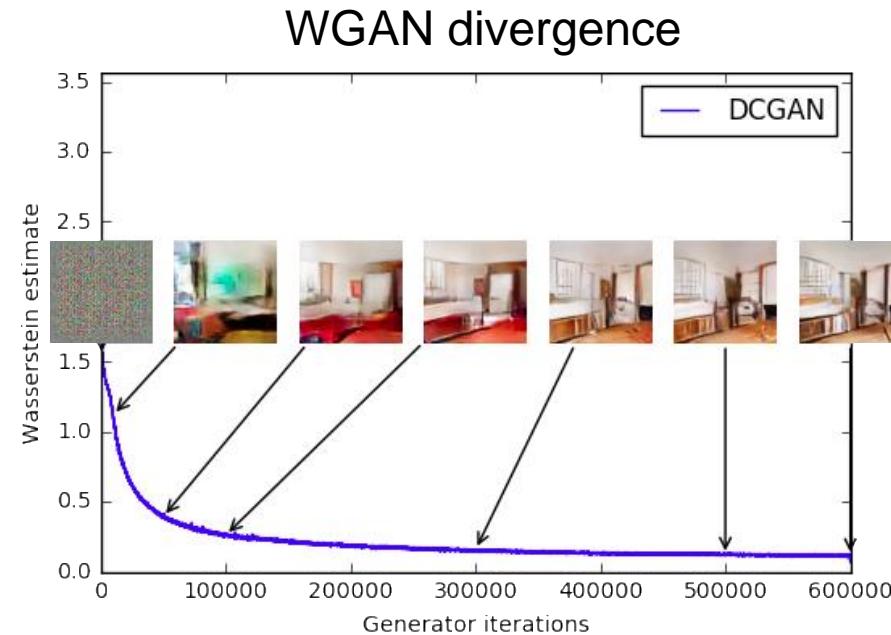
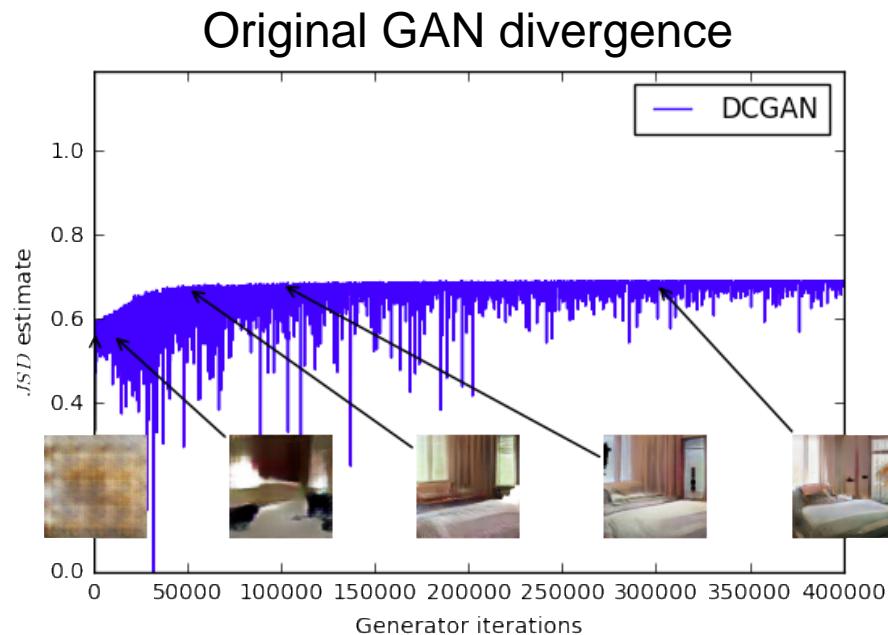
Wasserstein GAN (WGAN)

- Benefits (claimed)
 - Better gradients, more stable training



Wasserstein GAN (WGAN)

- Benefits (claimed)
 - Better gradients, more stable training
 - Objective function value is more meaningfully related to quality of generator output



Improved Wasserstein GAN (WGAN-GP)

- Weight clipping leads to problems with discriminator training
- Improved Wasserstein discriminator loss:

$$\mathbb{E}_{\tilde{x} \sim p_{\text{gen}}} D(\tilde{x}) - \mathbb{E}_{x \sim p_{\text{real}}} D(x)$$

$$+ \lambda \mathbb{E}_{\hat{x} \sim p_{\hat{x}}} [(\|\nabla_{\hat{x}} D(\hat{x})\|_2 - 1)^2]$$

Unit norm gradient penalty on
points \hat{x} obtained by interpolating
real and generated samples

Improved Wasserstein GAN: Results

| DCGAN | LSGAN | WGAN (clipping) | WGAN-GP (ours) |
|---|-------|-----------------|----------------|
| Baseline (G : DCGAN, D : DCGAN) | | | |
| G : No BN and a constant number of filters, D : DCGAN | | | |
| G : 4-layer 512-dim ReLU MLP, D : DCGAN | | | |
| No normalization in either G or D | | | |
| Gated multiplicative nonlinearities everywhere in G and D | | | |
| tanh nonlinearities everywhere in G and D | | | |
| 101-layer ResNet G and D | | | |

Least Squares GAN (LSGAN)

- Use least squares cost for generator and discriminator
 - Equivalent to minimizing Pearson χ^2 divergence

$$D^* = \arg \min_D [\mathbb{E}_{x \sim p_{\text{data}}} (D(x) - 1)^2 + \mathbb{E}_{z \sim p} (D(G(z)))^2]$$

Push discrim.
response on real
data close to 1 Push response on
generated data close to 0

$$G^* = \arg \min_G \mathbb{E}_{z \sim p} (D(G(z)) - 1)^2$$

Push response on
generated data close to 1

Least Squares GAN (LSGAN)

- Benefits (claimed)
 - Higher-quality images



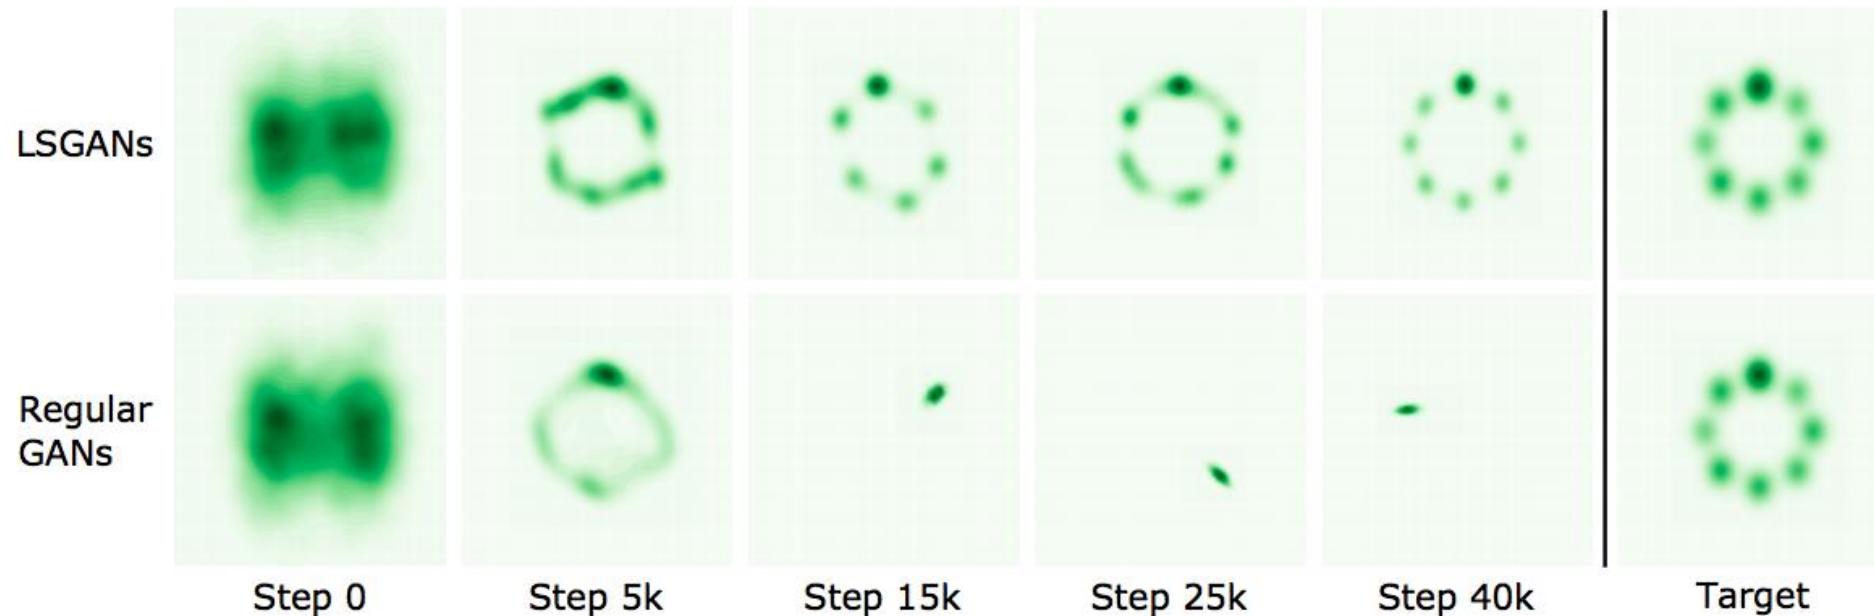
(a) Generated images (112×112) by LSGANs.



(b) Generated images (112×112) by DCGANs.

Least Squares GAN (LSGAN)

- Benefits (claimed)
 - Higher-quality images
 - More stable and resistant to mode collapse



Are GANs created equal?

- From the abstract:

“We find that most models can reach similar scores with enough hyperparameter optimization and random restarts. This suggests that improvements can arise from a higher computational budget and tuning more than fundamental algorithmic changes ... We did not find evidence that any of the tested algorithms consistently outperforms the non-saturating GAN introduced in Goodfellow et al. (2014)”

Outline

- Generative tasks
- Original GAN formulations
 - NSGAN
 - DCGAN
- Other popular formulations
 - WGAN, WGAN-GP
 - LSGAN
- State-of-the-art architectures
 - Progressive GAN, StyleGAN

Recent progress in GANs



Ian Goodfellow
@goodfellow_ian



4.5 years of GAN progress on face generation.

arxiv.org/abs/1406.2661 arxiv.org/abs/1511.06434

arxiv.org/abs/1606.07536 arxiv.org/abs/1710.10196

arxiv.org/abs/1812.04948



6:40 PM · Jan 14, 2019



Recent progress in GANs

EBGAN (2017)



ENERGY-
BASED
GAN

BigGAN (2018)



Progressive GANs

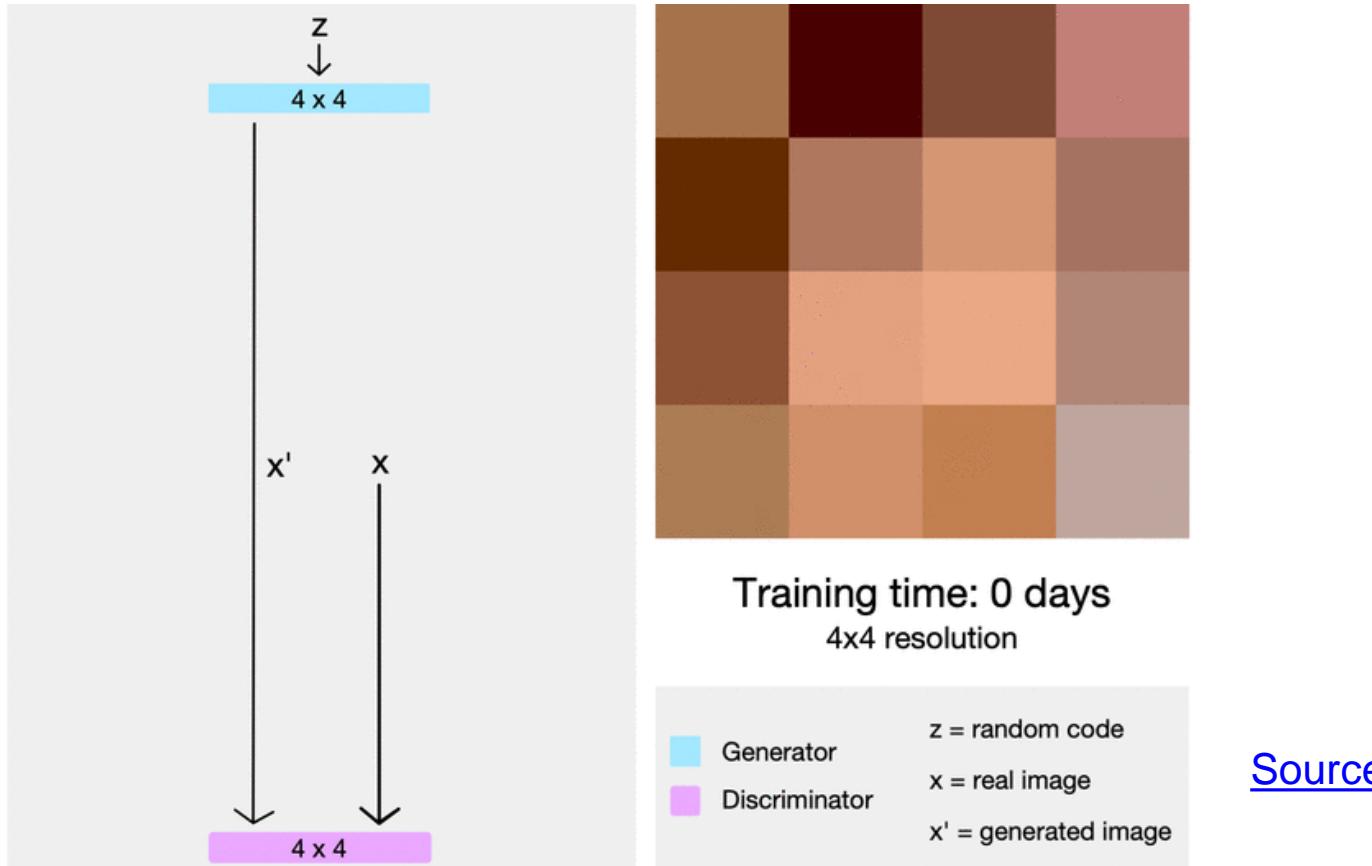
Realistic face images up to 1024 x 1024 resolution



T. Karras, T. Aila, S. Laine, J. Lehtinen. [Progressive Growing of GANs for Improved Quality, Stability, and Variation](#). ICLR 2018

Progressive GANs

- Key idea: train lower-resolution models, gradually add layers corresponding to higher-resolution outputs

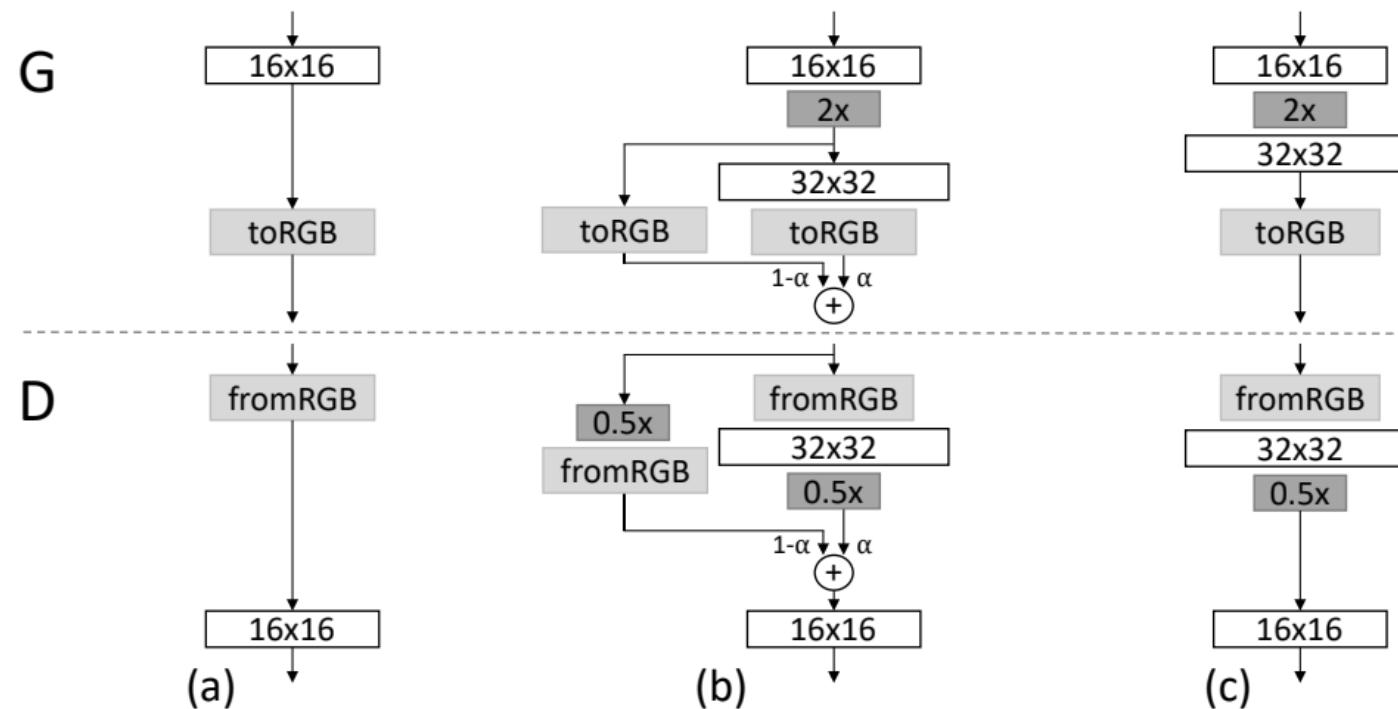


T. Karras, T. Aila, S. Laine, J. Lehtinen. [Progressive Growing of GANs for Improved Quality, Stability, and Variation](#). ICLR 2018

Progressive GANs

- Key idea: train lower-resolution models, gradually add layers corresponding to higher-resolution outputs

Transition from 16x16 to 32x32 images



Progressive GANs: Implementation details

- Loss: WGAN-GP loss (preferred) or LSGAN
- Architectures:
 - Nearest neighbor upsampling (2×2 replication) followed by regular convolutions instead of transposed conv layers
 - Average pooling instead of striding for downsampling in discriminator
 - Leaky ReLUs used in discriminator and generator
 - Per-pixel response normalization in generator: rescale feature vector in each pixel to unit length after each conv layer
- Use of minibatch standard deviation in discriminator (append to feature map)
- Exponential moving average of generator weights for display

Progressive GANs: Results

256 x 256 results for LSUN categories



POTTEDPLANT

HORSE

SOFA

BUS

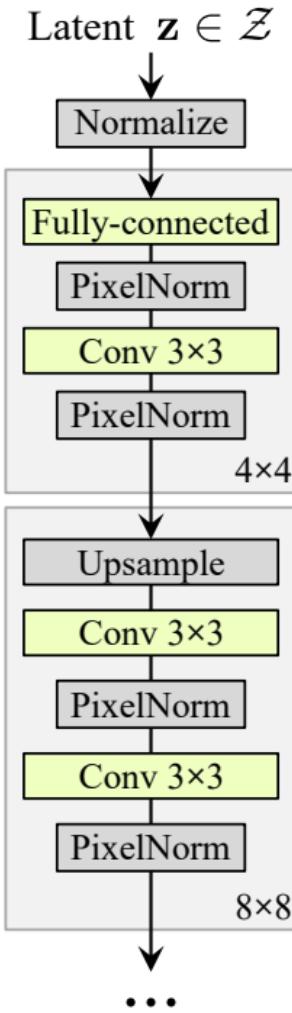
CHURCHOUTDOOR

BICYCLE

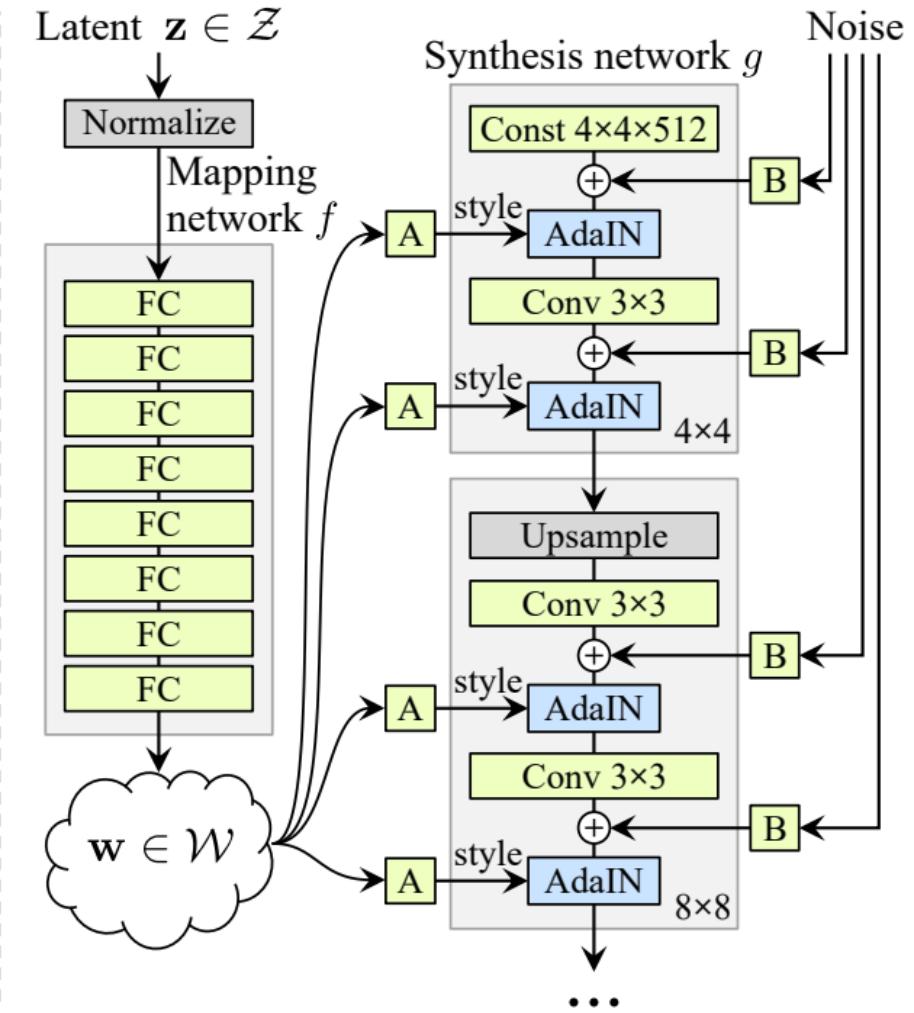
TVMONITOR

StyleGAN

- Built on top of Progressive GAN
- Start with learned constant (instead of noise vector)
- Use a mapping network to produce a *style code* w using learned affine transformations A
- Use *adaptive instance normalization* (AdaIN): scale and bias each feature map using learned style values
- Add noise after each convolution and before nonlinearity (enables stochastic detail)



(a) Traditional

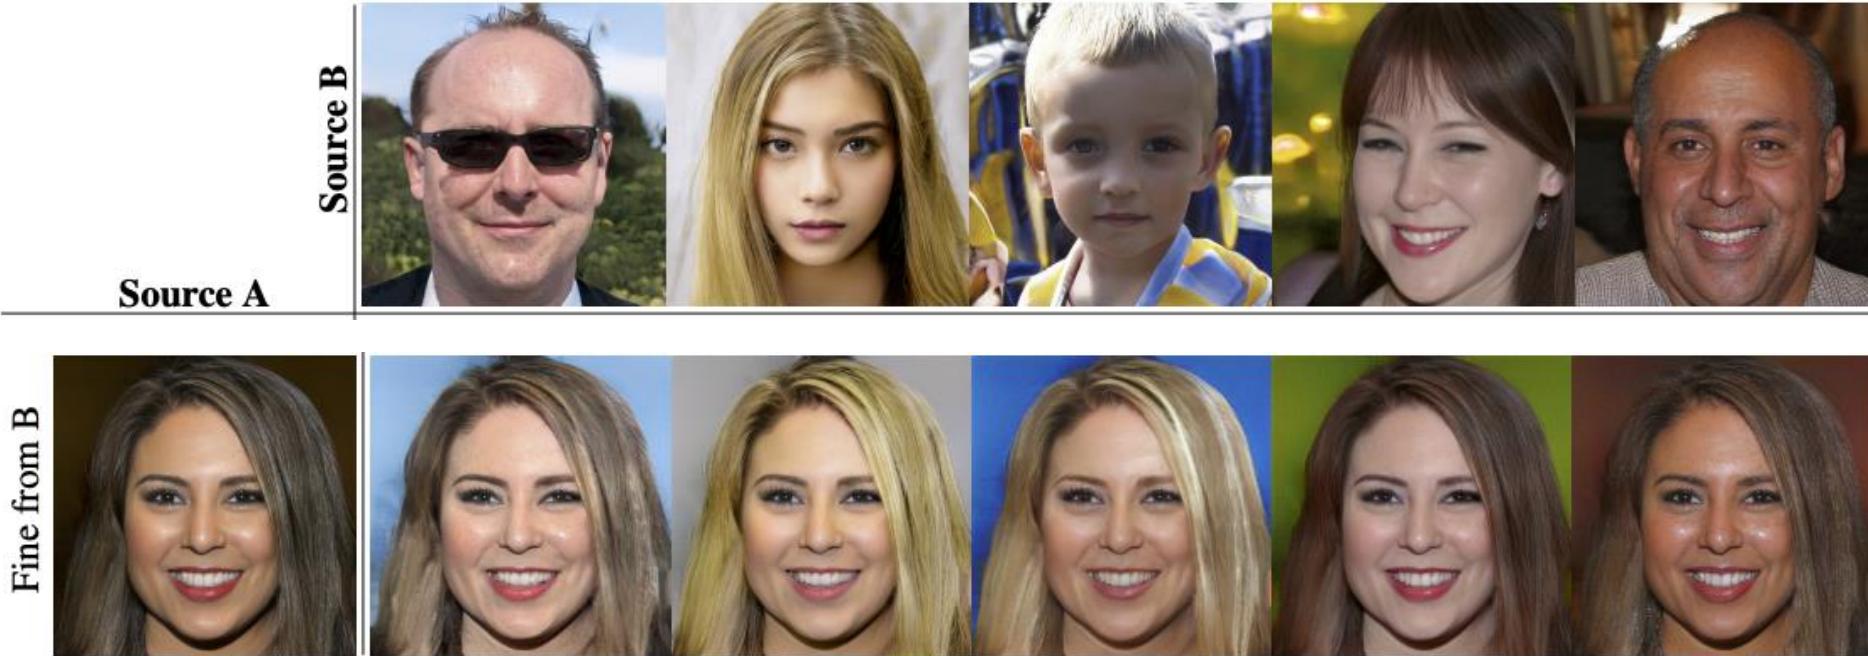


(b) Style-based generator

StyleGAN: Results



Mixing styles



“Two sets of images were generated from their respective latent codes (sources A and B); the rest of the images were generated by copying a specified subset of styles from source B and taking the rest from source A.”

Mixing styles

Middle styles from source B



Source A

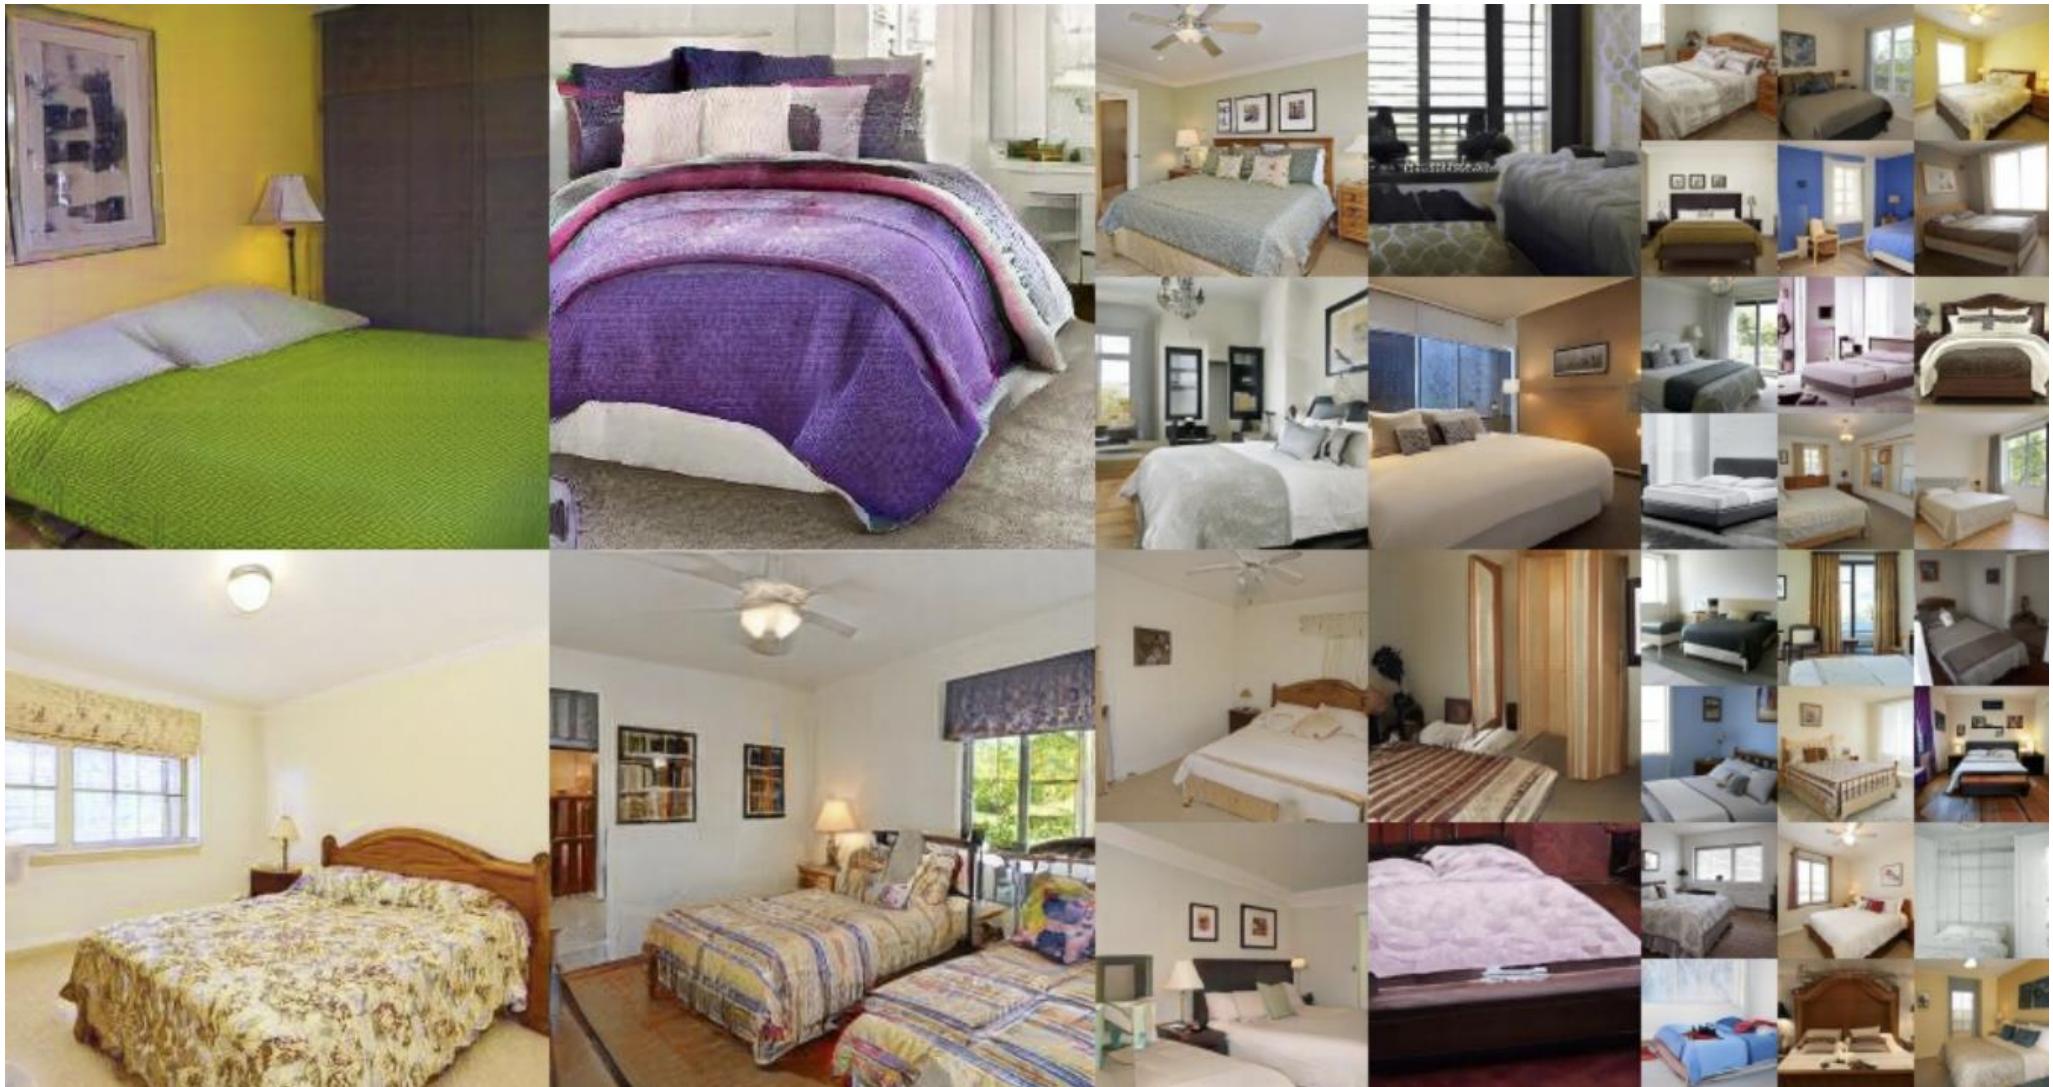
Source B

Mixing styles

Coarse styles from source B



StyleGAN: Bedrooms



StyleGAN: Cars



StyleGAN2

- Change normalization, remove progressive growing to address StyleGAN artifacts



Figure 1. Instance normalization causes water droplet -like artifacts in StyleGAN images. These are not always obvious in the generated images, but if we look at the activations inside the generator network, the problem is always there, in all feature maps starting from the 64x64 resolution. It is a systemic problem that plagues all StyleGAN images.



Figure 6. Progressive growing leads to “phase” artifacts. In this example the teeth do not follow the pose but stay aligned to the camera, as indicated by the blue line.

Outline

- Generative tasks
- Original GAN formulations
 - NSGAN
 - DCGAN
- Other popular formulations
 - WGAN, WGAN-GP
 - LSGAN
- State-of-the-art architectures
 - Progressive GAN, StyleGAN
- Evaluating GANs

How to evaluate GANs?

- Showing pictures of samples is not enough, especially for simpler datasets like MNIST, CIFAR, faces, bedrooms, etc.
- We cannot directly compute the likelihoods of high-dimensional samples (real or generated), or compare their distributions
- Many GAN approaches claim mainly to improve stability, which is hard to evaluate

GAN evaluation: Human studies

- Example: Turing test

Instructions

**Examples of real images**

**Examples of images generated by a computer**

We present you pictures that are either computer generated or are real photographs. Your task is to choose which one are which.

Images contain pictures of airplanes, automobiles, birds, cats, deer, dogs, frogs, horses, ships, and trucks. If you cannot clearly recognize what's the class of the object, then it's likely to be a generated image.

SET CHECKBOX ON IMAGES THAT LOOK LIKE GENERATED BY A COMPUTER.

| | | |
|--|--|--|
| <input type="checkbox"/>  | <input type="checkbox"/>  | <input type="checkbox"/>  |
| <input type="checkbox"/>  | <input type="checkbox"/>  | <input type="checkbox"/>  |
| <input type="checkbox"/>  | <input type="checkbox"/>  | <input type="checkbox"/>  |

Submit

GAN evaluation: Inception score (IS)

- Key idea: generators should produce images with a variety of recognizable object classes
- Defined as $IS(G) = \exp[\mathbb{E}_{x \sim G} KL(P(y|x) \parallel P(y))]$ where $P(y|x)$ is the posterior label distribution returned by an image classifier (e.g., InceptionNet) for sample x
 - If x contains a recognizable object, entropy of $P(y|x)$ should be low
 - If generator generates images of diverse objects, the marginal distribution $P(y)$ should have high entropy

GAN evaluation: Inception score (IS)

- Disadvantages
 - A GAN that simply memorizes the training data (overfitting) or outputs a single image per class (mode dropping) could still score well
 - Is sensitive to network weights, not necessarily valid for generative models not trained on ImageNet, can be gamed ([Barratt & Sharma 2018](#))



Figure 1. Sample of generated images achieving an Inception Score of 900.15. The maximum achievable Inception Score is 1000, and the highest achieved in the literature is on the order of 10.

GAN evaluation: Fréchet Inception Distance (FID)

- Key idea: fit simple distributions (Gaussians) to statistics of feature activations for real and generated data; estimate divergence parametrically
 - Pass generated samples through a network (InceptionNet), compute activations for a chosen layer
 - Estimate multivariate mean and covariance of activations, compute Fréchet distance to those of real data

GAN evaluation: Fréchet Inception Distance (FID)

- Key idea: fit simple distributions (Gaussians) to statistics of feature activations for real and generated data; estimate divergence parametrically
 - Pass generated samples through a network (InceptionNet), compute activations for a chosen layer
 - Estimate multivariate mean and covariance of activations, compute Fréchet distance to those of real data

In mathematics, the **Fréchet distance** is a measure of similarity between curves that takes into account the location and ordering of the points along the curves. It is named after Maurice Fréchet.

----- Source: en.wikipedia.org

GAN evaluation: Fréchet Inception Distance (FID)

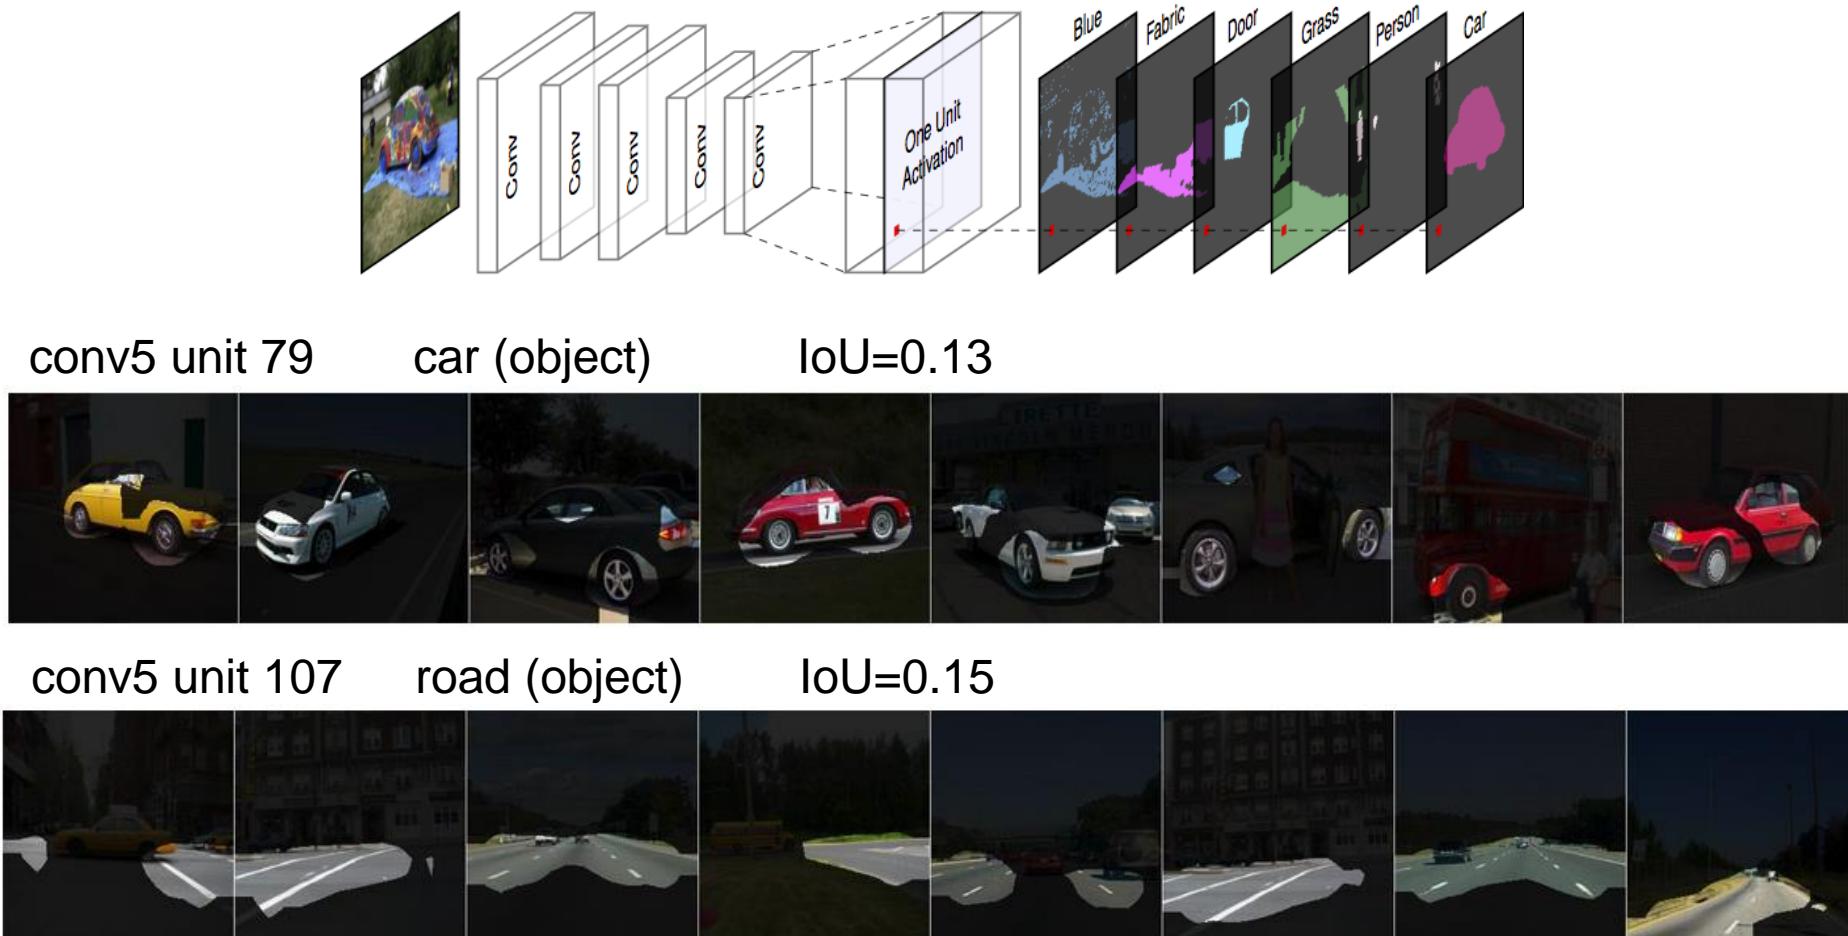
- Key idea: fit simple distributions (Gaussians) to statistics of feature activations for real and generated data; estimate divergence parametrically
 - Pass generated samples through a network (InceptionNet), compute activations for a chosen layer
 - Estimate multivariate mean and covariance of activations, compute Fréchet distance to those of real data
- Advantages: correlated with visual quality of samples and human judgment, can detect mode dropping (unlike IS)
- Disadvantage: cannot detect overfitting (like IS)

Outline

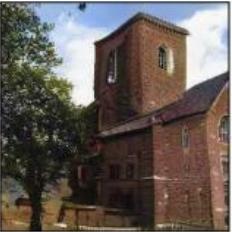
- Generative tasks
- Original GAN formulations
 - NSGAN
 - DCGAN
- Other popular formulations
 - WGAN, WGAN-GP
 - LSGAN
- State-of-the-art architectures
 - Progressive GAN, StyleGAN
- Evaluating GANs
- Visualizing and controlling GANs

GAN Dissection

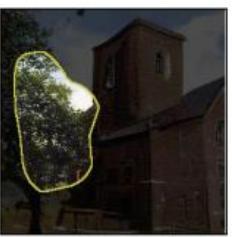
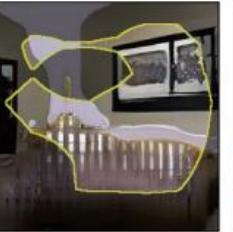
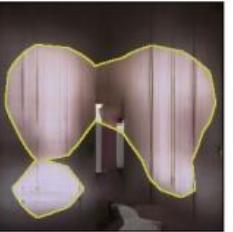
- Recall: network dissection



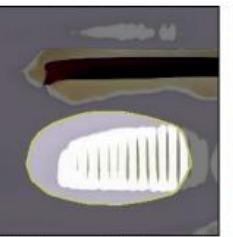
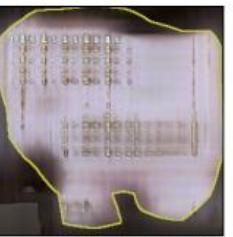
GAN Dissection



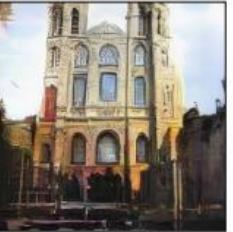
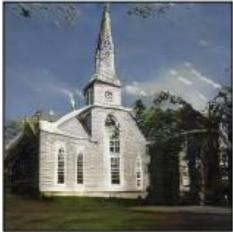
(a) Generate images of churches



(b) Identify GAN units that match trees



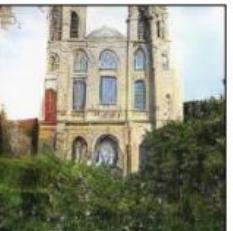
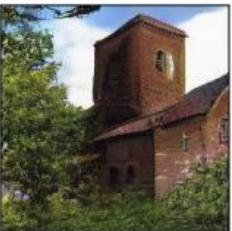
(e) Identify GAN units that cause artifacts



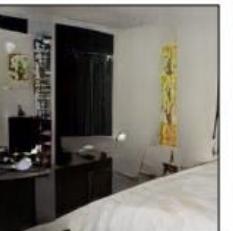
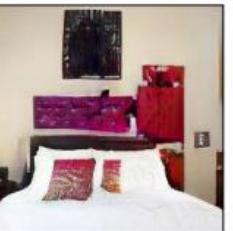
(c) Ablating units removes trees



(f) Bedroom images with artifacts



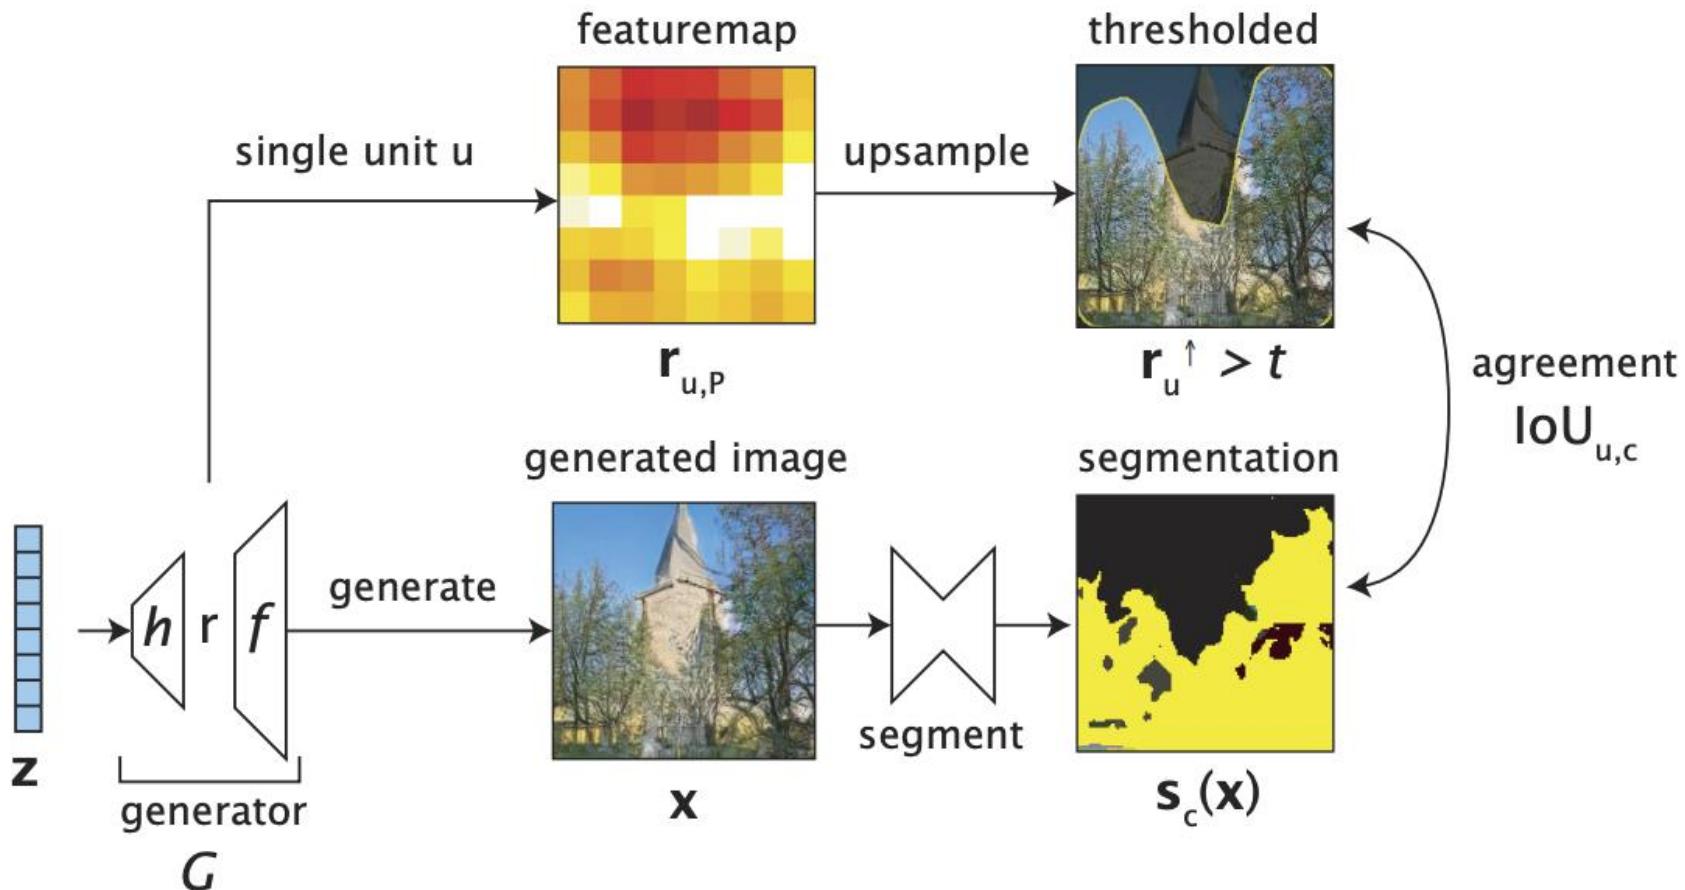
(d) Activating units adds trees



(g) Ablating “artifact” units improves results

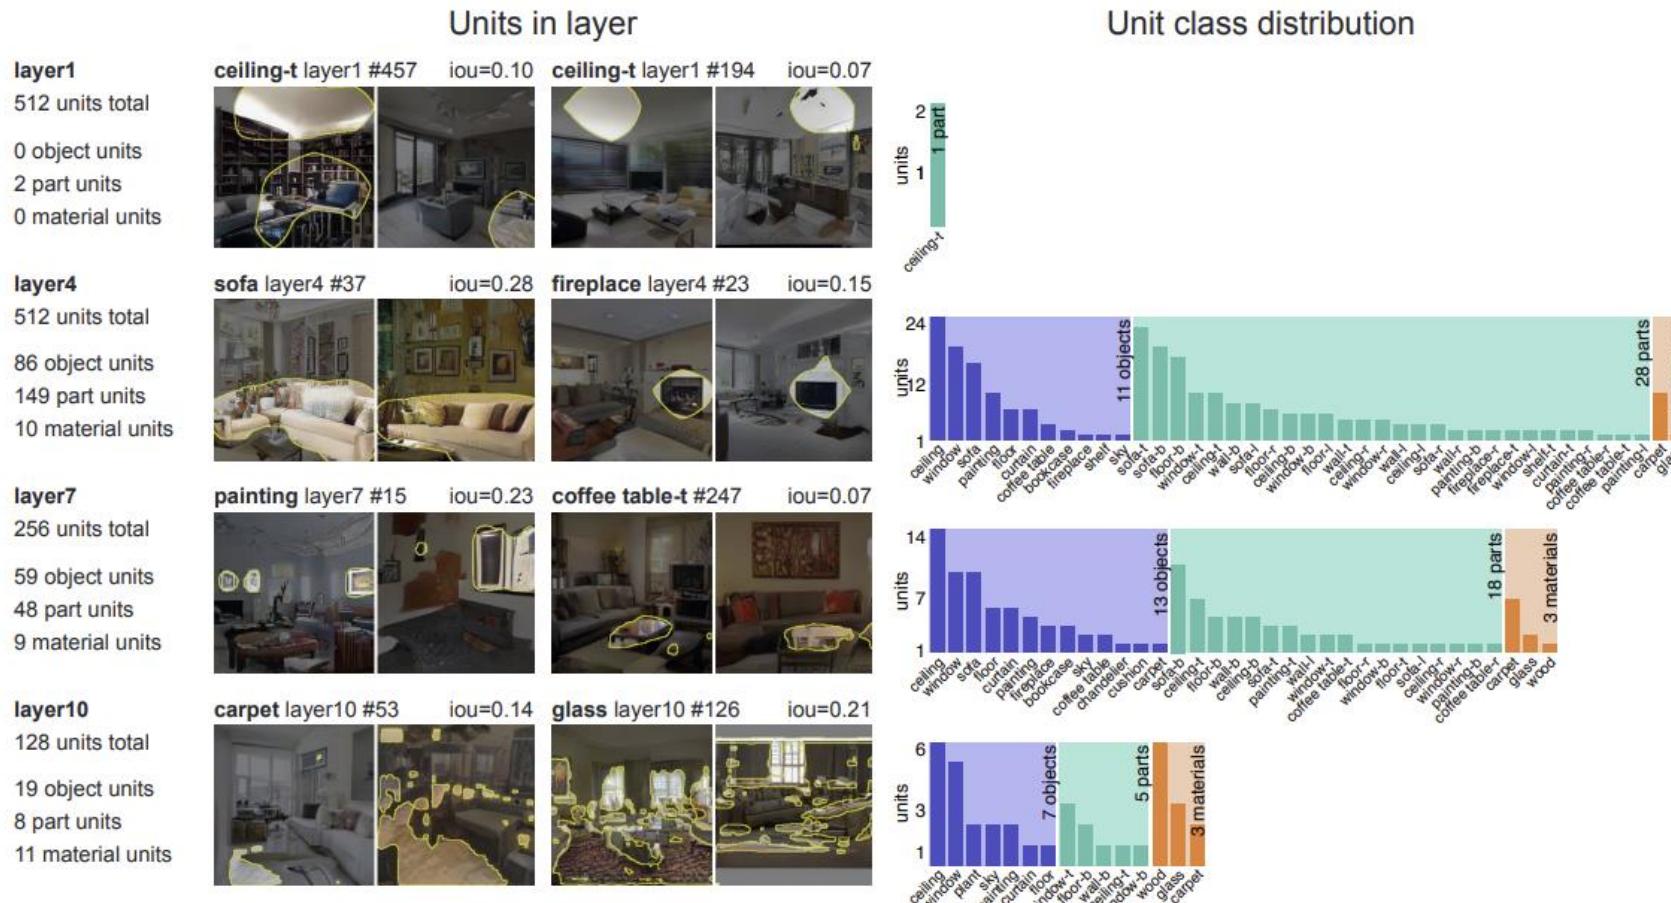
GAN Dissection

- Dissection:



GAN Dissection

- Interpreting units at different levels of a progressive GAN (trained on “bedroom”):



GANPaint demo



Input photo

Remove chairs

Output result

Input photo

Add windows

Output result



Input photo

Change rooftops

Output result

Input photo

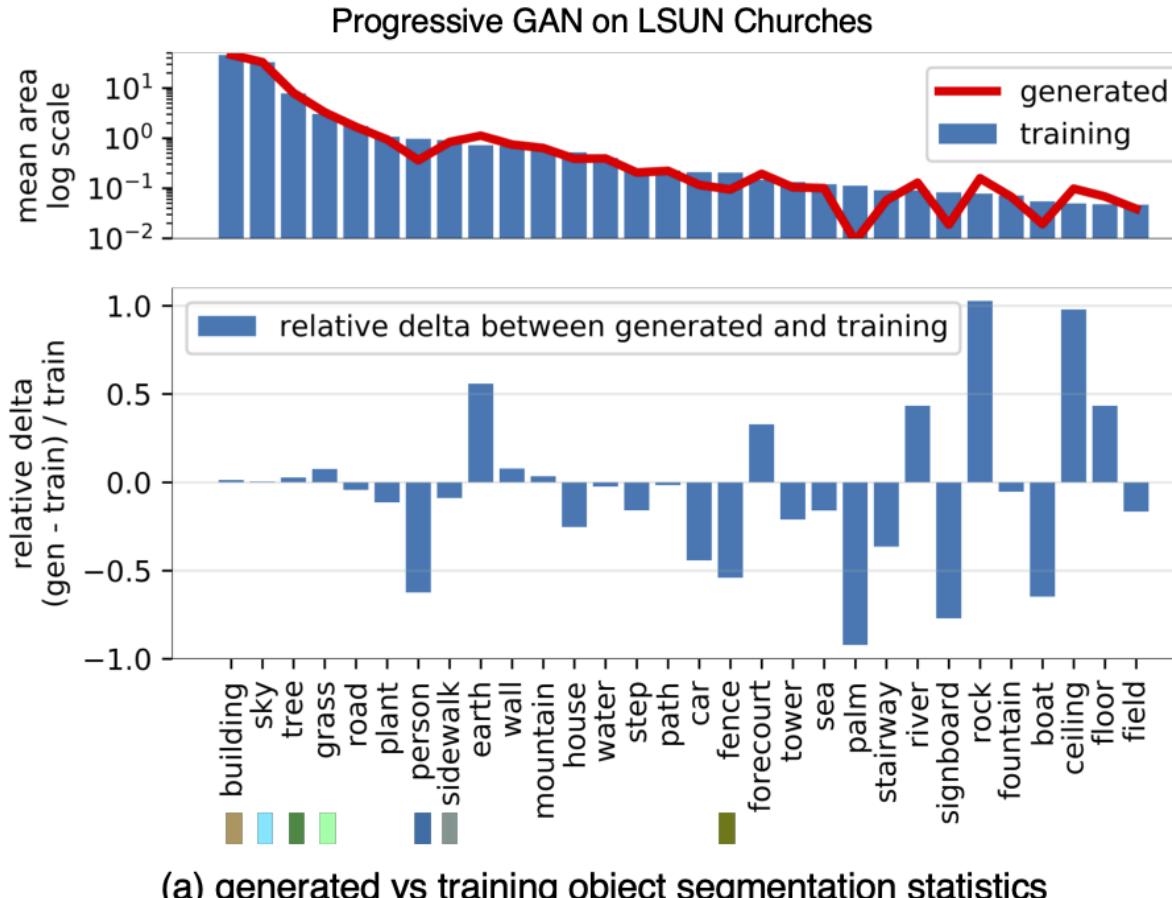
Restyle trees for spring

Restyle trees for autumn

<https://ganpaint.io/demo/?project=church>

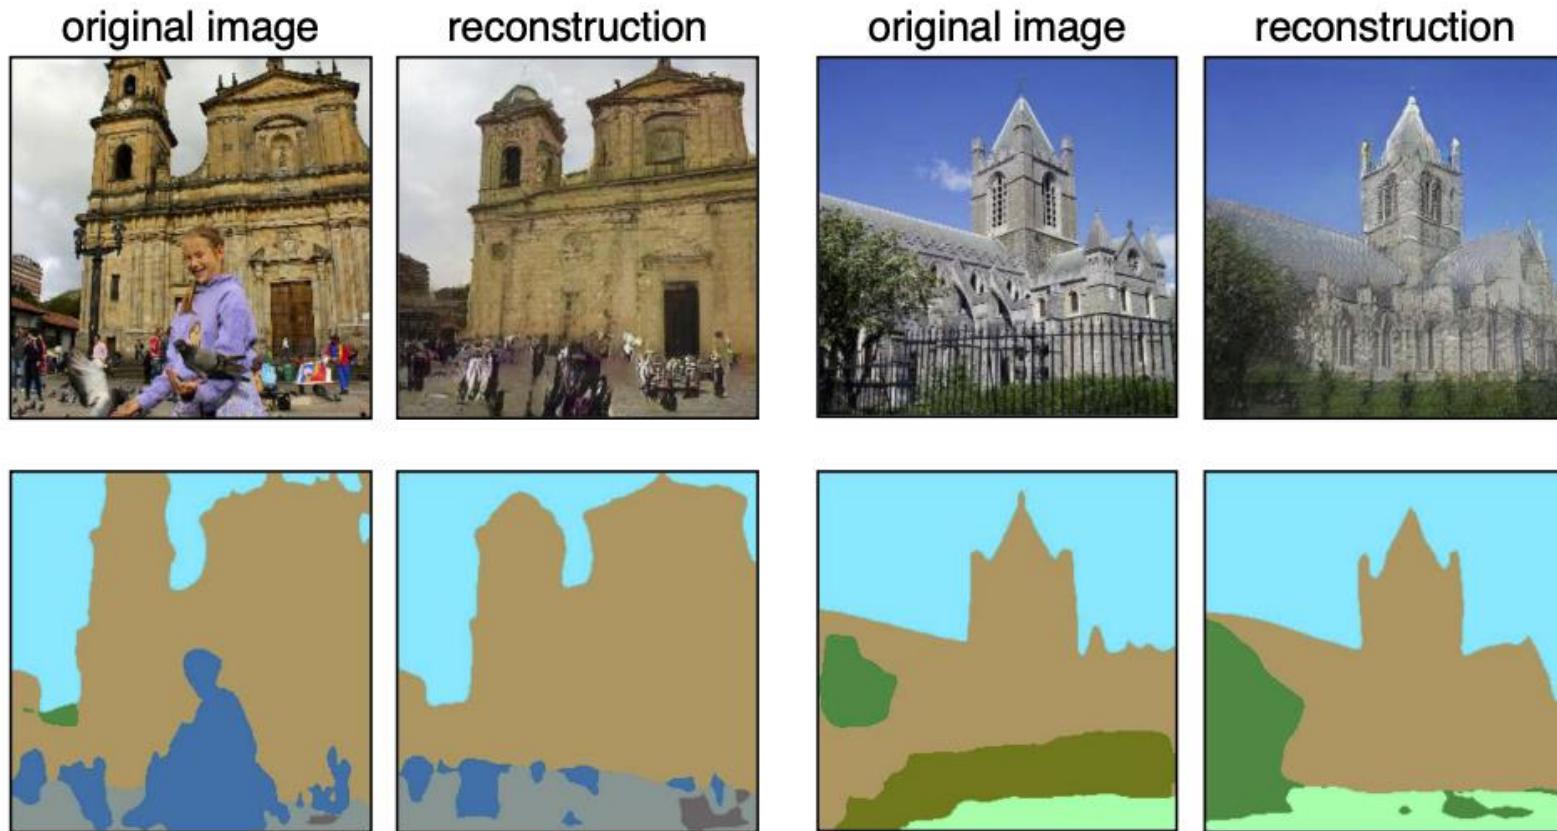
Seeing what a GAN cannot generate

- Compare semantic segmentations of real and generated images, see which classes get dropped by the GAN



Seeing what a GAN cannot generate

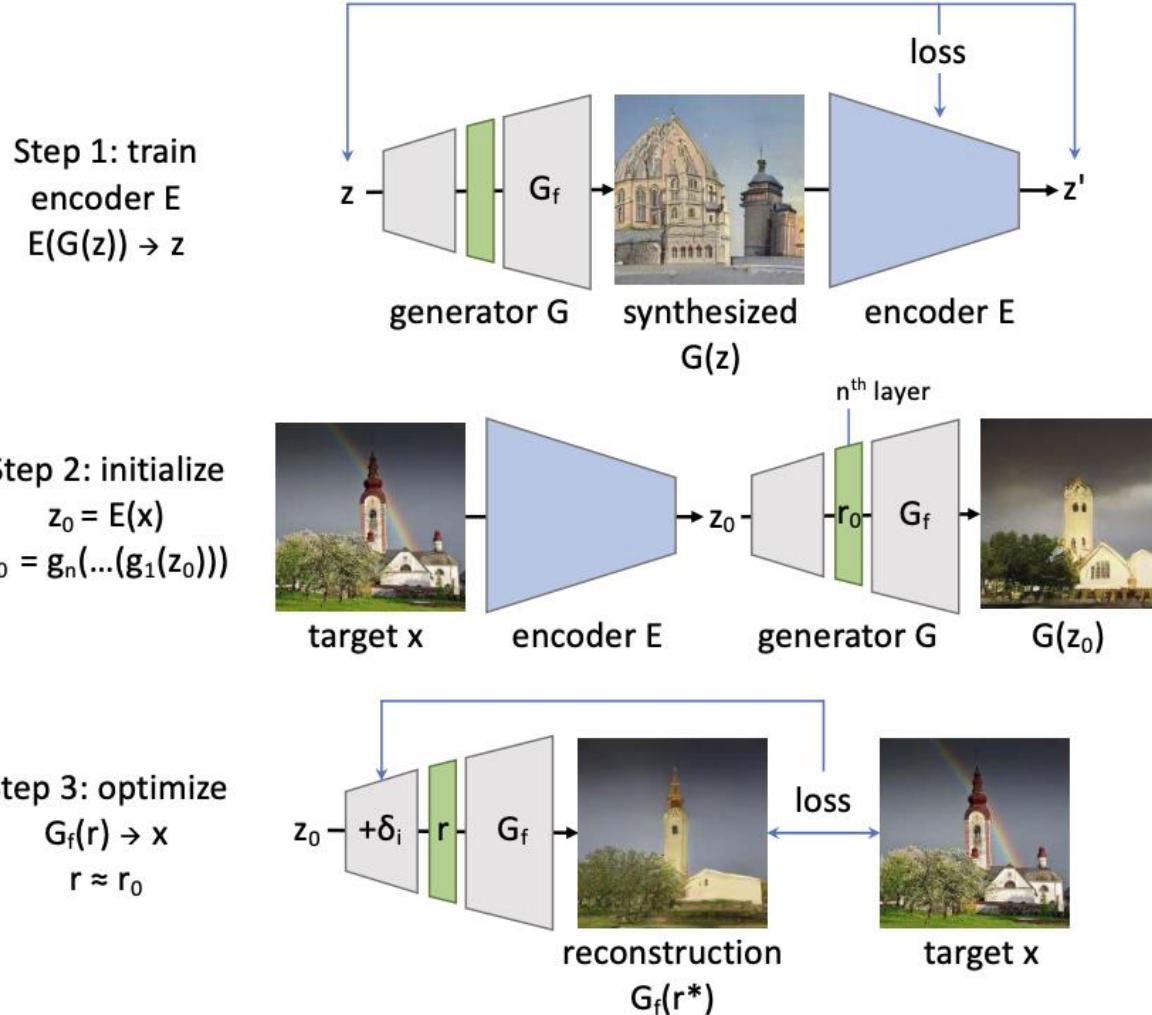
- Given real images, try to find “closest” generated images and see what gets omitted



(b) real images vs. reconstructions

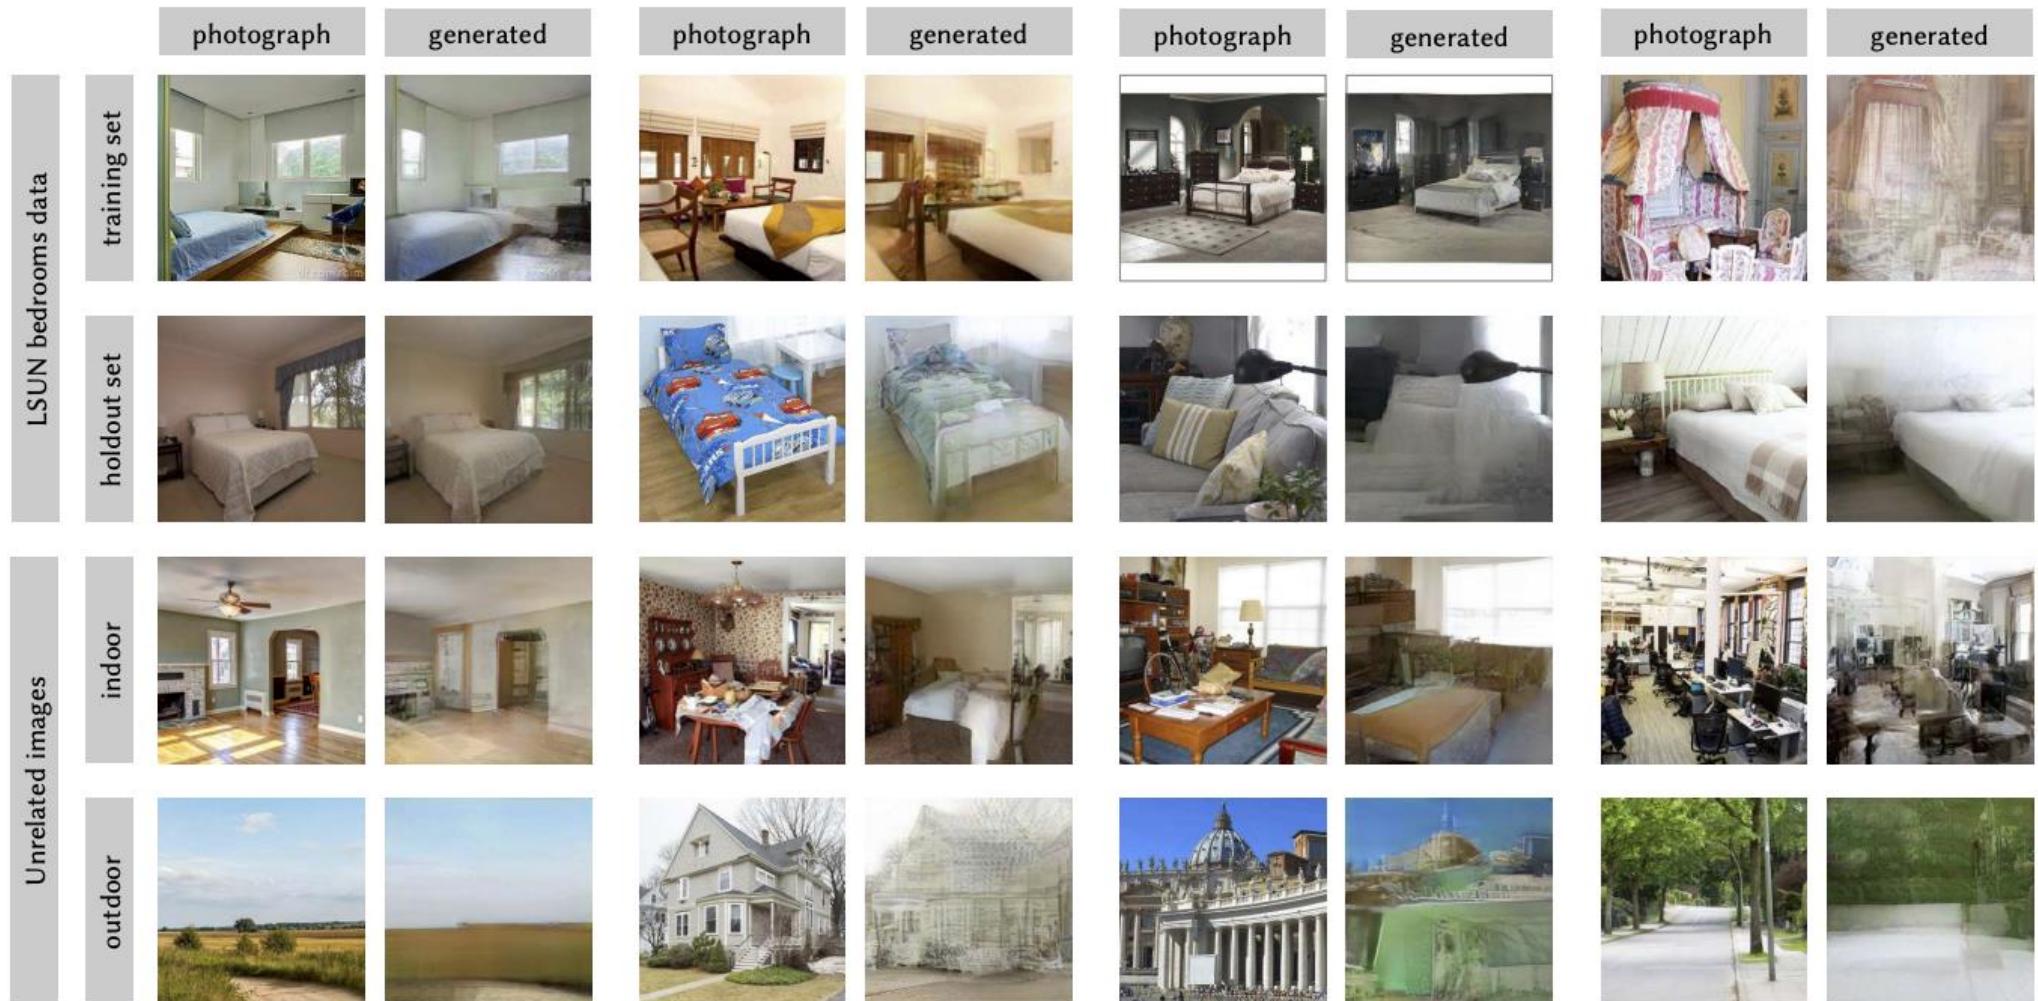
Seeing what a GAN cannot generate

- Layer inversion method:



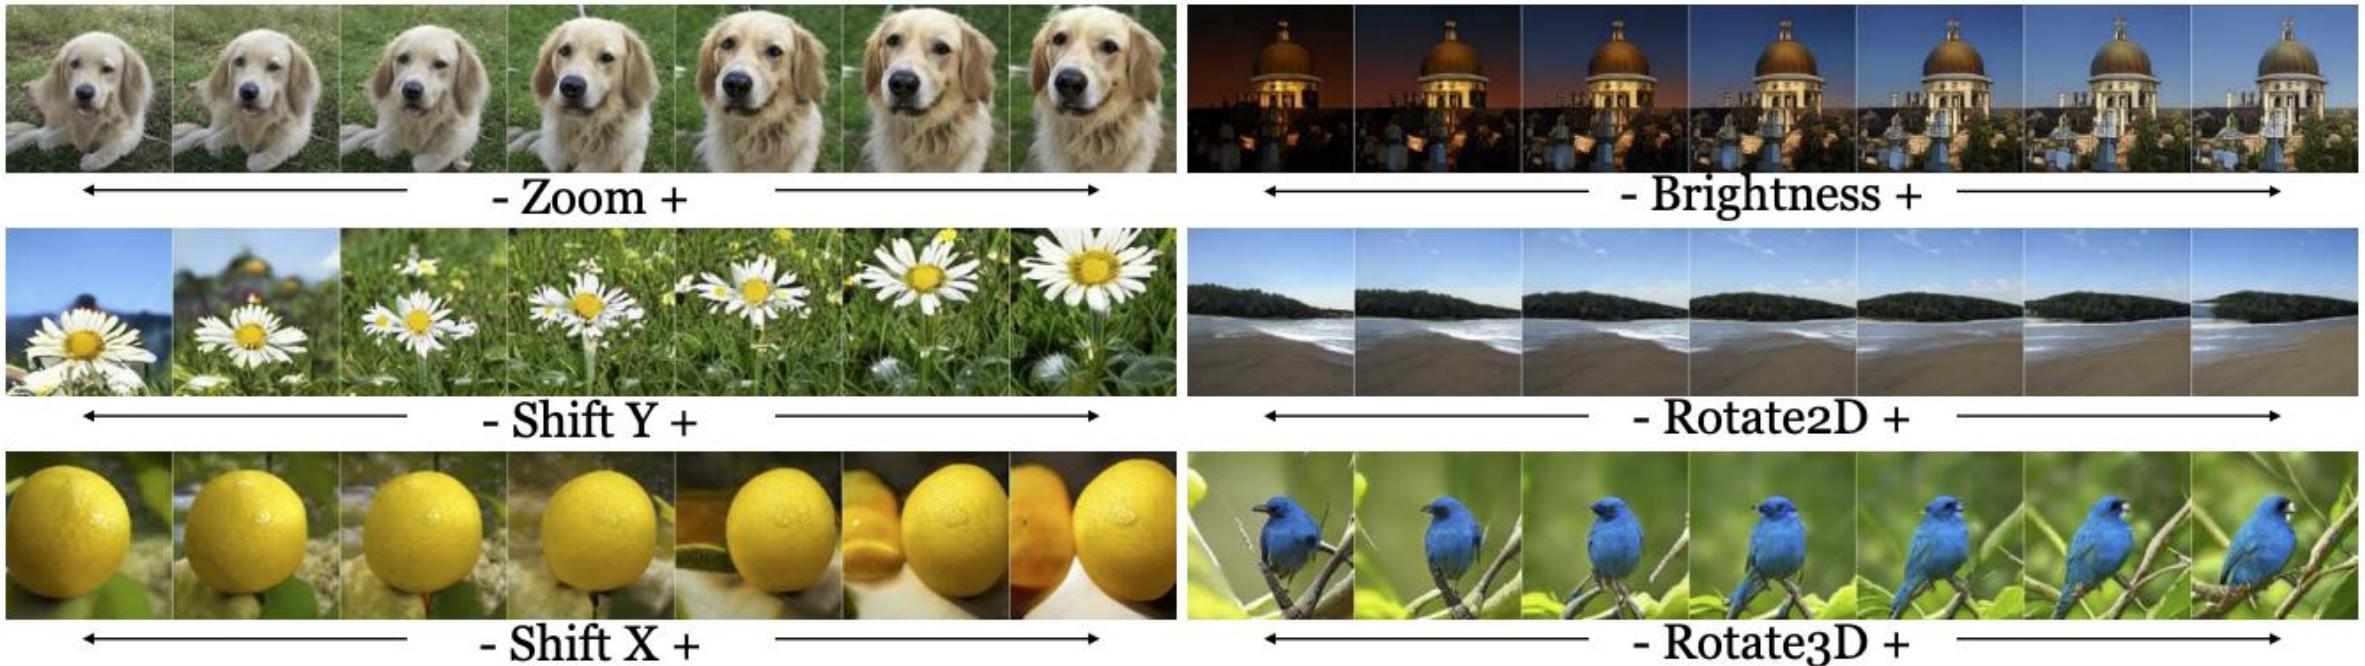
Seeing what a GAN cannot generate

- Reconstruction results for “bedroom” GAN



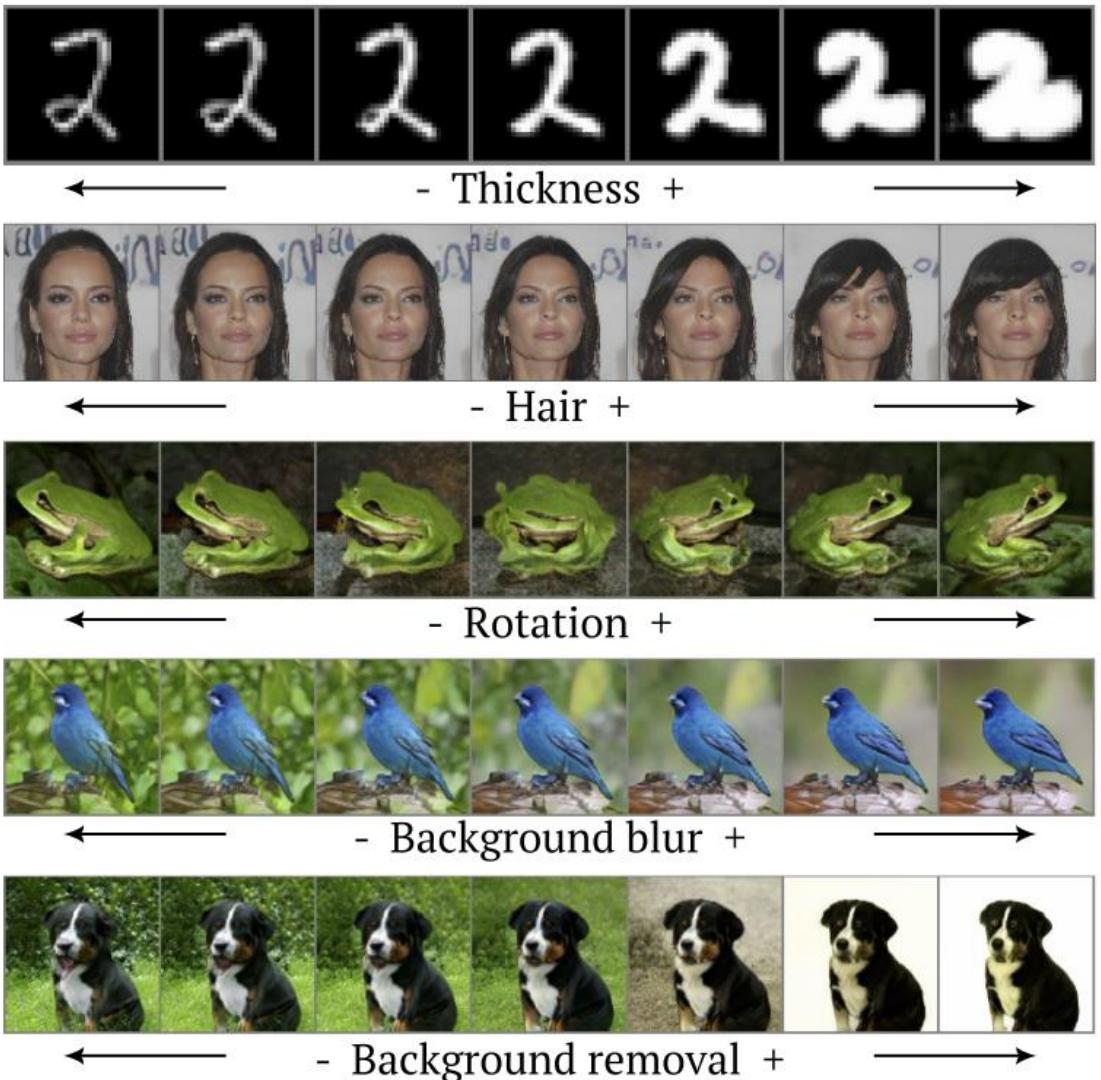
Traversing the GAN latent space

- Try to find simple “walks” in the latent space of GANs to achieve various meaningful transformations to explore structure of that space and test GANs’ ability to interpolate between training samples



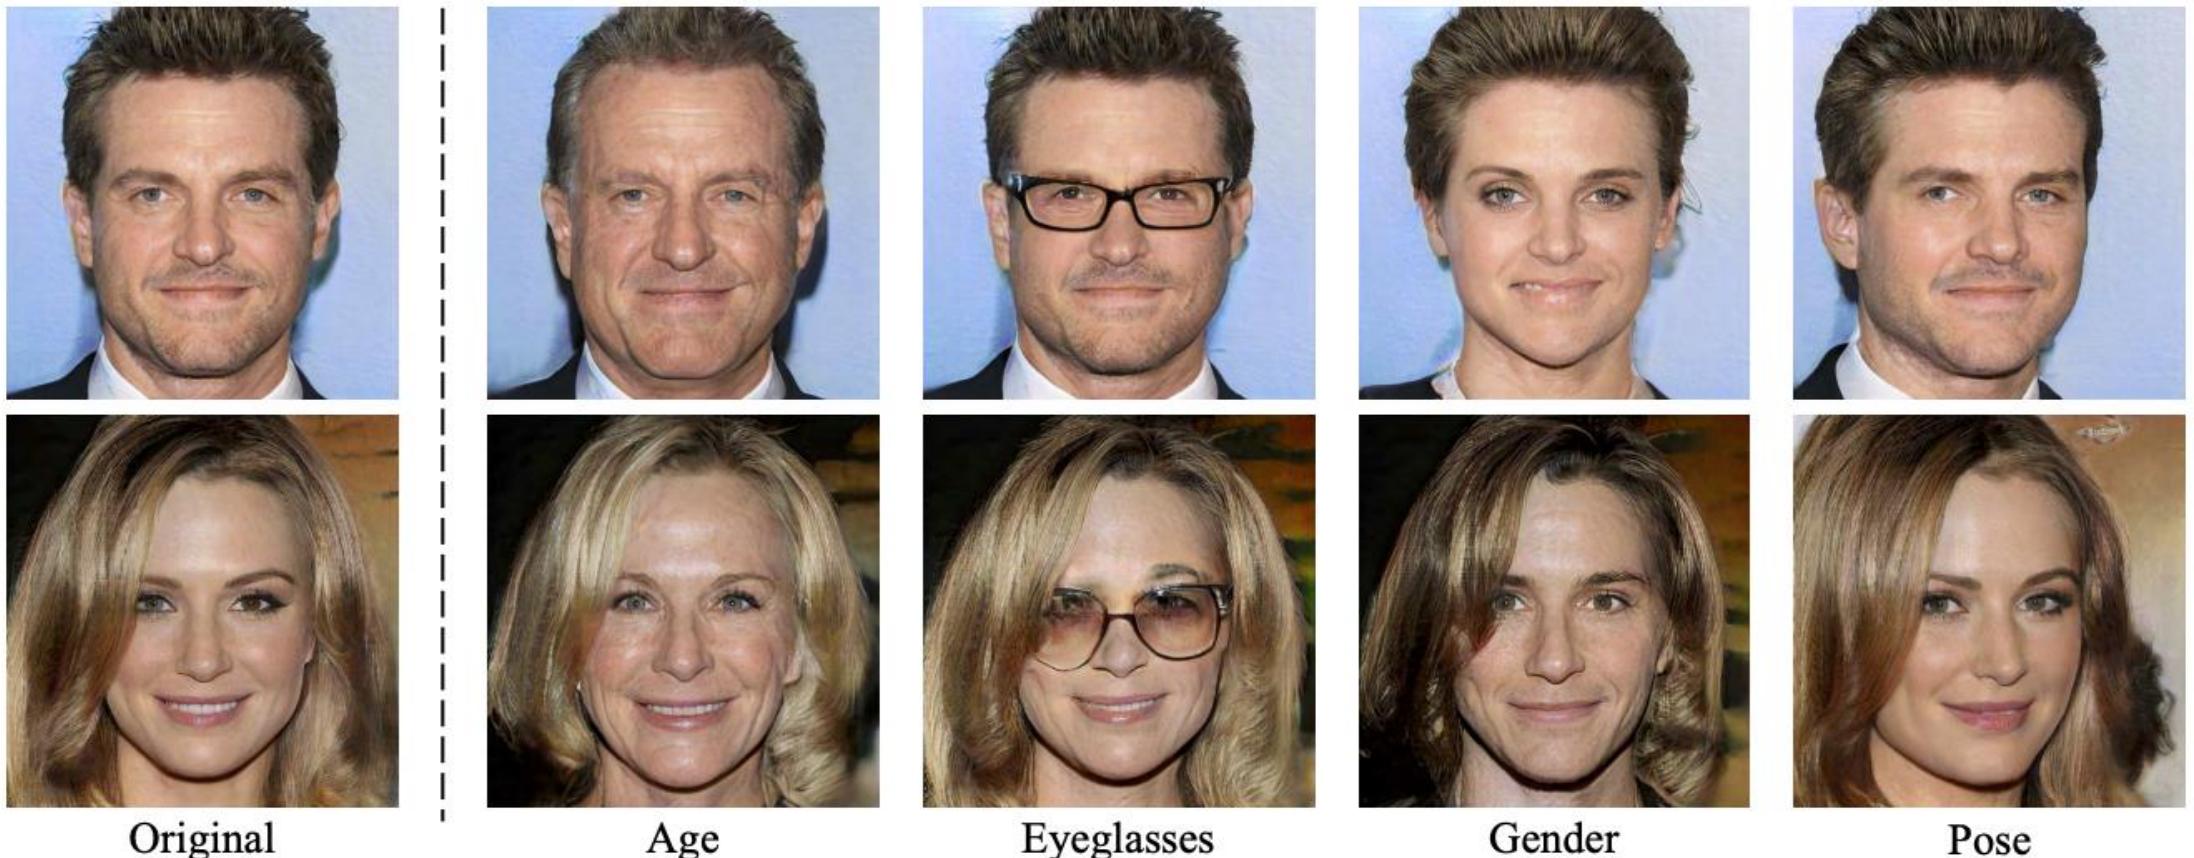
Traversing the GAN latent space

- Goal: learn a set of directions inducing “disentangled” image transformations that are easy to distinguish from each other



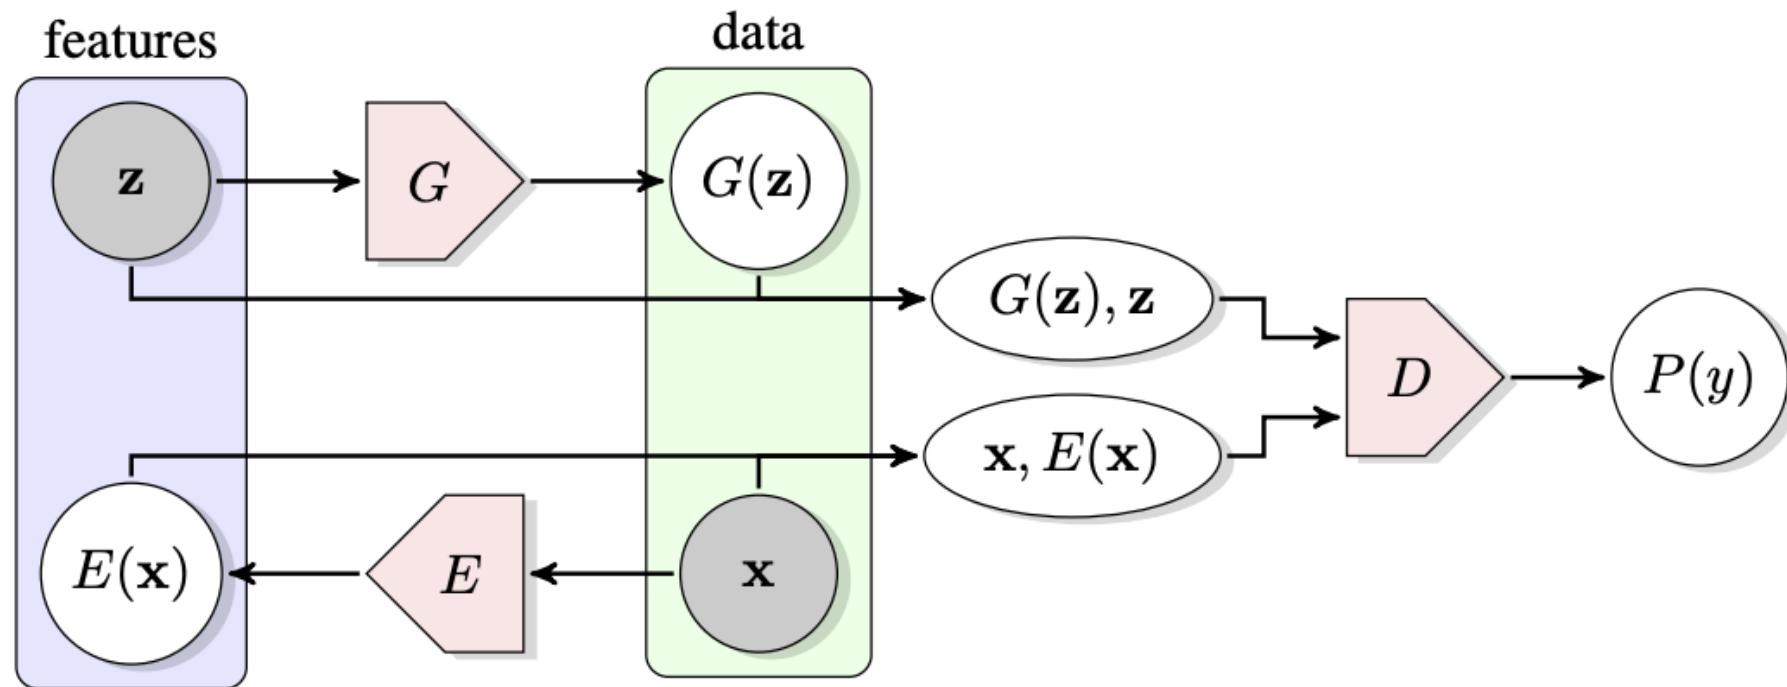
GAN editing

- Supervised training to find latent space directions corresponding to pose, smile, age, gender, eyeglasses



GANs for representation learning

- Bidirectional GAN (BiGAN): simultaneously train generator and encoder (mapping from images to z vectors or approximate inverse of the generator), show that the encoder creates a latent representation useful for other tasks



GANs for representation learning: BigBiGAN

- Train BiGAN with BigGAN architecture on ImageNet
- Show that bidirectional framework improves image generation
- Encoder representation gives results comparable to state-of-the-art self-supervised models on ImageNet classification

GANs for representation learning: BigBiGAN

Real images x



Reconstructions $G(E(x))$

Acknowledgement

Thanks to the following courses and corresponding researchers for making their teaching/research material online

- Deep Learning, Stanford University
- Introduction to Deep Learning, University of Illinois at Urbana-Champaign
- Introduction to Deep Learning, Carnegie Mellon University
- Convolutional Neural Networks for Visual Recognition, Stanford University
- Natural Language Processing with Deep Learning, Stanford University
- And Many More