

Assignment 1 Part 3A Introduction: Multi-label Image Classification

```

1 # Steps to load extra files and download dataset.
2 import os
3 !git clone "https://vipulrawat008:Invincible007!@github.com/the-visionary/assignment1"
4 os.chdir('/content')
5 os.chdir('assignment1')
6 os.chdir('assignment1-part3')
7 assert (os.getcwd() == '/content/assignment1/assignment1-part3')
8
9 !bash download_data.sh
10 assert (os.getcwd() == '/content/assignment1/assignment1-part3')

```

```

--2021-02-26 14:36:46-- http://host.robots.ox.ac.uk/pascal/VOC/voc2007/VOCTrainval_06-Nov-2007.tar
Resolving host.robots.ox.ac.uk (host.robots.ox.ac.uk)... 129.67.94.152
Connecting to host.robots.ox.ac.uk (host.robots.ox.ac.uk)|129.67.94.152|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 460032000 (439M) [application/x-tar]
Saving to: 'VOCTrainval_06-Nov-2007.tar.1'

.tgz.1          16%[==>]           70.20M  11.6MB/s   eta 39s    ^C

```

```

1 os.chdir('/content')
2 os.chdir('assignment1')
3 os.chdir('assignment1-part3')

```

```
1 !pip install torch==1.5.1 torchvision==0.6.1
```

```

Collecting torch==1.5.1
  Using cached https://files.pythonhosted.org/packages/a4/cf/007b6de316c9f3d4cb315a60c308342cc...
Collecting torchvision==0.6.1
  Downloading https://files.pythonhosted.org/packages/f4/ef/c32275c040f12f24adcd6302f0f3cd12cc...
    |██████████| 6.6MB 3.9MB/s
Requirement already satisfied: future in /usr/local/lib/python3.7/dist-packages (from torch==1.5.1)
Requirement already satisfied: numpy in /usr/local/lib/python3.7/dist-packages (from torch==1.5.1)
Requirement already satisfied: pillow>=4.1.1 in /usr/local/lib/python3.7/dist-packages (from torch==1.5.1)
Installing collected packages: torch, torchvision
  Found existing installation: torchvision 0.8.1+cu101
    Uninstalling torchvision-0.8.1+cu101:
      Successfully uninstalled torchvision-0.8.1+cu101
Successfully installed torch-1.5.1 torchvision-0.6.1

```

```

1 import os
2 import numpy as np
3 import torch
4 import torch.nn as nn
5 import torchvision
6
7 from torchvision import transforms
8 from sklearn.metrics import average_precision_score
9 from PIL import Image, ImageDraw

```

```
10 # from kaggle_submission import output_submission_csv
11 import matplotlib.pyplot as plt
12 from classifier import SimpleClassifier, Classifier#, AlexNet
13 from voc_dataloader import VocDataset, VOC_CLASSES
14
15 %matplotlib inline
16 %load_ext autoreload
17 %autoreload 2
```

```
The autoreload extension is already loaded. To reload it, use  
%reload_ext autoreload
```

▼ Multi-label Classification

In this assignment, you train a classifier to do multi-label classification on the PASCAL VOC 2007 dataset. This means that there are multiple classes which can appear in any given image. Your classifier will predict whether each class appears in an image. This is different from an exclusive multiclass classification like the ImageNet competition where only a single most appropriate class is predicted.

Part 3A

You will use this notebook to warm up with pytorch

What to do

In part 3A, You are asked to run below experiments. You don't need to change hyperparameters for this Part code provides everything that you will need.)

1. to train a simple network (defined in `classifiers.py`)
 2. to train the AlexNet (PyTorch built-in)
 - from scratch
 - finetuning AlexNet pretrained on ImageNet

What to submit

We ask you to run the following code and report the results in your homework submission. You may want to familiar with PyTorch.

You will need the numbers and plots this notebook outputs for reports, but you are not required to submit

▼ Reading Pascal Data

▼ Loading Training Data

In the following cell we will load the training data and also apply some transforms to the data.

```
4 train_transform = transforms.Compose([
5     transforms.Resize(227),
6     transforms.CenterCrop(227),
7     transforms.ToTensor(),
8     normalize
9 ])
10 ])
11
```

```
1 ds_train = VocDataset('VOCdevkit_2007/VOC2007/','train',train_transform)
```

```
/content/assignment1/assignment1-part3/voc_dataloader.py:109: VisibleDeprecationWarning: Creat
    return np.array(names), np.array(labels).astype(np.float32), np.array(box_indices), label_order
```

▼ Loading Validation Data

We will load the test data for the PASCAL VOC 2007 dataset. Do **NOT** add data augmentation transforms to

```
1 # Transforms applied to the testing data
2 test_transform = transforms.Compose([
3     transforms.Resize(227),
4     transforms.CenterCrop(227),
5     transforms.ToTensor(),
6     normalize,
7 ])
8
```

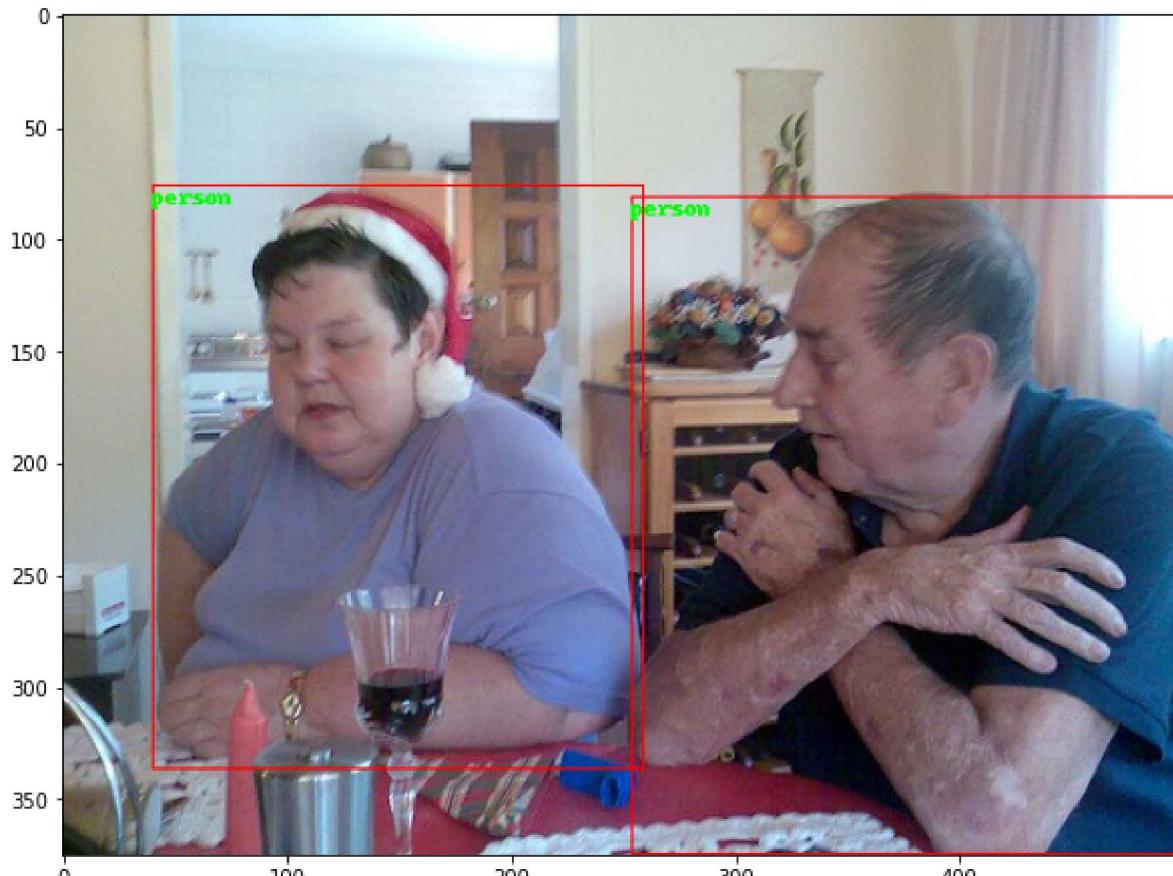
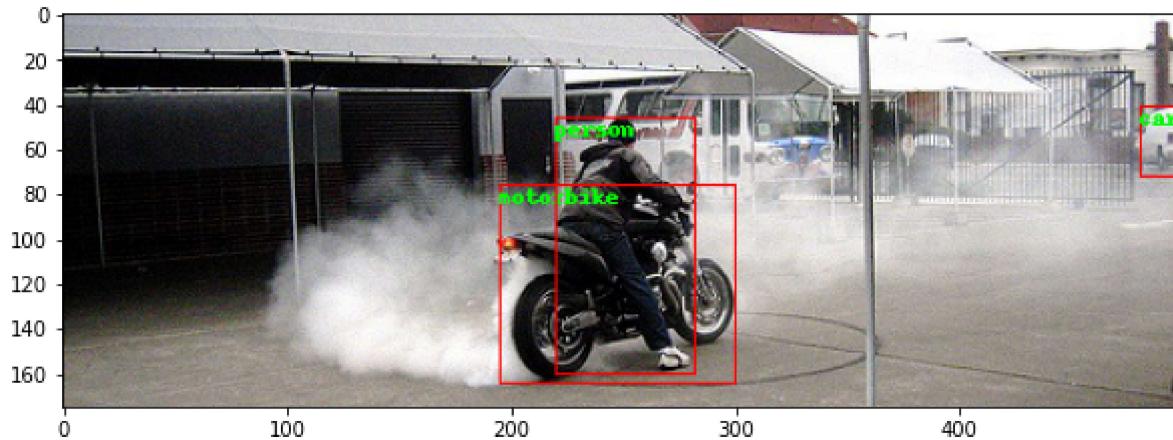
```
1 ds_val = VocDataset('VOCdevkit_2007/VOC2007/','val',test_transform)
```

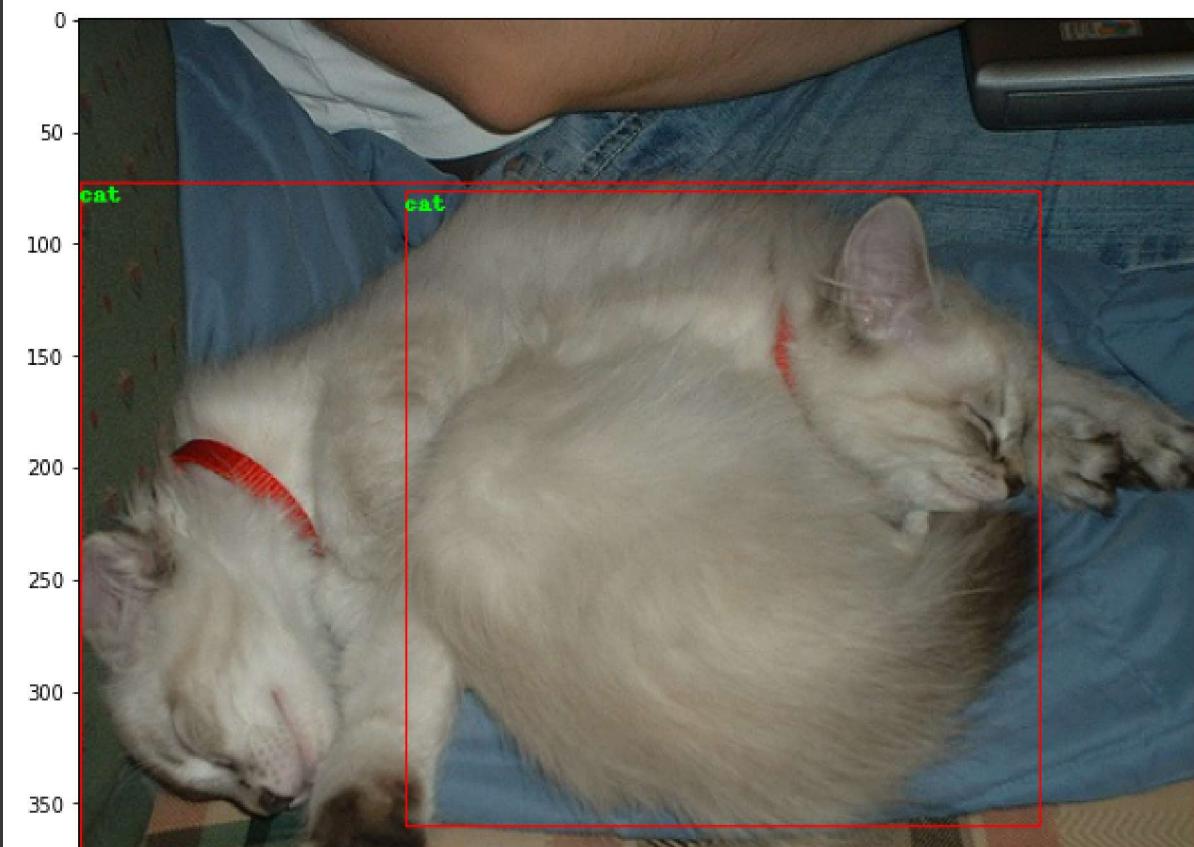
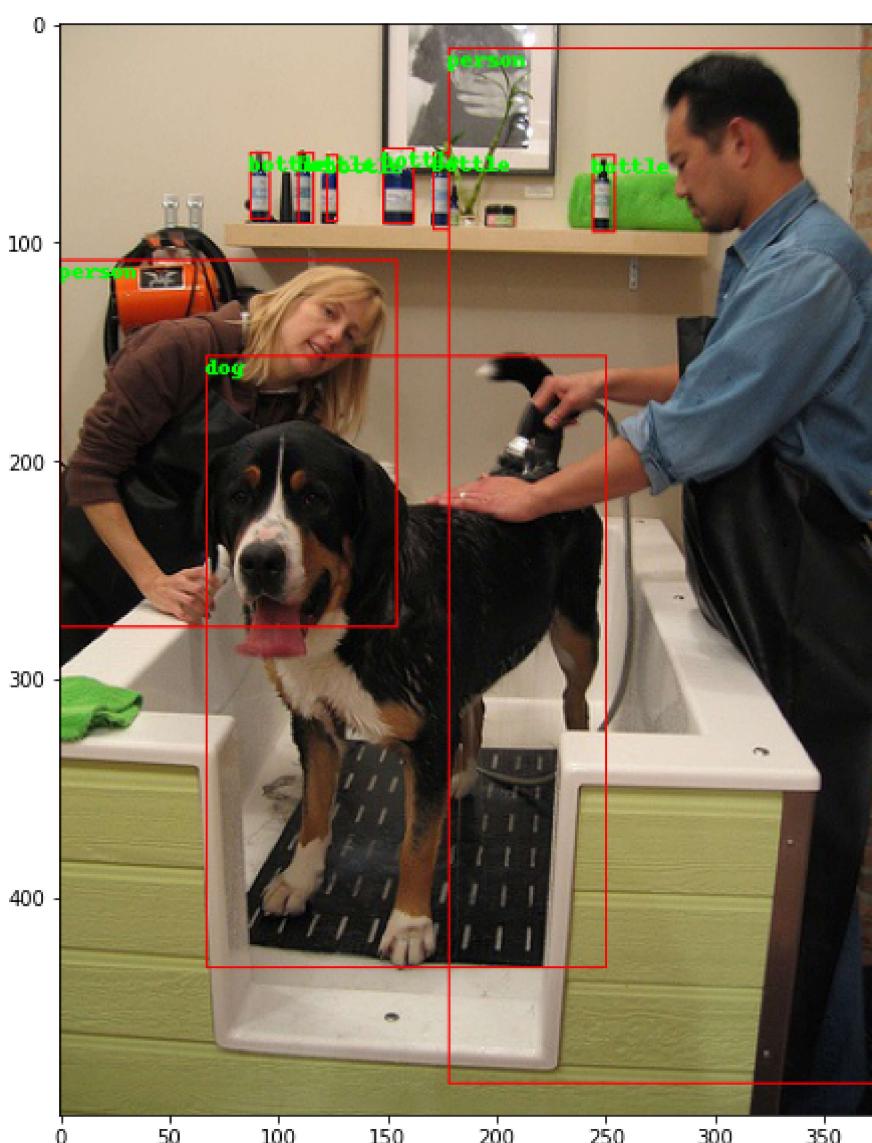
```
/content/assignment1/assignment1-part3/voc_dataloader.py:109: VisibleDeprecationWarning: Creat
    return np.array(names), np.array(labels).astype(np.float32), np.array(box_indices), label_order
```

▼ Visualizing the Data

PASCAL VOC has bounding box annotations in addition to class labels. Use the following code to visualize some corresponding annotations from the train set.

```
1 for i in range(5):
2     idx = np.random.randint(0, len(ds_train.names)+1)
3     _imgpath = os.path.join('VOCdevkit_2007/VOC2007/','JPEGImages', ds_train.names[
4     img = Image.open(_imgpath).convert('RGB')
5     draw = ImageDraw.Draw(img)
6     for j in range(len(ds_train.box_indices[idx])):
7         obj = ds_train.box_indices[idx][j]
8         draw.rectangle(list(obj), outline=(255,0,0))
9         draw.text(list(obj[0:2]), ds_train.classes[ds_train.label_order[idx][j]], fill='white')
10    plt.figure(figsize = (10,10))
11    plt.imshow(np.array(img))
```





▼ Classification

```

1 # def collate_fn(batch):
2 #     return tuple(zip(*batch))
3

1 # declare what device to use: gpu/cpu
2 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")

1 train_loader = torch.utils.data.DataLoader(dataset=ds_train,
2                                             batch_size=50,
3                                             shuffle=True,
4                                             #     collate_fn=lambda x: x,
5                                             num_workers=1)

1 val_loader = torch.utils.data.DataLoader(dataset=ds_val,
2                                         batch_size=50,
3                                         shuffle=True,
4                                         #     collate_fn=lambda x: x,
5                                         num_workers=1)

1 def train_classifier(train_loader, classifier, criterion, optimizer):
2     classifier.train()
3     loss_ = 0.0
4     losses = []
5     for i, (images, labels, _) in enumerate(train_loader):
6         images, labels = images.to(device), labels.to(device)
7         optimizer.zero_grad()
8         logits = classifier(images)
9         loss = criterion(logits, labels)
10        loss.backward()
11        optimizer.step()
12        losses.append(loss)
13    return torch.stack(losses).mean().item()

1 def test_classifier(test_loader, classifier, criterion, print_ind_classes=True, prin
2     classifier.eval()
3     losses = []
4     with torch.no_grad():
5         y_true = np.zeros((0,21))
6         y_score = np.zeros((0,21))
7         for i, (images, labels, _) in enumerate(test_loader):
8             images, labels = images.to(device), labels.to(device)
9             logits = classifier(images)
10            y_true = np.concatenate((y_true, labels.cpu().numpy()), axis=0)
11            y_score = np.concatenate((y_score, logits.cpu().numpy()), axis=0)
12            loss = criterion(logits, labels)
13            losses.append(loss.item())
14        aps = []
15        # ignore first class which is background
16        for i in range(1, y_true.shape[1]):
17            ap = average_precision_score(y_true[:, i], y_score[:, i])
18            if print_ind_classes:

```

```

19         print('----- Class: {:<12} AP: {:>8.4f} -----'.format(VOC_
20             aps.append(ap)
21
22     mAP = np.mean(aps)
23     test_loss = np.mean(losses)
24     if print_total:
25         print('mAP: {:.4f}'.format(mAP))
26         print('Avg loss: {}'.format(test_loss))
27
28     return mAP, test_loss, aps

```

```

1 # plot functions
2 def plot_losses(train, val, test_frequency, num_epochs):
3     plt.plot(train, label="train")
4     indices = [i for i in range(num_epochs) if ((i+1)%test_frequency == 0 or i ==0)]
5     plt.plot(indices, val, label="val")
6     plt.title("Loss Plot")
7     plt.ylabel("Loss")
8     plt.xlabel("Epoch")
9     plt.legend()
10    plt.show()
11
12 def plot_mAP(train, val, test_frequency, num_epochs):
13     indices = [i for i in range(num_epochs) if ((i+1)%test_frequency == 0 or i ==0)]
14     plt.plot(indices, train, label="train")
15     plt.plot(indices, val, label="val")
16     plt.title("mAP Plot")
17     plt.ylabel("mAP")
18     plt.xlabel("Epoch")
19     plt.legend()
20     plt.show()
21

```

▼ Training the network

The simple network you are given as is will allow you to reach around 0.15-0.2 mAP. In this project, you will fine-tune the network to reach higher mAP. Save plots and final test mAP scores as you will be adding these to the writeup.

```

1 def train(classifier, num_epochs, train_loader, val_loader, criterion, optimizer, test_frequency):
2     train_losses = []
3     train_mAPs = []
4     val_losses = []
5     val_mAPs = []
6
7     for epoch in range(1,num_epochs+1):
8         print("Starting epoch number " + str(epoch))
9         train_loss = train_classifier(train_loader, classifier, criterion, optimizer)
10        train_losses.append(train_loss)
11        print("Loss for Training on Epoch " +str(epoch) + " is "+ str(train_loss))
12        if(epoch%test_frequency==0 or epoch==1):
13            mAP_train, _, _ = test_classifier(train_loader, classifier, criterion, F

```

```
14     train_mAPs.append(mAP_train)
15     mAP_val, val_loss, _ = test_classifier(val_loader, classifier, criterion)
16     print('Evaluating classifier')
17     print("Mean Precision Score for Testing on Epoch " +str(epoch) + " is "+str(mAP_val))
18     val_losses.append(val_loss)
19     val_mAPs.append(mAP_val)
20
21 return classifier, train_losses, val_losses, train_mAPs, val_mAPs
```

```
1 classifier = SimpleClassifier().to(device)
2 # You can use this function to reload a network you have already saved previously
3 #classifier.load_state_dict(torch.load('voc_classifier.pth'))
```

```
1 criterion = nn.MultiLabelSoftMarginLoss()
2 optimizer = torch.optim.SGD(classifier.parameters(), lr=0.01, momentum=0.9)
```

```
1 # Training the Classifier
2 num_epochs = 20
3 test_frequency = 5
4
5 classifier, train_losses, val_losses, train_mAPs, val_mAPs = train(classifier, num_e
6
```

```
Loss for Training on Epoch 15 is 0.22407136857509613
----- Class: aeroplane AP: 0.2722 -----
----- Class: bicycle AP: 0.0608 -----
----- Class: bird AP: 0.1066 -----
----- Class: boat AP: 0.2111 -----
----- Class: bottle AP: 0.0804 -----
----- Class: bus AP: 0.0600 -----
----- Class: car AP: 0.1826 -----
----- Class: cat AP: 0.1271 -----
----- Class: chair AP: 0.2060 -----
----- Class: cow AP: 0.0485 -----
----- Class: diningtable AP: 0.1284 -----
----- Class: dog AP: 0.1444 -----
----- Class: horse AP: 0.0782 -----
----- Class: motorbike AP: 0.0528 -----
----- Class: person AP: 0.5101 -----
----- Class: pottedplant AP: 0.0605 -----
----- Class: sheep AP: 0.0459 -----
----- Class: sofa AP: 0.1509 -----
----- Class: train AP: 0.1132 -----
----- Class: tvmonitor AP: 0.0776 -----
```

mAP: 0.1359

Avg loss: 0.22459923753551408

Evaluating classifier

Mean Precision Score for Testing on Epoch 15 is 0.13586376672585931

Starting epoch number 16

Loss for Training on Epoch 16 is 0.224028080701828

Starting epoch number 17

Loss for Training on Epoch 17 is 0.22357989847660065

Starting epoch number 18

Loss for Training on Epoch 18 is 0.22061200439929962

Starting epoch number 19

Loss for Training on Epoch 19 is 0.2209300994873047

Starting epoch number 20

Loss for Training on Epoch 20 is 0.21946430206298828

----- Class: aeroplane AP: 0.3129 -----

```
----- Class: bicycle AP: 0.0734 -----
----- Class: bird AP: 0.1073 -----
----- Class: boat AP: 0.2140 -----
----- Class: bottle AP: 0.0912 -----
----- Class: bus AP: 0.0709 -----
----- Class: car AP: 0.2269 -----
----- Class: cat AP: 0.1345 -----
----- Class: chair AP: 0.2370 -----
----- Class: cow AP: 0.0630 -----
----- Class: diningtable AP: 0.1707 -----
----- Class: dog AP: 0.1389 -----
----- Class: horse AP: 0.1153 -----
----- Class: motorbike AP: 0.0736 -----
----- Class: person AP: 0.5127 -----
----- Class: pottedplant AP: 0.0728 -----
----- Class: sheep AP: 0.0414 -----
----- Class: sofa AP: 0.1775 -----
----- Class: train AP: 0.1517 -----
----- Class: tvmonitor AP: 0.1150 -----
```

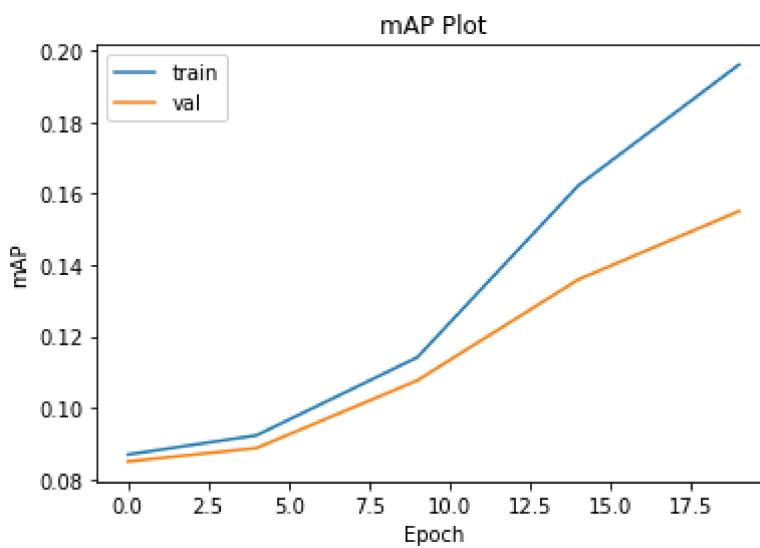
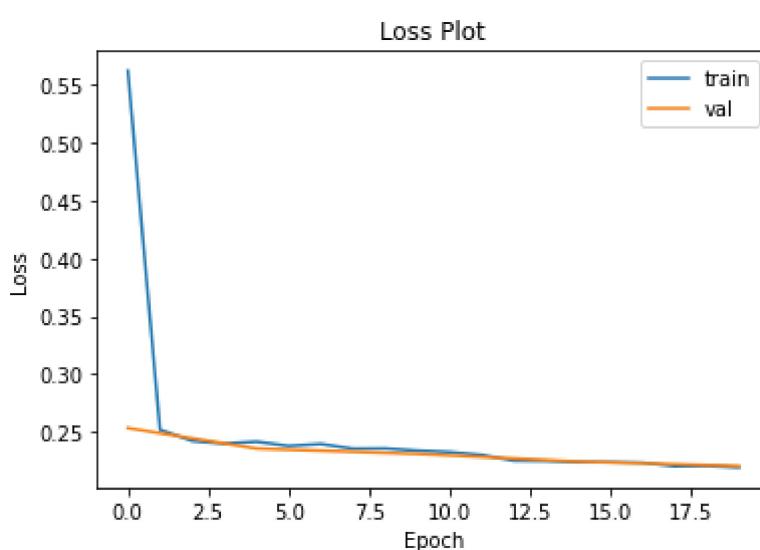
mAP: 0.1550

Avg loss: 0.22070468407051236

Evaluating classifier

Mean Precision Score for Testing on Epoch 20 is 0.155030745622131

```
1 # Compare train and validation metrics
2 plot_losses(train_losses, val_losses, test_frequency, num_epochs)
3 plot_mAP(train_mAPs, val_mAPs, test_frequency, num_epochs)
```



```
1 # Save the classifier network
```

```
2 # Suggestion: you can save checkpoints of your network during training and reload them
```

```
2 # Suggestion: you can save checkpoints of your network during training and reload them
3 torch.save(classifier.state_dict(), './voc_simple_classifier.pth')
```

▼ Evaluate on test set

```
1 ds_test = VocDataset('VOCdevkit_2007/VOC2007test/','test', test_transform)
2
3 test_loader = torch.utils.data.DataLoader(dataset=ds_test,
4                                              batch_size=50,
5                                              shuffle=False,
6                                              num_workers=1)
7
8 # Transforms applied to the testing data
9 test_transform = transforms.Compose([
10     transforms.Resize(227),
11     transforms.CenterCrop(227),
12     transforms.ToTensor(),
13     normalize,
14 ])
15
16 mAP_test, test_loss, test_aps = test_classifier(test_loader, classifier, criterion)
```

```
/content/assignment1/assignment1-part3/voc_dataloader.py:109: VisibleDeprecationWarning: Creating a
    return np.array(names), np.array(labels).astype(np.float32), np.array(box_indices), label_order
----- Class: aeroplane      AP:  0.3025 -----
----- Class: bicycle        AP:  0.0615 -----
----- Class: bird           AP:  0.1073 -----
----- Class: boat           AP:  0.1665 -----
----- Class: bottle          AP:  0.0851 -----
----- Class: bus            AP:  0.0614 -----
----- Class: car             AP:  0.2676 -----
----- Class: cat             AP:  0.1126 -----
----- Class: chair           AP:  0.2408 -----
----- Class: cow             AP:  0.0645 -----
----- Class: diningtable     AP:  0.1223 -----
----- Class: dog              AP:  0.1564 -----
----- Class: horse            AP:  0.1079 -----
----- Class: motorbike        AP:  0.0791 -----
----- Class: person           AP:  0.5287 -----
----- Class: pottedplant      AP:  0.0658 -----
----- Class: sheep            AP:  0.0713 -----
----- Class: sofa             AP:  0.1770 -----
----- Class: train            AP:  0.1716 -----
----- Class: tvmonitor         AP:  0.0955 -----
mAP: 0.1523
Avg loss: 0.2167697401344776
```

```
1 output_submission_csv('my_solution.csv', test_aps)
```

▼ AlexNet Baselines (From Scratch)

AlexNet was one of the earliest deep learning models to have success in classification. In this section we will implement AlexNet as a baseline. Furthermore, we will run an ImageNet-pretrained AlexNet to observe the impact of weight transfer learning.

final test mAP scores as you will be adding these to the writeup.

▼ Running AlexNet

In this section, we train AlexNet from scratch using the same hyperparameters as our previous experiment.

```
1 num_epochs = 20
2 test_frequency = 5
3
4 # Change classifier to AlexNet
5 classifier = torchvision.models.alexnet(pretrained=False)
6 classifier.classifier._modules['6'] = nn.Linear(4096, 21)
7 classifier = classifier.to(device)
8
9 criterion = nn.MultiLabelSoftMarginLoss()
10
11 optimizer = torch.optim.SGD(classifier.parameters(), lr=0.01, momentum=0.9)

1 classifier, train_losses, val_losses, train_mAPs, val_mAPs = train(classifier, num_e
1 plot_losses(train_losses, val_losses, test_frequency, num_epochs)
2 plot_mAP(train_mAPs, val_mAPs, test_frequency, num_epochs)

1 mAP_test, test_loss, test_aps = test_classifier(test_loader, classifier, criterion)
2 print("Test mAP: ", mAP_test)
```

You should notice somewhat poor performance. You could try running AlexNet with an Adam optimizer instead that makes a difference. This experiment is not required for the writeup, but it may show you the importance of the optimizer.

▼ Pretrained AlexNet

Here we look at the impact of pretrained features. This model's weights were trained on ImageNet, which is a large dataset. How do the pretrained features perform on VOC? Why do you think there is such a large difference in performance?

```
1 num_epochs = 20
2 test_frequency = 5
3
4 # Load Pretrained AlexNet
5 classifier = torchvision.models.alexnet(pretrained=True)
6 classifier.classifier._modules['6'] = nn.Linear(4096, 21)
7 classifier = classifier.to(device)
8 optimizer = torch.optim.SGD(classifier.parameters(), lr=0.01, momentum=0.9)

1 classifier, train_losses, val_losses, train_mAPs, val_mAPs = train(classifier, num_e
```